

Maintenance goals of agents in a dynamic environment: Formulation and policy construction[☆]

Chitta Baral^{a,*}, Thomas Eiter^c, Marcus Bjärelund^b, Mutsumi Nakamura^a

^a Department of Computer Science and Engineering, Arizona State University, Tempe, AZ 85287, USA

^b AstraZeneca R&D, S-43183 Mölndal, Sweden

^c Institute of Information Systems, Vienna University of Technology, A-1040 Vienna, Austria

ARTICLE INFO

Article history:

Received 25 September 2005

Received in revised form 3 March 2008

Accepted 18 March 2008

Available online 3 April 2008

Keywords:

Maintenance goals

k -maintainability

Agent control

Computational complexity of agent design

Answer set programming

Horn theories

SAT solving

Discrete event dynamic systems

Self-stabilization

ABSTRACT

The notion of maintenance often appears in the AI literature in the context of agent behavior and planning. In this paper, we argue that earlier characterizations of the notion of maintenance are not intuitive to characterize the maintenance behavior of certain agents in a dynamic environment. We propose a different characterization of maintenance and distinguish it from earlier notions such as *stabilizability*. Our notion of maintenance is more sensitive to a good-natured agent which struggles with an “adversary” environment, which hinders her by unforeseeable events to reach her goals (not in principle, but in case). It has a parameter k , referring to the length of non-interference (from exogenous events) needed to maintain a goal; we refer to this notion as *k-maintainability*. We demonstrate the notion on examples, and address the important but non-trivial issue of efficient construction of maintainability control functions. We present an algorithm which in polynomial time constructs a k -maintainable control function, if one exists, or tells that no such control is possible. Our algorithm is based on SAT Solving, and employs a suitable formulation of the existence of k -maintainable control in a fragment of SAT which is tractable. For small k (bounded by a constant), our algorithm is linear time. We then give a logic programming implementation of our algorithm and use it to give a standard procedural algorithm, and analyze the complexity of constructing k -maintainable controls, under different assumptions such as $k = 1$, and states described by variables. On the one hand, our work provides new concepts and algorithms for maintenance in dynamic environment, and on the other hand, a very fruitful application of computational logic tools. We compare our work with earlier works on control synthesis from temporal logic specification and relate our work to Dijkstra’s notion of self-stabilization and related notions in distributed computing.

© 2008 Elsevier B.V. All rights reserved.

[☆] A preliminary version of the formulation part, entitled “A formal characterization of maintenance goals”, has been presented at AAAI’00, and a preliminary version of the algorithm part entitled “A polynomial time algorithm for constructing k -maintainable policies” has been presented at ICAPS’04. The current version revises and combines both of them with additional elaborations, examples, results, and proofs. The major part of the algorithms was done when Chitta Baral was visiting Vienna University of Technology in 2003. Marcus Bjärelund carried out the major part of his work while he was with the Department of Computer and Information Science of Linköping University.

* Corresponding author.

E-mail addresses: chitta@asu.edu (C. Baral), eiter@kr.tuwien.ac.at (T. Eiter), marcus.bjareland@astrazeneca.com (M. Bjärelund), mutsumi@asu.edu (M. Nakamura).

1. Introduction and motivation

For an agent situated in a static environment, the goal is often to reach one out of several states where certain conditions are satisfied. Such a goal is usually expressed by a formula in propositional or first-order logic. Sometimes the goal requires constraining the path taken to reach one of the states. In that case, the goal can be expressed by a formula in temporal logic [3,7,54].

Our concern in this paper is about agents in a dynamic environment. In that case, things are more complex since the state of the world can change through both actions of the agent and of the environment. The agent's goal in a dynamic environment is then often more than just achieving a desired state, as after the agent has successfully acted to reach a desired state, the environment may change that state. In such a case, a common goal of an agent is to 'maintain' rather than just 'achieve' certain conditions. The goal of maintaining certain conditions (or a set of states that satisfy these conditions) is referred to as *maintenance goals*. Maintenance goals are well-known in the AI literature, e.g., [3,27,41,55,67], and have counterparts in other areas such as in stability theory of discrete event dynamic systems [56,58,61,62,66] and in active databases [17,51]. However, as we argue in this paper, earlier characterizations of maintenance goals are not adequate under all circumstances.

To see what is wrong with earlier definition of maintenance goals, suppose an agent's goal is to maintain a fluent f , i.e., the proposition f should be true. A straightforward attempt¹ to express it using temporal operators is the formula $\Box f$, where \Box is the temporal operator "Always" and $\Box f$ means that f is *true* in all the future states of the world. This is too strong a condition, as maintaining inherently means that things go out of shape and they have to be maintained back to shape. A better temporal logic representation of this goal is thus the formula $\Box \Diamond f$, where \Diamond is the temporal operator "Eventually". Intuitively, the formula $\Box \Diamond f$ is satisfied by an infinite trajectory of states of the form s_0, s_1, s_2, \dots , if at any stage $i \geq 0$, there exists some stage $j \geq i$ such that f is true in s_j . An agent's control is said to satisfy $\Box \Diamond f$ if all trajectories that characterize the evolution of the world due to the environment and the agent's control satisfy $\Box \Diamond f$. At first glance the formula $\Box \Diamond f$ seems to express the goal of maintaining f , as it encodes that if f becomes *false* in any state in the trajectory then it becomes *true* in a later state.

We consider $\Box \Diamond f$ to be also too strong a specification—in many situations—to express the intuitive notion of 'maintaining f ', if we take on a more refined view of the (sometimes nasty) part which the environment might play, which we illustrate by some examples. Suppose f denotes the condition that the Inbox of a customer service department be empty. Here the environment makes f false by adding new requests to the Inbox while the agent tries to make f true by processing the messages in the Inbox and removing them from it. If the agent is diligent in processing the message in the Inbox and makes it empty every chance the agent gets, we would then like to say that agent maintains the Inbox empty. But such a control does not satisfy the formula $\Box \Diamond f$ under all circumstances, because there will be trajectories where the agent is overwhelmed by the environment (flooding the Inbox) and f never becomes *true*.

Another example in support of our intuition behind maintainability is the notion of maintaining the consistency of a database [17,51,68]. When direct updates are made to a database, maintaining the consistency of the database entails the triggering of additional updates that may bring about additional changes to the database so that in the final state (after the triggering is done) the database reaches a consistent state. This does not mean that the database will reach consistency if continuous updates are made to it and it is not given a chance to recover. In fact, if continuous update requests are made we may have something similar to denial of service attacks. In this case we can not fault the triggers saying that they do not maintain the consistency of the database. They do. It is just that they need to be given a *window of opportunity* or a respite from continuous harassment from the environment to bring about the additional changes which are necessary to restore database consistency. The same holds for maintaining a room clean; we can not fault the cleaning person if he or she is continually sent away because the room is being continuously used.

Another example is a mobile robot [15,47] which is asked to 'maintain' a state where there are no obstacles in front of it. Here, if there is a belligerent adversary that keeps on putting an obstacle in front of the robot, there is no way for the robot to reach a state with no obstacle in front of it. But often we will be satisfied if the robot avoids obstacles in its front when it is not continually harassed. Of course, we would rather have the robot take a path that does not have such an adversary, but in the absence of such a path, it would be acceptable if it takes an available path and 'maintains' states where there are no obstacles in front.

The inadequacy of the expression $\Box \Diamond f$ in expressing our intuition about 'maintaining f ' is because $\Box \Diamond f$ is defined on trajectories which do not distinguish between transitions due to agent actions and environment actions. Thus we can not distinguish the cases

- (i) where the agent does its best to maintain f (and is sometimes thwarted by the environment) and can indeed make f true in some (say, k) steps if there is no interference from the environment during those steps; and
- (ii) where the agent really does not even try.

¹ All through the paper we consider the evaluation of linear temporal formulas with respect to all 'valid' trajectories. An alternative approach would be to use a variation of the branching time quantifier A , such as the operator A_π from [9], before the linear temporal formulas. Another alternative approach, referred to as boolean task specification, is used in [27,28,69].

We refer to (i) as *k-maintainability* in this paper. The expression $\Box \Diamond f$ can not express the idea of a *window of opportunity* (or *window of non-interference*) during which an agent can perform the actions necessary for maintaining. In fact, none of the standard notions of temporal logics [20,48], which are defined on trajectories that do not distinguish between the cause behind the transitions (whether they are due to agent's actions or due to the environment),² can express the idea behind *k-maintainability*.

The main contributions of this paper can be summarized as follows.

1. We introduce and formally define the notion of *k-maintainability*, and distinguish it from earlier notions of maintainability, in particular the specification $\Box \Diamond f$ and the similar notion of stabilizability from discrete event dynamic systems.
2. We provide polynomial time algorithms that can construct *k-maintainable* control policies, if one exists. (In the rest of the paper we will refer to 'control policy' simply by 'control'.) Our algorithm is based on SAT Solving, and employs a suitable formulation of the existence of *k-maintainable* control in a tractable fragment of SAT. We then give a logic programming implementation of this method, and finally distill from it a standard procedural algorithm. We briefly discuss earlier approaches to controller synthesis [1,10,19,41,49,60] with respect to temporal logic specifications and compare their complexity with that of our algorithms.
3. We analyze the computational complexity of constructing *k-maintainable* controls, under different settings of the environment and the windows of opportunity open to the agent, as well as under different forms of representation. We show that the problem is complete for **P**TIME in the standard setting, where the possible states are enumerated, and complete for **EX**PTIME in a STRIPS-style setting where states are given by value assignments to fluents. Furthermore, we elucidate the impact of the different factors and show, by our proofs of the hardness results, that the full problem complexity is inherent already to certain restricted cases.

Overall, our work not only provides new concepts and algorithms for realizing maintenance of an agent in dynamic environment, but also illustrates a very fruitful application of computational logic tools.

The rest of this paper is organized as follows. In Section 2 we present the background definitions of a system with an agent in an environment and define the notions of stability and stabilizability. In Section 3 we describe an example of a system with two buffers. We use this example for illustrating the concepts of stabilizability and *k-maintainability*, which is formally defined in Section 4. In Section 5 we present our algorithms for constructing *k-maintaining* controls, based on SAT Solving as well as a genuine algorithm extracted from it. In Section 6 we present an encoding for computing a control function using a logic programming engine and devote Section 7 to complexity analysis. Finally, in Section 8 we present some experimental results, discuss related work, and outline some future directions.

2. Background: Systems, goals, control, stability and stabilizability

In this paper, we are concerned with goal-directed agents in a dynamic world. Such agents can perform actions that change the state of the world. Because of the dynamic nature of the world, certain changes can happen to the state of the world beyond the control of an agent. The agent's job is thus to make the world evolve in a way coherent with a goal assigned to it. As for the agent control, we adopt here that an agent follows a Markovian control policy to do its job; that is, its control is a function from the set of states to the set of actions, detailed as follows.

Definition 1 (System). A system is a quadruple $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$, where

- \mathcal{S} is a finite set of system states;
- \mathcal{A} is a finite set of actions, which is the union of the set of agents actions, \mathcal{A}_{ag} , and the set of environmental actions, \mathcal{A}_{env} ;
- $\Phi : \mathcal{S} \times \mathcal{A} \rightarrow 2^{\mathcal{S}}$ is a non-deterministic transition function that specifies how the state of the world changes in response to actions; and
- $\text{poss} : \mathcal{S} \rightarrow 2^{\mathcal{A}}$ is a function that describes which actions are possible to take in which states.

The above notion of system is used in the discrete event dynamic systems community, for instance in [56,58,61,62,66]. In practice, the functions Φ and poss are required to be effectively (and efficiently) computable, and they may often be specified in a representation language such as in [32,34,63]. The possibility of an action has different meaning depending on whether it is an agent's action or whether it is an environmental action. In case of an agent's action, it is often dictated by the policy followed by the agent. For environmental actions, it encodes the various possibilities that are being accounted for in the model. We tacitly assume here that possible actions lead always to some successor state, i.e., the axiom that $\Phi(s, a) \neq \emptyset$ whenever $a \in \text{poss}(s)$ holds for any state s and action a , is satisfied by any system.

² If one distinguishes the cause behind transitions, then temporal logic can indeed be used to express maintainability. We discuss this further in Section 4.1.




Fig. 1. $S = \{b, c, d, f, g, h\}$, $\mathcal{A} = \{a, a', e\}$, $\mathcal{A}_{ag} = \{a, a'\}$, $\mathcal{A}_{env} = \{e\}$, Φ is as shown.

An example of a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, where $S = \{b, c, d, f, g, h\}$, $\mathcal{A} = \{a, a', e\}$, and the transition function Φ is shown in Fig. 1, where $s' \in \Phi(s, a)$ iff an arc $s \rightarrow s'$ labeled with a is present and $\text{poss}(s)$ are all actions that label arcs leaving s . Notice that in this example, $\Phi(s, a)$ is *deterministic*, i.e., $\Phi(s, a)$ is a singleton if nonempty.

The evolution of the world with respect to a system is characterized by the following definition.

Definition 2 (Trajectory). Given a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, an alternating infinite sequence of states and actions $s_0, a_1, s_1, a_2, \dots, s_k, a_{k+1}, s_{k+1}, \dots$ is said to be a *trajectory consistent with A*, if $s_{k+1} \in \Phi(s_k, a_{k+1})$, and $a_{k+1} \in \text{poss}(s_k)$.

The above notion of trajectory does not require that agent actions and environment actions be interleaved as is done in formulations of games. It allows one agent action to be followed by multiple environment actions and vice-versa, as in real worlds we do not have an arbiter who can enforce the alternation of the agent and environment actions. The notion assumes that the time points are fine enough that at any point only one action can occur. Thus, it does not allow explicit occurrence of agent and environment actions at the same time.

A common restriction on how the world evolves is defined using the notion of *stability*. The following definition of stability is adapted from [56] and has its origin in control theory and discrete event dynamic systems [56,58,61,62].

Definition 3 (Stable state 1). Given a system $A = (S, \mathcal{A}, \Phi, \text{poss})$ and a set of states E , a state s is said to be *stable* in A w.r.t. E if all trajectories consistent with A and starting from s visit E infinitely often. A set of states S is *stable* with respect to E if all states in S are stable with respect to E .

We say $A = (S, \mathcal{A}, \Phi, \text{poss})$ is a *stable system*, if all states in S are stable in A with respect to E .

Although the above definition of stability is with respect to a set of states E , it can be easily adapted to a propositional formula φ that can be evaluated at the states of system A . In that case $E = \{s \in S \mid A, s \models \varphi\}$, where $A, s \models \varphi$ denotes that s in A satisfies (in the propositional logic sense) φ . Thus, E is the set of states s at which φ is satisfied.

An alternative approach to characterize the evolution of states is through temporal operators. Some of the important temporal operators talking about the future are (cf. [30,48]): Next (\bigcirc), Always (\Box), Eventually (\Diamond), and Until (\mathcal{U}). Their meaning with respect a trajectory $\tau = s_0, a_1, s_1, \dots, s_k, a_{k+1}, s_{k+1}, \dots$ is defined as follows.

Let (τ, j) , for $j \geq 0$, denote the remainder of τ starting at s_j ; then

- $(\tau, j) \models p$ iff p is true in s_j , for any proposition p ;
- $(\tau, j) \models \bigcirc \phi$ iff $(\tau, j+1) \models \phi$;
- $(\tau, j) \models \Box \phi$ iff $(\tau, k) \models \phi$, for all $k \geq j$.
- $(\tau, j) \models \Diamond \phi$ iff $(\tau, k) \models \phi$, for some $k \geq j$.
- $(\tau, j) \models \phi_1 \mathcal{U} \phi_2$ iff there exists $k \geq j$ such that $(\tau, k) \models \phi_2$ and for all $i, j \leq i < k$, $(\tau, i) \models \phi_1$.

The standard Boolean connectives \wedge , \vee , and \neg are defined as usual. An alternative definition of stability can then be given as follows:

Definition 4 (Stable state 2). Given a system $A = (S, \mathcal{A}, \Phi, \text{poss})$ and an objective formula φ (i.e., without temporal operators), let $E_\varphi = \{s \in S \mid A, s \models \varphi\}$. A state s is then said to be *stable* in A w.r.t. E if for all trajectories τ of the form $\tau = s_0, a_1, s_1, \dots, s_k, a_{k+1}, s_{k+1}, \dots$ (where $s_0 = s$) consistent with A , it holds that $(\tau, 0) \models \Box \Diamond \varphi$.

In fact, this definition is equivalent to Definition 3. The advantage of using temporal operators, as in the above definition, instead of Definition 3 is that the former allows us to specify a larger class of goals and build on top of the notion of stability. For example, a notion similar to stability, referred to as a *response property* [48], is of the form $\Box(p \rightarrow \Diamond q)$.

2.1. Stabilizability

The notion of stability is defined with respect to a system and the evolution of the world consistent with the system. When we focus on an agent and its ability to make a system stable, we need a notion of *stabilizability* which intuitively means that there exists a control policy which the agent can use to fashion a stable system.

Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$, when discussing stabilizability of the system, we need to consider the following additional aspects:

- the set of actions \mathcal{A}_{ag} which the agent is capable of executing in principle (where $\mathcal{A}_{ag} \subseteq \mathcal{A}$);
- the set of *exogenous actions* that may occur in the state s , beyond the agent's control, modeled by a function $\text{exo} : \mathcal{S} \rightarrow 2^{\mathcal{A}_{env}}$, where $\text{exo}(s) \subseteq \text{poss}(s)$ for each state s (recall that \mathcal{A}_{env} are the environmental actions). We call any such exo an *exogenous function*.

Intuitively, given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$, \mathcal{A}_{ag} , exo , and E , a state s is stabilizable with respect to E , if we are able to find a policy or *control function* such that it makes the resulting system stable and the agent starting from s following that policy will not reach a state where no further actions are possible.

The last condition is referred to as *aliveness*. It is formally defined by the following two definitions, the first of which defines the set $R(A, s)$ of states that can be reached from s in the system A .

Definition 5 (Closure). Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$ and a state s , $R(A, s) \subseteq \mathcal{S}$ is the smallest set of states that satisfying the following conditions:

1. $s \in R(A, s)$,
2. If $s' \in R(A, s)$, and $a \in \text{poss}(s')$, then $\Phi(s', a) \subseteq R(A, s)$.

For any set of states $S \subseteq \mathcal{S}$, the closure of A w.r.t. S is defined by

$$\text{Closure}(S, A) = \bigcup_{s \in S} R(A, s).$$

Example 1. In the system A in Fig. 1, we have that $R(A, d) = \{d, h\}$ and $R(A, f) = \{f, g, h\}$, and therefore $\text{Closure}(\{d, f\}, A) = \{d, f, g, h\}$. This is illustrated in Fig. 2.

Note that $\text{Closure}(S, A)$ satisfies the Kuratowski closure axioms [43], that is, $\text{Closure}(\emptyset, A) = \emptyset$; $S \subseteq \text{Closure}(S, A)$; $\text{Closure}(\text{Closure}(S, A), A) = \text{Closure}(S, A)$; and, moreover, $\text{Closure}(S_1 \cup S_2, A) = \text{Closure}(S_1, A) \cup \text{Closure}(S_2, A)$. Furthermore, $\Phi(s, a) \subseteq \text{Closure}(S, A)$ holds for each state $s \in \text{Closure}(S, A)$ and $a \in \text{poss}(s)$.

Definition 6 (Aliveness). Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$ and a state s , we say s is *alive* if $\text{poss}(s') \neq \emptyset$, for all $s' \in R(A, s)$. We say $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$ is *alive* if all states in \mathcal{S} are alive.

The notion of control function is formally defined as follows.

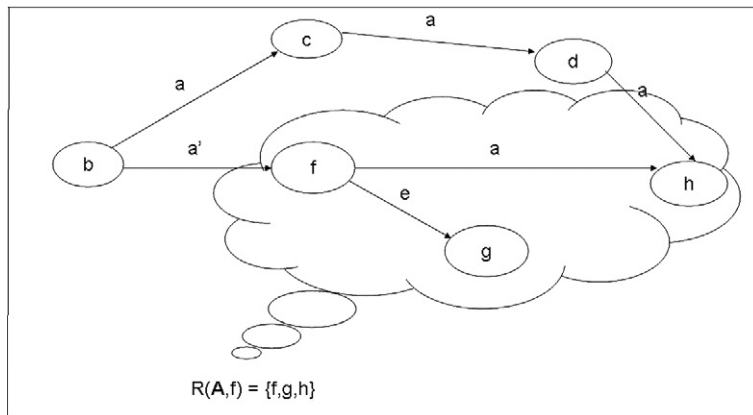


Fig. 2. $R(A, d) = \{d, h\}$ and $R(A, f) = \{f, g, h\}$, and $\text{Closure}(\{d, f\}, A) = \{d, f, g, h\}$.




Fig. 3. Policy K is doing a in states b, c, d and f ; $\text{poss}(b) = \{a, a'\}$; $\text{poss}_{K, \text{exo}}(b) = \{a\}$; $\text{Closure}(\{b, c\}, A) = \{b, c, d, f, g, h\}$; $\text{Closure}(\{b, c\}, A_{K, \text{exo}}) = \{b, c, d, h\}$.

Definition 7 (Control). Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$ and a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$ of agent actions, a *control function* for A w.r.t. \mathcal{A}_{ag} is a partial function

$$K : \mathcal{S} \rightarrow \mathcal{A}_{ag},$$

such that $K(s) \in \text{poss}(s)$ whenever $K(s)$ is defined.³

We are now ready to formally define the notion of stabilizability.

Definition 8 (Stabilizability). Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$, a function exo as above, and a set of states E , we say that $s \in \mathcal{S}$ is *stabilizable* with respect to E , if there exists a control function $K : \mathcal{S} \rightarrow \mathcal{A}_{ag}$ for A w.r.t. \mathcal{A}_{ag} with the following properties:

1. s is stable with respect to E in the system $A_{K, \text{exo}} = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss}_{K, \text{exo}})$, where, for any state s' , $\text{poss}_{K, \text{exo}}(s') = \{K(s')\} \cup \text{exo}(s')$; and
2. s is alive in $A_{K, \text{exo}}$.

A set of states $S \subseteq \mathcal{S}$ is *stabilizable with respect to E* , if there is a control function K for \mathcal{A} w.r.t. \mathcal{A}_{ag} such that every state $s \in S$ is stabilizable with respect to E witnessed by K .

Having provided this definition, we shall illustrate it on an elaborated example in the next section, where we describe an intuitive control function for the management of two finite buffers.

Before closing this section, we introduce for later use the notion of a non-deterministic control.

Definition 9 (Non-deterministic control). Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$ and a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$ of agent actions, a partial function $K : \mathcal{S} \rightarrow 2^{\mathcal{A}_{ag}}$ such that $K(s) \subseteq \text{poss}(s)$ and $K(s) \neq \emptyset$ whenever $K(s)$ is defined, is called *non-deterministic control* for A w.r.t. \mathcal{A}_{ag} .

Informally, a non-deterministic control leaves the agent a choice to execute one out of several actions. It is an envelope for multiple control functions, which result by refining K to some arbitrary action in $K(s)$ whenever $K(s)$ is defined; the notion of stabilizability is defined similar as for control functions, with the only change that in $A_{K, \text{exo}}$, we set $\text{poss}_{K, \text{exo}}(s') = K(s') \cup \text{exo}(s')$ in place of $\text{poss}_{K, \text{exo}}(s') = \{K(s')\} \cup \text{exo}(s')$.

The following proposition is immediate.

Proposition 1. Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$, and a function exo , a set of states $S \subseteq \mathcal{S}$ is stabilizable w.r.t. a set of states $E \subseteq \mathcal{S}$ under a control function K for A w.r.t. \mathcal{A}_{ag} iff S is stabilizable w.r.t. E under a non-deterministic control K^+ for A w.r.t. \mathcal{A}_{ag} . Furthermore, each such K is a refinement of some K^+ with this property (i.e., for each s , $K(s) \in K^+(s)$ and $K(s)$ is defined iff $K^+(s)$ is defined), and each refinement K of K^+ is a control function witnessing stabilizability of S w.r.t. E .

3. Example scenario: Two finite buffers

In this section, we introduce an example which we will use in illustrating the notion of stabilizability and also other concepts in some of the later sections of the paper.

³ In the planning literature, in Markov Decision Planning a control function is also called a (control) policy, which is usually assumed to be a total function [35]; in Model-Based Planning, [18], it is called a deterministic state-action table, and non-deterministic controls (introduced below) are called (non-deterministic) state-action tables; [35] refers to them also as (non-deterministic) plans or policies, respectively.

We imagine a system with two finite buffers, b_1 and b_2 , where objects are added to b_1 in an uncontrollable way. An agent moves objects from b_1 to b_2 and processes them there. When an object has been processed, it is automatically removed from b_2 . This is a slight modification of a finite buffer example from [58] and generalizes problems such as ftp agents maintaining a clean ftp area by moving submitted files to other directories, or robots moving physical objects from one location to another.

In our framework, we shall describe a system A_b which models this scenario. For simplicity, we assume that the agent has three control actions \mathbf{M}_{12} that moves an object from b_1 to b_2 (if such an object exists), the opposite action, \mathbf{M}_{21} that moves an object from b_2 to b_1 , and **Proc** that processes and removes an object in b_2 . There is one exogenous action, *Ins*, that inserts an object into buffer b_1 . The capacities of b_1 and b_2 are assumed to be equal.

Let us assume that the control goal of this system is to keep b_1 empty. Then, the system is not stabilizable, since objects can be continually inserted before the agent has a chance to empty the buffer. However, if no insertions are performed for a certain window of non-interference, the agent can always empty b_1 . This implies that the system is maintainable but not stabilizable. We now make the above argument explicit by using a concrete instance of A_b .

Example 2 (Buffer example). We assume that the maximum capacity of the buffers b_1 and b_2 is 3. The components of $A_b = (\mathcal{S}_b, \mathcal{A}_b, \Phi_b, \text{poss}_b)$ are then as follows.

- We model every state by the current number of objects in b_1 and b_2 . That is, a state s is identified by a pair of integers $\langle i, j \rangle$ where i denotes the number of objects in b_1 and j the number of objects in b_2 . With the maximum capacity of 3, the set of states, \mathcal{S}_b , consists of $4 \times 4 = 16$ states and is given by

$$\mathcal{S}_b = \{0, 1, 2, 3\} \times \{0, 1, 2, 3\}.$$

- The set of actions is $\mathcal{A}_b = \{\mathbf{M}_{12}, \mathbf{M}_{21}, \mathbf{Proc}, \text{Ins}\}$.
- We assume that the transition function Φ_b is deterministic, i.e., $|\Phi_b(s, a)| \leq 1$, defined as follows, where we write $\Phi_b(s, a) = s'$ for $\Phi_b(s, a) = \{s'\}$. For every $i, j \in \{0, \dots, 3\}$, let

$$\Phi_b(\langle i, j \rangle, \mathbf{M}_{12}) = \langle i - 1, j + 1 \rangle,$$

$$\Phi_b(\langle i, j \rangle, \mathbf{M}_{21}) = \langle i + 1, j - 1 \rangle,$$

$$\Phi_b(\langle i, j \rangle, \mathbf{Proc}) = \langle i, j - 1 \rangle,$$

$$\Phi_b(\langle i, j \rangle, \text{Ins}) = \langle i + 1, j \rangle,$$

where addition and subtraction are modulo 3, and in all other cases $\Phi_b(s, a) = \emptyset$.

- The enabling function, poss_b , is defined by

$$\mathbf{M}_{12} \in \text{poss}_b(\langle i, j \rangle) \quad \text{iff} \quad i \geq 1 \text{ and } j \leq 2,$$

$$\mathbf{M}_{21} \in \text{poss}_b(\langle i, j \rangle) \quad \text{iff} \quad i \leq 2 \text{ and } j \geq 1,$$

$$\mathbf{Proc} \in \text{poss}_b(\langle i, j \rangle) \quad \text{iff} \quad j \geq 1,$$

$$\text{Ins} \in \text{poss}_b(\langle i, j \rangle) \quad \text{iff} \quad i \leq 2.$$

It is easy to see that for $S = \{\langle 0, 0 \rangle\}$ (no objects in the buffers) and $E = \{\langle 0, 0 \rangle, \langle 0, 1 \rangle, \langle 0, 2 \rangle, \langle 0, 3 \rangle\}$ (that is, we want to keep b_1 empty) S is not stabilizable w.r.t. E , since the exogenous action *Ins* can always interfere in the task of bringing the system back to E . For example, consider the control K_b defined as follows:

$$K_b(\langle i, j \rangle) = \mathbf{M}_{12} \quad \text{when } i \geq 1 \text{ and } j < 3, \text{ and}$$

$$K_b(\langle i, j \rangle) = \mathbf{Proc} \quad \text{when } (i = 0 \text{ and } j \geq 1) \text{ or } j = 3.$$

Intuitively, the above control directs the transfer of objects from buffer 1 to 2 whenever possible, and if that is not possible it directs processing of objects in buffer 2 if that is possible. In Fig. 4, which shows the transition diagram between states, the transitions by the control K_b are marked with \mathbf{M}_{12} and **Proc**.

Consider the following trajectory consistent with the control system $A_{K, \text{exo}} = (\mathcal{S}_b, \mathcal{A}_b, \Phi_b, \text{poss}_{b, K_b, \text{exo}})$:

$$\tau = \langle 0, 0 \rangle, \text{Ins}, \langle 1, 0 \rangle, \text{Ins}, \langle 2, 0 \rangle, \mathbf{M}_{12}, \langle 1, 1 \rangle, \text{Ins}, \langle 2, 1 \rangle, \mathbf{M}_{12}, \langle 1, 2 \rangle, \text{Ins}, \langle 2, 2 \rangle, \mathbf{M}_{12}, \langle 1, 3 \rangle, \mathbf{Proc}.$$

It consists of a prefix $\langle 0, 0 \rangle, \text{Ins}, \dots, \mathbf{M}_{12}$ and a cycle $\langle 1, 2 \rangle, \dots, \mathbf{Proc}$. In τ , no state in E is ever reached after the starting state $\langle 0, 0 \rangle$. Similar trajectories can be found for any control and hence S is not stabilizable with respect to E .

On the other hand, $S = \{\langle 0, 0 \rangle\}$ is stabilizable w.r.t. $E' = \{0, 1, 2\} \times \{0, 1, 2, 3\}$ (that is, we want to have at most two objects in b_1 at any time): Following K_b we can go from any of the states in $\mathcal{S}_b \setminus E' = \{\langle 3, 0 \rangle, \langle 3, 1 \rangle, \langle 3, 2 \rangle, \langle 3, 3 \rangle\}$ to E' with the execution of at most two control actions, while no exogenous actions are possible for those states.




Fig. 4. The transition diagram of the buffer system A_b for the concrete instance (buffer capacity 3).

4. Limited interference and k -maintainability

As we mentioned in Section 1, our main intuition behind the notion of maintainability is that maintenance becomes possible only if there is a window of non-interference from the environment during which maintenance is performed by the agent. In other words, an agent k -maintains a condition c if its control (or its reaction) is such that if we allow it to make the controlling actions without interference from the environment for at least k steps, then it gets to a state that satisfies c within those k steps.

Our definition of maintainability has the following parameters:

- (i) a system $A = (S, \mathcal{A}, \Phi, \text{poss})$,
- (ii) a set of initial states S that the system may be initially in,
- (iii) a set of desired states E that we want to maintain,
- (iv) a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$ of agent actions,
- (v) a function $\text{exo} : S \rightarrow 2^{\mathcal{A}_{env}}$ detailing exogenous actions, such that $\text{exo}(s) \subseteq \text{poss}(s)$, and
- (vi) a control function K (mapping a relevant part of S to \mathcal{A}_{ag}) such that $K(s) \in \text{poss}(s)$.

The next step is to formulate when the control K maintains E assuming that the system is initially in one of the states in S . For that, we require that if in the system $A_{K, \text{exo}} = (S, \mathcal{A}, \Phi, \text{poss}_{K, \text{exo}})$, where $\text{poss}_{K, \text{exo}}(s) = \{K(s)\} \cup \text{exo}(s)$ restricts the agent actions to the control K , the agent is in a state s that has been reached from any state in S (i.e., $s \in \text{Closure}(S, A_{K, \text{exo}})$), then given a window of non-interference from exogenous actions, it must get into some desired state during that window. One of the importance of using the notion of closure here is that one can focus only on a possibly smaller set of states, rather than all the states, thus limiting the possibility of an exponential blow-up—as warned in [36]—of the number of control rules.

Now a next question might be: Suppose the above condition of maintainability is satisfied, and while the control is leading the system towards a desired state, an exogenous action happens and takes the system off that path. What then? The answer is that the state the system will reach after the exogenous action will be a state from the closure. Thus, if the system is then left alone (without interference from exogenous actions) it will be again on its way to a desired state. So in our notion of maintainability, the control is always taking the system towards a desired state, and after any disturbance from an exogenous action, the control again puts the system back on a path to a desired state.

We define the notion of unfolding a control as follows.

Definition 10 ($\text{Unfold}_k(s, A, K)$). Let $A = (S, \mathcal{A}, \Phi, \text{poss})$ be a system, let $s \in S$, and let K be a control for A . Then $\text{Unfold}_k(s, A, K)$ is the set of all sequences $\sigma = s_0, s_1, \dots, s_l$ where $l \leq k$ and $s_0 = s$ such that $K(s_j)$ is defined for all $j < l$, $s_{j+1} \in \Phi(s_j, K(s_j))$, and if $l < k$, then $K(s_l)$ is undefined.

Intuitively, an element of $\text{Unfold}_k(s, A, K)$ is a sequence of states of length at most $k + 1$ that the system may go through if it follows the control K starting from the state s . If the length of the sequence is less than $k + 1$, it means that at the last state of the sequence K is undefined and thus the sequence can not unfold further.

Fig. 5 illustrates this.

The above definition of $\text{Unfold}_k(s, A, K)$ is easily extended to the case when K is a non-deterministic control, meaning $K(s)$ is a set of actions instead of a single action. In that case, we overload Φ and for any set of actions a^* , define $\Phi(s, a^*) = \bigcup_{a \in a^*} \Phi(s, a)$.




Fig. 5. Let K be the policy of doing action a in states b, c, d and f . $Unfold_3(b, A, K) = \{\langle b, c, d, h \rangle, \langle b, g \rangle\}$ and $Unfold_3(c, A, K) = \{\langle c, d, h \rangle\}$.




Fig. 6. Let $S = \{b\}$ and $E = \{h\}$. A 3-maintainable policy with respect to them will be to do a in states b, c and d .

We now define the notion of k -maintainability, which can be used to verify the correctness of a control.

Definition 11 (k -maintainability). Given a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, a set of agents action $\mathcal{A}_{ag} \subseteq \mathcal{A}$, and a specification of exogenous action occurrence exo , we say that a control⁴ K for A w.r.t. \mathcal{A}_{ag} k -maintains $S \subseteq S$ with respect to $E \subseteq S$, where $k \geq 0$, if for each state $s \in \text{Closure}(S, A_{K,exo})$ and each sequence $\sigma = s_0, s_1, \dots, s_l$ in $Unfold_k(s, A, K)$ with $s_0 = s$, it holds that $\{s_0, \dots, s_l\} \cap E \neq \emptyset$.

We say that a set of states $S \subseteq S$ (resp. A , if $S = S$) is k -maintainable, $k \geq 0$, with respect to a set of states $E \subseteq S$, if there exists a control K which k -maintains S w.r.t. E . K is then referred to as the witnessing control function. Furthermore, S (resp. A) is called *maintainable* w.r.t. E , if S (resp. A) is k -maintainable w.r.t. E for some $k \geq 0$.

In the following we will omit explicit mention of \mathcal{A}_{ag} , S , and E for control functions and maintainability if they are clear from the context.

Intuitively, the condition $\{s_0, s_1, \dots, s_l\} \cap E \neq \emptyset$ above means that we can get from a state s_0 outside E to a state in E within at most k transitions—where each transition is dictated by the control K —if the world were to unfold as in s_0, s_1, \dots, s_l , where $s_0 = s$. In particular, 0-maintainability means that the agent has nothing to do: after any exogenous action happening, the system will be in a state from E . Therefore, a trivial control K will do which is undefined on every state.

Note that in the above definition we no longer require aliveness. If a non-alive state is reached while unfolding, the unfolding stops there and the definition of k -maintainability requires that a goal state (from E) is reached by then.

Example 3. Reconsider the system A in Fig. 1. Let us assume that $\mathcal{A}_{ag} = \{a, a'\}$, that $exo(s) = \{e\}$ iff $s = f$ and that $exo(s) = \emptyset$ otherwise. Suppose now that we want a 3-maintainable control policy for $S = \{b\}$ w.r.t. $E = \{h\}$. Clearly, such a control policy K is to take a in b, c , and d . Indeed, $\text{Closure}(\{b\}, A_{K,exo}) = \{b, c, d, h\}$ and $Unfold_3(b, A, K) = \{\langle b, c, d, h \rangle\}$, $Unfold_3(c, A, K) = \{\langle c, d, h \rangle\}$, and $Unfold_3(d, A, K) = \{\langle d, h \rangle\}$; furthermore, each sequence contains h . See Fig. 6.

Suppose now, as shown in Fig. 7, there is an exogenous action e that can take from c to f . Then, no k -maintainable control policy for $S = \{b\}$ w.r.t. $E = \{h\}$ exists for any $k \geq 0$. Indeed, the agent can always end up in the dead-end g . If, however, in addition $\Phi(g, a') = \{f, h\}$ and $a' \in \text{poss}(g)$, a 3-maintainable control policy K is $K(s) = a$ for $s \in \{b, c, d, f\}$ and $K(g) = a'$.

⁴ Note that here only $K(s)$ for $s \in \text{Closure}(S, A_{K,exo})$ is of relevance. For all other s , $K(s)$ can be arbitrary or undefined.




Fig. 7. Let $S = \{b\}$ and $E = \{h\}$. No 3-maintainable policy with respect to them exists.

Example 4 (Buffer example (cont'd)). Earlier we showed that in A_b , $S = \{(0, 0)\}$ is not stabilizable w.r.t. $E = \{(0, 0), \langle 0, 1 \rangle, \langle 0, 2 \rangle, \langle 0, 3 \rangle\}$. Thus, we might ask whether S is at least maintainable w.r.t. E ? The answer is positive: For the worst case system state, $\langle 3, 3 \rangle$, a control can move the system to $\langle 3, 0 \rangle$ (by three transitions executing **Proc**) without interfering occurrences of exogenous actions. If there then are three further transitions without interference, the control can apply \mathbf{M}_{12} three times and effect the state $\langle 0, 3 \rangle$. This implies that S is 6-maintainable w.r.t. E . We can, with a similar argument show that A is 9-maintainable w.r.t. $\{(0, 0)\}$. A similar argument can be made with respect to the control K_b of Example 2.

However, we have that A is *not* maintainable w.r.t., for example, $\{(0, 3)\}$ (Since we cannot go from, for example, $\{(0, 0)\}$, to $\{(0, 3)\}$ with control actions only).

As the above example points out, it is possible that S is maintainable but not stabilizable with respect to E . The converse is also possible. In other words, in certain cases we may have a system where a given S is stabilizable with respect to a set E , but yet is not maintainable. This happens when every path between a state in S and a state in E involves at least one exogenous action. In that case the agent, who does not have control over the exogenous actions, can not on its own make the transition from a state in S to a state in E . However, often for each exogenous action there are equivalent (in terms of effects) agent actions. In that case, any stabilizable system is also maintainable.

We note the following monotonicity property of k -maintainability, which is an easy consequence of the definition:

Proposition 2. Suppose that for a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, a set of agents action $\mathcal{A}_{ag} \subseteq \mathcal{A}$, and a specification of exogenous action occurrence exo , the control function K k -maintains $S \subseteq S$ w.r.t. $E \subseteq S$. Then, K also k -maintains any set $S' \subseteq \text{Closure}(S, A_{K, \text{exo}})$ with respect to any set $E' \supseteq E$.

4.1. An alternative characterization of k -maintainability

The characterization of stability and stabilizability in Section 2 is based on imposing conditions on trajectories obtained from the transition graph of a system. Such a characterization has the advantage that it is amenable to developing temporal operators that can express more general conditions.

In contrast, the definition of maintainability in Definition 11 is not based on trajectories. Nonetheless, one can give an alternative characterization based on trajectories, which we do next. To bridge from finite trajectories (which are relevant with respect to maintainability), to infinite ones as in Definition 2, we consider for each system $A = (S, \mathcal{A}, \Phi, \text{poss})$ an extension, A^∞ , which results by adding a fresh environmental action a_{nop} such that in A^∞ , for each state s we have $\Phi(s, a_{nop}) = \{s\}$ and $a_{nop} \in \text{poss}(s)$ if $\text{poss}(s) = \emptyset$ in A . Informally, A^∞ adds infinite loops to halting states of A .

Proposition 3. Given a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, a set of agents action $\mathcal{A}_{ag} \subseteq \mathcal{A}$, a specification of exogenous action occurrence exo , and a set of states E , a set of states S is k -maintainable with respect to E , $k \geq 0$, if and only if there exists a control K for A w.r.t. \mathcal{A}_{ag} such that for each state s in S and every trajectory of form $\tau = s_0, a_1, s_1, a_2, \dots, a_j, s_j, a_{j+1}, \dots$ consistent with $A_{K, \text{exo}}^\infty$ and $s_0 = s$, it holds that $\{a_{i+1}, \dots, a_{i+k}\} \subseteq \mathcal{A}_{ag}$ or $a_{i+k} = a_{nop}$ for some $i \geq 0$ implies that $\{s_i, \dots, s_{i+k}\} \cap E \neq \emptyset$.

Proof. For the only if direction, suppose that S is k -maintainable w.r.t. E , witnessed by the control function K . Let $s \in S$ and $\tau = s_0, a_1, s_1, a_2, \dots, a_j, s_j, a_{j+1}, \dots$ be consistent with $A_{K, \text{exo}}^\infty$ such that $s_0 = s$ and $\{a_{i+1}, \dots, a_{i+k}\} \subseteq \mathcal{A}_{ag}$ or $a_{i+k} = a_{nop}$, for some $i \geq 0$. Then, we have $s_i \in \text{Closure}(S, A_{K, \text{exo}}^\infty)$. If $k = 0$, then since K is a witnessing control, we have $s_i \in E$, and thus $\{s_i, s_{i+1}, \dots, s_{i+k}\} \cap E \neq \emptyset$ holds. Consider thus $k > 0$. If $a_{i+k} \in \mathcal{A}_{ag}$ (which implies $\{a_{i+1}, \dots, a_{i+k}\} \subseteq \mathcal{A}_{ag}$), then the sequence $s_i, s_{i+1}, \dots, s_{i+k}$ belongs to $\text{Unfold}_k(s_i, A, K)$. Since K is a witnessing control function, we again have $\{s_i, s_{i+1}, \dots, s_{i+k}\} \cap E \neq \emptyset$. Otherwise, if $a_{i+k} = a_{nop}$, let $l \geq 1$ be the least index such that $a_l = a_{nop}$. By definition of $A_{K, \text{exo}}^\infty$, we have that $K(s_{l-1})$ is undefined. Hence, the sequence $\sigma = s_{l-1}$ belongs to $\text{Unfold}_k(s_{l-1}, A, K)$. Since K is a control, it follows that $s_{l-1} \in E$. Since $s_j = s_{l-1}$ for each $j \geq l$, and in particular $s_{i+k} = s_{l-1}$, it follows again that $\{s_i, s_{i+1}, \dots, s_{i+k}\} \cap E \neq \emptyset$. This proves the only if direction.

Conversely, suppose K is a control for A w.r.t. \mathcal{A}_{ag} such that for each $s \in S$ and trajectory $\tau = s_0, a_1, s_1, a_2, \dots, a_j, s_j, a_{j+1}, \dots$ consistent with $A_{K,exo}^\infty$ and $s_0 = s$, it holds that $\{a_{i+1}, \dots, a_{i+k}\} \subseteq \mathcal{A}_{ag}$ or $a_{i+k} = a_{nop}$ for some $i \geq 0$ implies that $\{s_i, s_{i+1}, \dots, s_{i+k}\} \cap E \neq \emptyset$. We claim that K witnesses k -maintainability of S w.r.t. E . Towards a contradiction, suppose the contrary. Hence, it follows from the definition of $A_{K,exo}^\infty$, that there is some state $s \in S$ and trajectory $\tau = s_0, a_1, s_1, a_2, \dots, a_j, s_j, a_{j+1}, \dots$ consistent with $A_{K,exo}^\infty$ and $s_0 = s$, such that for some $j \geq 0$ we have $s_j \in \text{Closure}(S, A_{K,exo}^\infty)$ and $s_j, s_{j+1}, \dots, s_{j+l}$ is in $\text{Unfold}_k(s_j, A, K)$, where $l \leq k$, but $E \cap \{s_j, \dots, s_{j+l}\} = \emptyset$.

By definition of $\text{Unfold}_k(s_j, A, K)$, we have that $\{a_{j+1}, \dots, a_{j+l-1}\} \subseteq \mathcal{A}_{ag}$ and that $a_{j+l} = a_{j+l+1} = \dots = a_{j+k} = a_{nop}$. By hypothesis, $E \cap \{s_j, \dots, s_{j+k}\} \neq \emptyset$ holds. Thus, we conclude that $E \cap \{s_{j+l+1}, \dots, s_{j+k}\} \neq \emptyset$ must hold, and hence $l < k$. However, by definition of $\Phi(s, a_{nop})$ we have $s_{j+l} = s_{j+l+1} = \dots = s_{j+k}$. This implies that $E \cap \{s_j, \dots, s_{j+l}\} \neq \emptyset$, which is a contradiction.

This proves that K witnesses k -maintainability of S w.r.t. E . \square

While this result shows that we could equally well have developed our notion of k -maintainability on the basis of trajectories, in the rest of this paper we shall stick to the setting which uses closure and unfolding. We find the latter more intuitive, as well as more convenient for designing algorithms and for proofs. Furthermore, this setting requires no special handling of possible finite trajectories, which complicates matters as becomes apparent from Proposition 3.

This alternative characterization suggests how to express the notion of k -maintainability using existing temporal logics when a distinction can be made⁵ between states that are reached through an agent's actions and states that are reached by an exogenous action. To make this distinction, let us assume that the latter kind of states have the fluent *interfered* as true and the former have it as false. Now, let $\text{Step}[k](\phi)$ be a shorthand for the formula

$$\phi \vee \bigcirc \phi \vee \bigcirc \bigcirc \phi \vee \dots \vee \underbrace{\bigcirc \dots \bigcirc}_k \phi$$

where the last subformula involves k consecutive \bigcirc 's. Intuitively, the formula $\text{Step}[k](\phi)$ means that ϕ is true within k steps. Now the last part of Proposition 3, ignoring the issue of a_{nop} actions (for simplicity), is as follows:

for each state s in S and every trajectory of form $\tau = s_0, a_1, s_1, a_2, \dots, a_j, s_j, a_{j+1}, \dots$ consistent with $A_{K,exo}^\infty$ and $s_0 = s$, $\{a_{i+1}, \dots, a_{i+k}\} \subseteq \mathcal{A}_{ag}$ for some $i \geq 0$ implies that $\{s_i, \dots, s_{i+k}\} \cap E \neq \emptyset$.

This can be expressed in LTL as $\phi_S \Rightarrow \square(\neg \text{Step}[k](\text{interfered}) \Rightarrow \text{Step}[k](\phi_E))$, or equivalently $\phi_S \Rightarrow \square(\text{Step}[k](\text{interfered} \vee \phi_E))$, where ϕ_S and ϕ_E are propositional formulas which described the states in S and E , respectively.

Using the above formula, k -maintainability can be written as follows:

There exists a control K for A w.r.t. \mathcal{A}_{ag} such that for each state s in S and every trajectory of form $\tau = s_0, a_1, s_1, a_2, \dots, a_j, s_j, a_{j+1}, \dots$ consistent with $A_{K,exo}^\infty$ and $s_0 = s$, the trajectory satisfies the temporal formula $\phi_S \Rightarrow (\square \text{Step}[k](\text{interfered} \vee \phi_E))$.

To capture the complete definition by a temporal formula, one needs branching time temporal operators akin to A , like the operator A_π in [9] meaning that “for all paths following the policy under consideration”. In that case, the specification would be $A_\pi(\square(\phi_S \Rightarrow \square(\text{Step}[k](\text{interfered} \vee \phi_E))))$. Note that here we need \square in between A_π and ϕ_S to indirectly account for the phrase “for each state s in S ”.

In upcoming sections, we discuss the above characterization in the context of general control generation algorithms that work with arbitrary temporal specifications. However, we use the characterization of k -maintainability in Definition 11 as it matches more closely with our algorithms, and does not necessitate defining a compilation that eliminates exogenous actions and introduces the new fluent *interfered* and proving the equivalence of that compilation.

5. Polynomial time methods to construct k -maintainable controls

Now that we have defined the notion of k -maintainability, our next step is to show how some k -maintainable control can be constructed in an automated way. We start with some historical background. In the program specification and synthesis literature, there have been a number of works, e.g., [19,49,59,60] to automatically synthesize programs similar to our control policies from given temporal logic specifications. Although these algorithms can accept richer goals, they all either allude to worst case exponential nature of their algorithms or prove that the complexity is exponential or even higher. None of them study special classes of specifications with lower complexity of constructing control. We discuss these papers and their complexity results in Section 8.2.

⁵ Such distinctions can be made by compiling a system, a set of agent actions, a specification of exogenous action occurrence, a set of goal states and a set of initial states to an automata that eliminates transitions due to exogenous actions but records their presence through the fluent *interfered*. We discuss such a compilation in Section 8.2.1.

There has been extensive use of situation control rules [26] and reactive control in the AI literature. But there have been less efforts in the AI literature to define correctness of such control rules [8],⁶ and to automatically construct correct control rules for general goals [10,37,41].⁷ In [42], it is suggested that in a control rule of the form: “if condition c is satisfied then do action a ”, the action a is the action that *leads to* the goal from any state where the condition c is satisfied. In [8] a formal meaning of “leads to” is given as: for all states s that satisfy c , a is the first action of a minimal cost plan from s to the goal. Using this definition, an algorithm is presented in [52] to construct k -maintainable controls. This algorithm is sound but not complete, in the sense that it generates correct controls only, but there is no guarantee that it will find always a control if one exists.

In [10,41], an algorithm to construct control with respect to linear temporal logic goals is given. This algorithm is based on progressing linear temporal formulas. The worst case complexity of the algorithm is given as double exponential with respect to the number of subformulas of the goal specification f , assuming a fixed number of states. Note that the temporal representation of k -maintainability will have subformulas with k nested operators \bigcirc . However, no studies regarding the complexity of specific goal specifications were done earlier. With the help of one of the authors of [10,41] we explored how well the algorithms in [10,41] will do with respect to our goal and found that by the using the algorithm as given in [10,41] the complexity will be exponential in the size of k . In Section 8.2.1 we will discuss this in more detail.

In the following, we overcome the problems one faces in the above mentioned approaches and give a sound and complete polynomial time algorithm for constructing k -maintainable control policies. In fact, the algorithm works in linear time for k bounded by a constant, and can be adapted to a linear time algorithm for maintainability, i.e., where k is arbitrary but finite.

We provide it in two steps: First we consider the case when the transition function Φ is deterministic, and then we generalize to the case where Φ may be non-deterministic. In each case, we present different methods, which illustrates our discovery process and also gives a better grasp of the final algorithm. We first present an encoding of our problem as a propositional theory and appeal to propositional SAT solvers to construct the control. As it turns out, this encoding is in a tractable fragment of SAT, for which specialized solvers (in particular, Horn SAT solvers) can be used easily. Finally, we present a direct algorithm distilled from the previous methods.

The reasoning behind this line of presentation is the following:

- (i) It illustrates the methodology of using SAT and Horn SAT encodings to solve problems;
- (ii) the encodings allow us to quickly implement and test algorithms;
- (iii) the proof of correctness mimics the encodings; and
- (iv) we can exploit known complexity results for Horn SAT to determine the complexity of our algorithm, and in particularly to establish tractability.

As for (ii), we can make use of answer set solvers such as DLV [29,45] or Smodels [53,65] which extend Horn logic programs by non-monotonic negation. These solvers allow efficient computation of the least model and some maximal models of a Horn theory, and can be exploited to construct robust or “small” controls, respectively.

Just for clarification, our approach of going from a SAT encoding to a Horn SAT encoding to a procedural algorithm is not to suggest that one is better or more efficient. Rather, it shows the usefulness of the general methodology to go from a logical specification (i.e., a SAT encoding in this case) to a procedural algorithm via transformations (Horn SAT encoding) that help us to prove the polynomial nature of the algorithm.

The problem we want to solve, which we refer to as k -MAINTAIN, has the following input and output:

Input: An input I is a system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$, sets of states $E \subseteq \mathcal{S}$ and $S \subseteq \mathcal{S}$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$, a function exo , and an integer $k \geq 0$.

Output: A control K such that S is k -maintainable with respect to E (using the control K), if such a control exists. Otherwise the output is the answer that no such control exists.

We assume here that the functions $\text{poss}(s)$ and $\text{exo}(s)$ can be efficiently evaluated; e.g., when both functions are given by their graphs (i.e., in a table).

5.1. Deterministic transition function $\Phi(s, a)$

We start with the case of deterministic transitions, i.e., $\Phi(s, a)$ is a singleton set $\{s'\}$ whenever nonempty. In abuse of notation, we simply will write $\Phi(s, a) = s'$ in this case.

⁶ Here we exclude the works related to MDPs as it is not known how to express the kind of goal we are interested in—such as k -maintenance goals—using reward functions.

⁷ In recent years, some planning algorithms and systems have been developed [13,14,18,39,40,44] that generate control rules for particular classes of goals.

Our first algorithm to solve k -MAINTAIN will be based on a reduction to propositional SAT solving. Given an input I for k -MAINTAIN, we construct a SAT instance $sat(I)$ in polynomial time such that $sat(I)$ is satisfiable if and only if the input I allows for a k -maintainable control, and that the satisfying assignments for $sat(I)$ encode possible such controls.

In our encoding, we shall use for each state $s \in S$ propositional variables s_0, s_1, \dots, s_k . Intuitively, s_i will denote that there is a path from state s to some state in E using only agent actions and at most i of them, to which we refer as “there is an a -path from s to E of length at most i ”.

The encoding $sat(I)$ contains the following formulas:

(0) For all $s \in S$, and for all j , $0 \leq j < k$:

$$s_j \Rightarrow s_{j+1}$$

(1) For all $s \in E \cap S$:

$$s_0$$

(2) For any two states $s, s' \in S$ such that $\Phi(a, s) = s'$ for some action $a \in \text{exo}(s)$:

$$s_k \Rightarrow s'_k$$

(3) For any state $s \in S \setminus E$ and all i , $1 \leq i \leq k$:

$$s_i \Rightarrow \bigvee_{s' \in PS(s)} s'_{i-1},$$

where $PS(s) = \{s' \in S \mid \exists a \in \mathcal{A}_{ag} \cap \text{poss}(s): s' = \Phi(a, s)\}$;

(4) For all $s \in S \setminus E$:

$$s_k$$

(5) For all $s \in S \setminus E$:

$$\neg s_0$$

The intuition behind the above encoding is as follows. The clauses in (0) state that if there is an a -path from s to E of length at most j then, logically, there is also an a -path of length at most $j + 1$. The clauses in (1) say that for states s in $S \cap E$, there is an a -path of length 0 from s to E . Next, (4) states that for any starting state s in S outside E , there is an a -path from s to E of length at most k , and (5) states that for any state s outside E , there is no a -path from s to E of length 0. The clauses in (3) state that if, for any state s , there is an a -path from s to E of length at most i , then for some possible agent action a and successor state $s' = \Phi(a, s)$, there is an a -path from s' to E of length at most $i - 1$. When looking for k -maintainable controls the clauses in (2) take into account the possibility that s may be in the closure. If indeed s is in the closure and there is an a -path from s to E of length at most k , then the same must be true with respect to the states s' reachable from s using exogenous actions. When looking for non-deterministic control they play a role in computing maximal non-deterministic controls. The role of each of the above clauses become more clear when relating the models of $sat(I)$ with controls that k -maintain.

Given any model M of $sat(I)$, we can extract a desired control K from it by defining $K(s) = a$ for all s outside E with s_k true in M , where a is a possible agent action in s such that $s' = \Phi(s, a)$ and s' is closer to E than s is. In case of multiple possible a and s' , one a can be arbitrarily picked. Otherwise, $K(s)$ is left undefined.

In particular, for $k = 0$, only the clauses from (1), (2), (4) and (5) do exist. As easily seen, $sat(I)$ is satisfiable in this case if and only if $S \subseteq E$ and no exogenous action leads outside E , i.e., the closure of S under exogenous actions is contained in E . This means that no actions of the agent are required at any point in time, and we thus obtain the trivial 0-control K which is undefined on all states, as desired.

The next result states that the SAT encoding works properly in general.

Proposition 4. Let I consist of a system $A = (S, \mathcal{A}, \Phi, \text{poss})$ where Φ is deterministic, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$, sets of states $E \subseteq S$ and $S \subseteq S$, an exogenous function exo , and an integer k . For any model M of $sat(I)$, let $C_M = \{s \in S \mid M \models s_k\}$, and for any state $s \in C_M$ let $\ell_M(s)$ denote the smallest index j such that $M \models s_j$ (i.e., s_0, s_1, \dots, s_{j-1} are false and s_j is true), which we call the level of s w.r.t. M . Then,

- (i) S is k -maintainable w.r.t. E iff $sat(I)$ is satisfiable.
- (ii) Given any model M of $sat(I)$, the partial function $K_M^+ : S \rightarrow 2^{\mathcal{A}_{ag}}$ defined on $C_M \setminus E$ such that

$$K_M^+(s) = \{a \in \mathcal{A}_{ag} \cap \text{poss}(s) \mid \Phi(s, a) = s', s' \in C_M, \ell_M(s') < \ell_M(s)\},$$

is a valid non-deterministic control for A w.r.t. \mathcal{A}_{ag} ;

- (iii) any control K which refines K_M^+ for some model M of $sat(I)$ k -maintains S w.r.t. E .

The proof of this proposition can be easily obtained by adapting the proof of Proposition 6.

5.1.1. Horn SAT encoding

While $\text{sat}(I)$ is constructible in polynomial time from I , we can not automatically infer that solving k -MAINTAIN is polynomial, since SAT is a canonical NP-hard problem. However, a closer look at the structure of the clauses in $\text{sat}(I)$ reveals that this instance is solvable in polynomial time. Indeed, it is a *reverse Horn* theory; i.e., by reversing the propositions, we obtain a Horn theory. Let us use propositions \bar{s}_i whose intuitive meaning is converse of the meaning of s_i . Then the Horn theory corresponding to $\text{sat}(I)$, denoted $\overline{\text{sat}}(I)$, is as follows:

(0) For all $s \in \mathcal{S}$ and $j, 0 \leq j < k$:

$$\overline{s_{j+1}} \Rightarrow \overline{s_j}.$$

(1) For all $s \in E \cap \mathcal{S}$:

$$\overline{s_0} \Rightarrow \perp.$$

(2) For any states $s, s' \in \mathcal{S}$ such that $s' = \Phi(a, s)$ for some action $a \in \text{exo}(s)$:

$$\overline{s'_k} \Rightarrow \overline{s_k}.$$

(3) For any state $s \in \mathcal{S} \setminus E$, and for all $i, 1 \leq i \leq k$:

$$\left(\bigwedge_{s' \in PS(s)} \overline{s'_{i-1}} \right) \Rightarrow \overline{s_i},$$

where $PS(s) = \{s' \in \mathcal{S} \mid \exists a \in \mathcal{A}_{ag} \cap \text{poss}(s): s' = \Phi(a, s)\}$.

(4) For all $s \in \mathcal{S} \setminus E$:

$$\overline{s_k} \Rightarrow \perp.$$

(5) For all $s \in \mathcal{S} \setminus E$:

$$\overline{s_0}.$$

Here, \perp denotes falsity. We then obtain a result similar to Proposition 4, and the models M of $\overline{\text{sat}}(I)$ lead to k -maintainable controls, which we can construct similarly; just replace in part (ii) C_M with $\bar{C}_M = \{s \in \mathcal{S} \mid M \not\models \overline{s_k}\}$. Notice that \bar{C}_M coincides with the set of states $C_{\bar{M}}$ for the model \bar{M} of $\text{sat}(I)$ such that $\bar{M} \models p$ iff $M \not\models \bar{p}$, for each atom p .

We now illustrate the above Horn encoding with respect to an example.

Example 5. Consider the system $A = (\mathcal{S}, \mathcal{A}, \Phi, \text{poss})$, where $\mathcal{S} = \{b, c, d, f, g, h\}$, $\mathcal{A} = \{\mathbf{a}, \mathbf{a}', \mathbf{e}\}$, and the (deterministic) transition function Φ was shown in Fig. 1, where $\Phi(s, a) = s'$ iff an arc $s \rightarrow s'$ labeled with a is present and $\text{poss}(s)$ are all actions that label arcs leaving s .

For $\mathcal{A} = \{\mathbf{a}, \mathbf{a}'\}$ and $\text{exo}(s) = \{\mathbf{e}\}$ iff $s = f$ and $\text{exo}(s) = \emptyset$ otherwise, this leads for $\mathcal{S} = \{b\}$, $E = \{h\}$, and $k = 3$ to the following Horn encoding $\overline{\text{sat}}(I)$:

(From 0)

$$\begin{array}{llllll} \overline{b_1} \Rightarrow \overline{b_0} & \overline{b_2} \Rightarrow \overline{b_1} & \overline{b_3} \Rightarrow \overline{b_2} & \overline{c_1} \Rightarrow \overline{c_0} & \overline{c_2} \Rightarrow \overline{c_1} & \overline{c_3} \Rightarrow \overline{c_2} \\ \overline{d_1} \Rightarrow \overline{d_0} & \overline{d_2} \Rightarrow \overline{d_1} & \overline{d_3} \Rightarrow \overline{d_2} & \overline{f_1} \Rightarrow \overline{f_0} & \overline{f_2} \Rightarrow \overline{f_1} & \overline{f_3} \Rightarrow \overline{f_2} \\ \overline{g_1} \Rightarrow \overline{g_0} & \overline{g_2} \Rightarrow \overline{g_1} & \overline{g_3} \Rightarrow \overline{g_2} & \overline{h_1} \Rightarrow \overline{h_0} & \overline{h_2} \Rightarrow \overline{h_1} & \overline{h_3} \Rightarrow \overline{h_2} \end{array}$$

(From 1)

(From 2)

$$\overline{g_3} \Rightarrow \overline{f_3}.$$

(From 3)

$$\begin{array}{lll} \overline{c_0} \wedge \overline{f_0} \Rightarrow \overline{b_1} & \overline{c_1} \wedge \overline{f_1} \Rightarrow \overline{b_2} & \overline{c_2} \wedge \overline{f_2} \Rightarrow \overline{b_3} \\ \overline{d_0} \Rightarrow \overline{c_1} & \overline{d_1} \Rightarrow \overline{c_2} & \overline{d_2} \Rightarrow \overline{c_3} \\ \overline{h_0} \Rightarrow \overline{d_1} & \overline{h_1} \Rightarrow \overline{d_2} & \overline{h_2} \Rightarrow \overline{d_3} \\ \overline{h_0} \Rightarrow \overline{f_1} & \overline{h_1} \Rightarrow \overline{f_2} & \overline{h_2} \Rightarrow \overline{f_3} \\ \overline{g_1} & \overline{g_2} & \overline{g_3} \end{array}$$

(From 4)

$$\overline{b_3} \Rightarrow \perp.$$

(From 5)

$$\overline{b_0} \quad \overline{c_0} \quad \overline{d_0} \quad \overline{f_0} \quad \overline{g_0}.$$




Fig. 8. Computing the least model: [6] From 5 we get $\bar{b}_0, \bar{c}_0, \bar{d}_0, \bar{f}_0, \bar{g}_0$; [7] From 3 we get $\bar{g}_1, \bar{g}_2, \bar{g}_3$; [8] From 6 and 3 we get \bar{b}_1, \bar{c}_1 ; [9] From 7 and 2 we get \bar{f}_3 ; [10] From 0 and 9 we get \bar{f}_2 ; [11] From 0 and 10 we get \bar{f}_1 ; [12] From 3, 8 and 11 we get \bar{b}_2 .

This theory has the least model

$$M = \{\bar{g}_3, \bar{g}_2, \bar{g}_1, \bar{g}_0, \bar{f}_3, \bar{f}_2, \bar{f}_1, \bar{f}_0, \bar{b}_2, \bar{b}_1, \bar{b}_0, \bar{c}_1, \bar{c}_0, \bar{d}_0\};$$

Hence, $\bar{C}_M = \{b, c, d, h\}$, which gives rise to the non-deterministic control K^+ such that $K^+(s) = \{a\}$ for $s \in \{b, c, d\}$ and $K^+(s)$ is undefined for $s \in \{f, g, h\}$. In this case, there is a single control K refining K^+ , which has $K(s) = a$ for $s \in \{b, c, d\}$ and is undefined otherwise. This is intuitive: The agent must reach h , and has to avoid taking a' in b since then it might arrive at the no-good state g . Thus, she has to take a in b and, as the only choice, in the subsequent states c and d . Also, we might not add any state apart from b, c , and d without losing 3-maintainability. In this particular case, M is also maximal on the propositions s_3 , where $s \in S \setminus E = \{b, c, d, f, g\}$: By (4), we can not add \bar{b}_3 , and by (0) and the clauses $\bar{c}_2 \wedge \bar{f}_2 \Rightarrow \bar{b}_3$ and $\bar{d}_1 \Rightarrow \bar{c}_2$ in (3) then also neither \bar{c}_3 nor \bar{d}_3 . Thus, the above control K is also smallest and, in fact, the only one possible for 3-maintainability.

As computing a model of a Horn theory is a well-known polynomial problem [25], we thus obtain the following result.

Theorem 5. Under deterministic state transitions, problem k -MAINTAIN is solvable in polynomial time. \square

An interesting aspect of the above is that, as well-known, each satisfiable Horn theory T has the least model, M_T , which is given by the intersection of all its models. Moreover, the least model is computable in linear time, cf. [25,50]. This model not only leads to a k -maintainable control, but also leads to a *maximal* control, in the sense that the control is defined on a greatest set of states outside E among all possible k -maintainable controls for S' w.r.t. E such that $S \subseteq S'$. This gives a clear picture of which other states may be added to S while k -maintainability is preserved; namely, any states in \bar{C}_{M_T} . Furthermore, any control K computed from M_T applying the method in Proposition 4 (using \bar{C}_{M_T}) works for such an extension of S as well.

On the other hand, intuitively a k -maintainable control constructed from some maximal model of $\bar{sat}(I)$ with respect to the propositions \bar{s}_k is undefined to a largest extent, and works merely for a smallest extension. We may generate, starting from M_T , such a maximal model of T by trying to flip first, step by step all propositions \bar{s}_k which are *false* to *true*, as well as other propositions entailed. In this way, we can generate a maximal model of T on $\{\bar{s}_k \mid s \in S \setminus E\}$ in polynomial time, from which a “lean” control can also be computed in polynomial time.

5.2. Non-deterministic transition function $\Phi(s, a)$

We now generalize our method for constructing k -maintainable controls to the case in which transitions due to Φ may be non-deterministic. As before, we first present a general propositional SAT encoding, and then rewrite to a propositional Horn SAT encoding. To explain some of the notations, we need the following definition, which generalizes the notion of an a -path to the non-deterministic setting.

Definition 12 (a -path). We say that there exists an a -path of length at most $k \geq 0$ from a state s to a set of states S' , if either $s \in S'$, or $s \notin S'$, $k > 0$ and there is some action $a \in \mathcal{A}_{ag} \cap \text{poss}(s)$ such that for every $s' \in \Phi(s, a)$ there exists an a -path of length at most $k - 1$ from s' to S' .

In the following encoding of an instance I of problem k -MAINTAIN to SAT, referred to as $sat'(I)$, s_i will again intuitively denote that there is an a -path from s to E of length at most i . The proposition s_{-a_i} , $i > 0$, will denote that for such s there is an a -path from s to E of length at most i starting with action a ($\in \text{poss}(s)$). The encoding $sat'(I)$ has again groups (0)–(5) of clauses as follows:

(0), (1), (4) and (5) are the same as in $\text{sat}(I)$.

(2) For any state $s \in S$ and s' such that $s' \in \Phi(a, s)$ for some action $a \in \text{exo}(s)$:

$$s_k \Rightarrow s'_k$$

(3) For every state $s \in S \setminus E$ and for all $i, 1 \leq i \leq k$:

$$(3.1) \quad s_i \Rightarrow \bigvee_{a \in \mathcal{A}_{ag} \cap \text{poss}(s)} s_{-a_i};$$

(3.2) for every $a \in \mathcal{A}_{ag} \cap \text{poss}(s)$ and $s' \in \Phi(s, a)$:

$$s_{-a_i} \Rightarrow s'_{i-1};$$

(3.3) for every $a \in \mathcal{A}_{ag} \cap \text{poss}(s)$, if $i < k$:

$$s_{-a_i} \Rightarrow s_{-a_{i+1}}.$$

Group (2) above is very similar to group (2) of $\text{sat}(I)$ in the previous subsection. The only change is that we now have $s' \in \Phi(a, s)$ instead of $s' = \Phi(a, s)$. The main difference is in group (3). We now explain those clauses. The clauses in (3.1) and (3.2) together state that if there is an a-path from s to E of length at most i , then there is some possible action a for the agent, such that for each state s' that potentially results by taking a in s , there must be an a-path from s' to E of length at most $i - 1$. The clauses $s_{-a_i} \Rightarrow s_{-a_{i+1}}$ in (3.3) say that on a longer a-path from s the agent must be able to pick a also. Notice that there are no formulas in $\text{sat}'(I)$ which forbid to pick different actions a and a' in the same state s , and thus we have a non-deterministic control; however, we can always refine it easily to a control.

Proposition 6. Let I consist of a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$, sets of states $E, S \subseteq S$, an exogenous function exo , and an integer k . For any model M of $\text{sat}'(I)$, let $C_M = \{s \in S \mid M \models s_k\}$, and for any state $s \in C_M \setminus E$ let $\ell_M(s)$ denote the smallest index j such that $M \models s_{-a_j}$ for some action $a \in \mathcal{A}_{ag} \cap \text{poss}(s)$, which we call the a-level of s w.r.t. M . Then,

- (i) S is k -maintainable w.r.t. E iff $\text{sat}'(I)$ is satisfiable;
- (ii) given any model M of $\text{sat}'(I)$, the partial function $K_M^+ : S \rightarrow 2^{\mathcal{A}_{ag}}$ which is defined on $C_M \setminus E$ by

$$K_M^+(s) = \{a \mid M \models s_{-a_{\ell_M(s)}}\}$$

is a valid non-deterministic control; and

- (iii) any control K which refines K_M^+ for some model M of $\text{sat}'(I)$ k -maintains S w.r.t. E .

Proof. Since the if direction of (i) follows from (ii) and (iii), it is sufficient to show the only if direction of (i) and both (ii) and (iii).

As for the only if direction of (i), suppose S is k -maintainable w.r.t. E . Then there exists a control K such that for each state $s \in \text{Closure}(S, A_{K, \text{exo}})$, and for each sequence $\sigma = s^{(0)}, s^{(1)}, \dots, s^{(l)}$ in $\text{Unfold}_k(s, A, K)$ where $s^{(0)} = s$, $\{s^{(0)}, \dots, s^{(l)}\} \cap E \neq \emptyset$. We now construct an interpretation M for $\text{sat}'(I)$ as follows.

For each $s \in \text{Closure}(S, A_{K, \text{exo}})$, and each sequence $\sigma = s^{(0)}, s^{(1)}, \dots, s^{(l)}$ in $\text{Unfold}_k(s, A, K)$ with $s = s^{(0)}$, let $i_\sigma (\geq 0)$ be the smallest index i such that $s^{(i)} \in E$, and let i^* be the maximum over all i_σ for s . Intuitively, i^* is the length of the longest path in the tree with root s where each node n not in E is sprouted by taking the control action $K(n)$ and adding each state in $\Phi(n, K(n))$ as a child. Then, we assign *true* to $s_{i^*}, s_{i^*+1}, \dots, s_k$ and, if $i^* > 0$, to $s_{-a_{i^*}}, s_{-a_{i^*+1}}, \dots, s_{-a_k}$, where $K(s) = a$. All other propositions are assigned *false* in M . We now argue that M is a model of $\text{sat}(I)$.

It is straightforward to see that M satisfies the formulas generated by (0), (1), (4) and (5). Now consider the formulas $s_k \Rightarrow s'_k$ generated in (2). If s_k is true, then $s \in \text{Closure}(S, A_{K, \text{exo}})$ by construction. In this case, for any $s' \in \Phi(a, s)$ of an exogenous action a , we have $s' \in \text{Closure}(S, A_{K, \text{exo}})$, and since K k -maintains S w.r.t. E , s'_i is *true* in M for some $i \leq k$ which implies, by construction, that s'_k is assigned *true* in M . Let us finally consider the formulas generated in (3). If s_i , where $s \in S \setminus E$, is assigned *true* in M for some $i \in \{1 \leq i \leq k\}$, then $s \in \text{Closure}(S, A_{K, \text{exo}})$ holds by construction of M . Since K is a k -maintaining control and $s \notin E$, we must have $K(s)$ defined and thus, by construction of M , we have s_{-a_i} assigned *true* in M . Since $K(s) \in \mathcal{A}_{ag} \cap \text{poss}(s)$, the clause (3.1) is thus satisfied. Furthermore, each clause in (3.2) is satisfied when $a \neq K(s)$, since then s_{a_i} is assigned *false* in M . For $a = K(s)$, proposition s_{a_i} is *true* in M and thus, by construction, also s_i . Since K is k -maintaining control, every state $s' \in \Phi(s, a)$ belongs to $\text{Closure}(S, A_{K, \text{exo}})$. Let, for each sequence $\sigma' = s^{(0)}, s^{(1)}, \dots, s^{(l)}$ in $\text{Unfold}_k(s, A, K)$ such that $s^{(0)} = s'$, the sequence $P(\sigma) = s^{(0)}, s^{(1)}, \dots, s^{(i)}$ be the shortest prefix of σ such that $s^{(i)} \in E$ (notice that $i < k$). Then, the sequence $s, P(\sigma)$ is a prefix of some sequence in $\text{Unfold}(s, A, K)$. Hence, it follows that in the construction of M , the number i^* for s is larger than the one for s' . Thus, by construction of M , it follows that s'_{i-1} is assigned *true* in M . This means that the formulas in (3.2) are satisfied in M . Finally, the clauses (3.3) are clearly satisfied in M by construction of M . Thus, M is a model of $\text{sat}'(I)$, which means that $\text{sat}'(I)$ is satisfiable.

To show (ii), let us assume that $\text{sat}'(I)$ has a model M , and consider the partial function $K_M^+ : S \rightarrow 2^{\mathcal{A}_{ag}}$ which is defined on $C_M \setminus E$ by $K_M^+(s) = \{a \mid M \models s_{-a_{\ell_M(s)}}\}$. We thus have to show that $K_M^+(s) \subseteq \text{poss}(s)$ and $K_M^+(s) \neq \emptyset$ when $K_M^+(s)$ is defined. By clause (3.1), and the definition of C_M , ℓ_M , and K_M^+ this is immediate.

To show (iii), let K be any control which refines K_M^+ for some model M of $\text{sat}'(I)$. Let the distance $d_K(s, S)$ of a state s from the set of states S be as in the proof of Proposition 4, i.e., the minimum number of transitions—through exogenous actions and/or control actions dictated by the control K —needed to reach s from any state in S .

We will show, by using induction on $d(s, S) \geq 0$, that for every state $s \in \text{Closure}(S, A_{K, \text{exo}})$ and every sequence $\sigma = s^{(0)}, s^{(1)}, \dots, s^{(l)}$ with $s = s^{(0)}$ in $\text{Unfold}_k(s, A, K)$, the set $\{s^{(0)}, \dots, s^{(l)}\}$ intersects with E and that $M \models s_k$ (i.e., $s \in C_M$). This proves that K k -maintains S w.r.t. E .

The base case, $d(s, S) = 0$, is about states $s \in S$. From the formulas in (0), (1), and (4) we have $M \models s_k$ for every such state s . Consider any sequence $\sigma = s^{(0)}, s^{(1)}, \dots, s^{(l)}$ in $\text{Unfold}_k(s, A, K)$ such that $s = s^{(0)}$. If $s \in E$, then we must have $l = 0$, and $\{s^{(0)}, \dots, s^{(l)}\} \cap E \neq \emptyset$. Otherwise, $M \models s_{a_k}$ where $a = K(s)$. We then have $s^{(1)} \in \Phi(s, a)$, and thus by our construction of K and the clauses in (3.2) we have that $M \models s_{k-1}^{(1)}$. Repeating this argument, we can infer that $s_k^{(0)}, s_{k-1}^{(1)}, \dots, s_{k-l}^{(l)}$ are all assigned *true* in M . If $k = l$, it follows from the clauses in (5) that $s^{(l)} \in E$. Otherwise, if $l < k$, then K must be undefined on $s^{(l)}$; by the clauses (1), this again means $s^{(l)} \in E$. Hence, $\{s^{(0)}, \dots, s^{(l)}\} \cap E \neq \emptyset$.

Thus the statement holds in the base case. Now for the induction step, let us assume that it holds for every state $s \in \text{Closure}(S, A_{K, \text{exo}})$ at distance $d(s, S) = d \geq 0$ from S . Let us now consider a state $s \in \text{Closure}(S, A_{K, \text{exo}})$ at distance $d(s', S) = d + 1$ from S . Then there is a state s' at distance $d(s, S) = d$ from S such that $s \in \Phi(a, s')$ and either (i) $a \in \text{exo}(s')$ or (ii) $a \in K(s')$. In both cases, we have by the induction hypothesis that $M \models s'_k$, and we can conclude $M \models s_k$ from the clauses in (2) in case (i) and from our construction of K and the clauses in (3.2), (1), and (0) in case (ii), respectively. Furthermore, by similar argumentation as in the case $d = 0$ above, we obtain that for each sequence $\sigma = s^{(0)}, s^{(1)}, \dots, s^{(l)}$ in $\text{Unfold}_k(s, A, K)$ with $s = s^{(0)}$ it holds that $\{s^{(0)}, \dots, s^{(l)}\} \cap E \neq \emptyset$. This concludes the induction and the proof of (iii). \square

One advantage of the encoding $\text{sat}'(I)$ over the encoding $\text{sat}(I)$ for deterministic transition function Φ above is that it directly gives us the possibility to read off a suitable control from the s_{a_i} propositions, $a \in \text{poss}(s)$, which are true in any model M that we have computed, without looking at the transition function $\Phi(s, a)$ again. On the other hand, the encoding is more involved, and uses a larger set of propositions. Nonetheless, the structure of the formulas in $\text{sat}'(I)$ is benign for computation and allows us to compute a model, and from it a k -maintainable control in polynomial time.

5.2.1. Horn SAT encoding (general case)

The encoding $\text{sat}'(I)$ is, like $\text{sat}(I)$, a reverse Horn theory. We thus can rewrite $\text{sat}'(I)$ similarly to a Horn theory, $\overline{\text{sat}}'(I)$ by reversing the propositions, where the intuitive meaning of $\overline{s_i}$ and $\overline{s_{a_i}}$ is the converse of the meaning of s_i and s_{a_i} respectively. The encoding $\overline{\text{sat}}'(I)$ is as follows:

- (0), (1), (4) and (5) are as in $\overline{\text{sat}}(I)$
 (2) For every states $s, s' \in S$ such that $s' \in \Phi(a, s)$ for some action $a \in \text{exo}(s)$:

$$\overline{s'_k} \Rightarrow \overline{s_k}.$$

- (3) For every state $s \in S \setminus E$ and for all $i, 1 \leq i \leq k$:

$$(3.1) (\bigwedge_{a \in \mathcal{A}_{ag} \cap \text{poss}(s)} \overline{s_{a_i}}) \Rightarrow \overline{s_i};$$

$$(3.2) \text{ for every } a \in \mathcal{A}_{ag} \cap \text{poss}(s) \text{ and } s' \in \Phi(s, a):$$

$$\overline{s'_{i-1}} \Rightarrow \overline{s_{a_i}};$$

$$(3.3) \text{ for every } a \in \mathcal{A}_{ag} \cap \text{poss}(s), \text{ if } i < k:$$

$$\overline{s_{a_{i+1}}} \Rightarrow \overline{s_{a_i}}.$$

We obtain from Proposition 6 easily the following result, which is the main result of this section so far.

Theorem 7. Let I consist of a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$, sets of states $E, S \subseteq S$, an exogenous function exo , and an integer k . Let, for any model M of $\overline{\text{sat}}'(I)$, $\overline{C}_M = \{s \mid M \not\models \overline{s_k}\}$, and let $\overline{\ell}_M(s) = \min\{j \mid M \not\models \overline{s_{a_j}}, a \in \mathcal{A}_{ag} \cap \text{poss}(s)\}$ for every $s \in S$. Then,

- (i) S is k -maintainable w.r.t. E iff the Horn SAT instance $\overline{\text{sat}}'(I)$ is satisfiable;
 (ii) Given any model M of $\overline{\text{sat}}'(I)$, every control K such that $K(s)$ is defined iff $s \in \overline{C}_M \setminus E$ and

$$K(s) \in \{a \in \mathcal{A}_{ag} \cap \text{poss}(s) \mid M \not\models \overline{s_{a_j}}, j = \overline{\ell}_M(s)\},$$

k -maintains S w.r.t. E .

Corollary 8. Problem k -MAINTAIN is solvable in polynomial time. More precisely, it is solvable in time $O(k\|I\|)$, where $\|I\|$ denotes the size of input I .

Proof. A straightforward analysis yields that the size of $\overline{sat}'(I)$, measured by the number of atoms in it, is $O(k(|S| + |\Phi| + |poss|))$, if \mathcal{A}_{ag} , S , E , Φ , $poss$ and exo are stored in a standard way as bitmaps, i.e., a (multi-dimensional) array with value range $\{0,1\}$ (thus, $\|I\| = O(|S|^2|\mathcal{A}| + \log k)$). Furthermore, the clauses in $\overline{sat}'(I)$ can be easily generated within the same time bound. Since the least model of any Horn theory T is computable in time $O(|T|)$ where $|T|$ is the number of atoms in it [25,50], deciding satisfiability and computing some model M of $\overline{sat}'(I)$ is feasible in $O(k\|I\|)$ time. Furthermore, \overline{C}_M and $\{(s, \tilde{e}_M(s)) \mid s \in S\}$ are computable from M in linear time in the number of atoms, using suitable data structures, and from this a control K as in Theorem 7(ii) in the same time. Hence, a k -maintaining control for S w.r.t. E is computable in $O(k\|I\|)$ time.

Note that a more economic representation stores S , E , \mathcal{A}_{ag} as sets (i.e., lists) and Φ , $poss$, and exo by their graphs in tables, i.e., sets of tuples $\{(s, a, \Phi(s, a)) \mid s \in S, a \in \mathcal{A}\}$, $\{(s, poss(s)) \mid s \in S\}$, and $\{(s, exo(s)) \mid s \in S\}$. Also under this representation, and if moreover tuples where $\Phi(s, a) = \emptyset$ (resp., $poss(s) = \emptyset$ and $exo(s) = \emptyset$) are not stored (which is of the same order as storing the sets of tuples $\{(s, a, s') \mid s' \in \Phi(a, s)\}$, $\{(s, a) \mid a \in poss(s)\}$, $\{(s, a) \mid a \in exo(s)\}$), the $O(k\|I\|)$ time bound holds. Indeed, arrays storing S , E , and \mathcal{A}_{ag} for lookup in $O(1)$ time are constructible in time $O(|S| + |\mathcal{A}|)$. Then, $poss_{ag} = \{(s, a) \in poss \mid a \in \mathcal{A}_{ag}\}$ storing $\mathcal{A}_{ag} \cap poss(s)$ for all s is constructible in $O(|poss|)$ time. From this, all clauses of $\overline{sat}'(I)$ except (2) and (3.2) can be readily generated in time $O(k(|S| + |poss_{ag}|))$. The clauses (2) and (3.2) can be easily constructed from $\Phi_{exo} = \{(s, a, s') \in \Phi \mid a \in exo(s)\}$ and $\Phi_{poss} = \{(s, a, s') \in \Phi \mid a \in poss(s)\}$ in time $O(|\Phi_{exo}|)$ and $O(k|\Phi_{poss}|)$, respectively. The sets Φ_{exo} and Φ_{poss} can be generated from Φ and exo in time $O(|\Phi| + |exo| + |poss|)$, using an auxiliary array $aux[\mathcal{A}, S]$ to enable random access to $\Phi(a, s)$; notice that $aux[a, s]$ needs not be defined if $\Phi(a, s) = \emptyset$. In total, $\overline{sat}'(I)$ is constructible in $O(|\mathcal{A}| + |exo| + k(|S| + |\Phi| + |poss|)) = O(k\|I\|)$ time. \square

Thus in particular, finding a maintaining control under a small window of opportunity for maintenance, i.e., a k -maintaining control for k bounded by a constant, is feasible in *linear time* in the size of the input.

Similar as in Section 5.1.1, the least model of the theory given by $\overline{sat}'(I)$, $M_{\overline{sat}'(I)}$, leads to a *maximal* control in the sense that the pre-image of K outside E , i.e., the states outside E in which K is defined, is greatest among all possible k -maintaining controls which include S . Furthermore, a smallest k -maintaining control can be similarly computed from any maximal model of $\overline{sat}'(I)$ with respect to the propositions \overline{s}_k where s is outside E , which can be generated from $M_{\overline{sat}'(I)}$ by stepwise maximization. Again, both maximal and smallest controls can be computed in polynomial time.

Example 6. Reconsider the system $A = (S, \mathcal{A}, \Phi, poss)$ from Example 5. Let us modify the transition function Φ such that $\Phi(c, \mathbf{a}) = \{d, f\}$ instead of $\Phi(c, \mathbf{a}) = \{d\}$. Then, for the respective modified instance I of 3-MAINTAIN, denoted I_1 , the encoding $\overline{sat}'(I_1)$ looks as follows.

(0), (1), (2), (4), and (5) are as in $sat(I_1)$ in Example 5;

$$(3.1): \begin{array}{lll} \overline{b_a1} \wedge \overline{b_a'1} \Rightarrow \overline{b_1} & \overline{b_a2} \wedge \overline{b_a'2} \Rightarrow \overline{b_2} & \overline{b_a3} \wedge \overline{b_a'3} \Rightarrow \overline{b_3} \\ \overline{c_a1} \Rightarrow \overline{c_1} & \overline{c_a2} \Rightarrow \overline{c_2} & \overline{c_a3} \Rightarrow \overline{c_3} \\ \overline{d_a1} \Rightarrow \overline{d_1} & \overline{d_a2} \Rightarrow \overline{d_2} & \overline{d_a3} \Rightarrow \overline{d_3} \\ \overline{f_a1} \Rightarrow \overline{f_1} & \overline{f_a2} \Rightarrow \overline{f_2} & \overline{f_a3} \Rightarrow \overline{f_3} \\ \overline{g_1} & \overline{g_2} & \overline{g_3} \end{array}$$

$$(3.2): \begin{array}{llllll} \overline{h_0} \Rightarrow \overline{d_a1} & \overline{h_1} \Rightarrow \overline{d_a2} & \overline{h_2} \Rightarrow \overline{d_a3} & \overline{h_0} \Rightarrow \overline{f_a1} & \overline{h_1} \Rightarrow \overline{f_a2} & \overline{h_2} \Rightarrow \overline{f_a3} \\ \overline{d_0} \Rightarrow \overline{c_a1} & \overline{d_1} \Rightarrow \overline{c_a2} & \overline{d_2} \Rightarrow \overline{c_a3} & \overline{f_0} \Rightarrow \overline{c_a1} & \overline{f_1} \Rightarrow \overline{c_a2} & \overline{f_2} \Rightarrow \overline{c_a3} \\ \overline{c_0} \Rightarrow \overline{b_a1} & \overline{c_1} \Rightarrow \overline{b_a2} & \overline{c_2} \Rightarrow \overline{b_a3} & \overline{f_0} \Rightarrow \overline{b_a'1} & \overline{f_1} \Rightarrow \overline{b_a'2} & \overline{f_2} \Rightarrow \overline{b_a'3} \end{array}$$

$$(3.3): \begin{array}{lllll} \overline{d_a2} \Rightarrow \overline{d_a1} & \overline{d_a3} \Rightarrow \overline{d_a2} & \overline{f_a2} \Rightarrow \overline{f_a1} & \overline{f_a3} \Rightarrow \overline{f_a2} & \overline{c_a2} \Rightarrow \overline{c_a1} \\ \overline{c_a3} \Rightarrow \overline{c_a2} & \overline{b_a2} \Rightarrow \overline{b_a1} & \overline{b_a3} \Rightarrow \overline{b_a2} & \overline{b'_a2} \Rightarrow \overline{b'_a1} & \overline{b'_a3} \Rightarrow \overline{b'_a2} \end{array}$$

It turns out that $\overline{sat}'(I)$ has no models: From $\overline{g_3}$, the clause $\overline{g_3} \Rightarrow \overline{f_3}$ in (2), and clauses in (0), we obtain that $\overline{f_i}$, $i \in \{0, \dots, 3\}$, is true in every model M of $\overline{sat}'(I_1)$. Hence, by the clause $\overline{f_2} \Rightarrow \overline{b_a3}$ in (3.2), also $\overline{b_a3}$ is true in M . On the other hand, from the formula $\overline{f_1} \Rightarrow \overline{c_a2}$ in (3.2), we obtain that $\overline{c_a2}$ must be true in M , and thus by the clauses $\overline{c_a2} \Rightarrow \overline{c_2}$ in (3.1) and $\overline{c_2} \Rightarrow \overline{b_a3}$ in (3.2) that $\overline{b_a3}$ is true in M . The clause $\overline{b_a3} \wedge \overline{b_a'3} \Rightarrow \overline{b_3}$ thus implies that $\overline{b_3}$ is true in M . However, by the formula $\overline{b_3} \Rightarrow \perp$ in (4), $\overline{b_3}$ must be false in M . Thus, no model M of $\overline{sat}'(I_1)$ can exist, which by Theorem 7 means that there is no 3-maintaining control for $S = \{b\}$ w.r.t. $E = \{h\}$. Indeed, regardless of whether a control function K selects \mathbf{a} or \mathbf{a}' in state b , within at most 2 steps from b the state f might be reached, from which the exogenous function might move the system to the no-good state g .

Suppose now again that $\Phi(c, \mathbf{a}) = \{d, f\}$ and that the agent can take \mathbf{a}' in g , which results in either h or f (i.e., $\Phi(g, \mathbf{a}') = \{f, h\}$ and $\mathbf{a}' \in poss(g)$). Then the Horn encoding $\overline{sat}'(I_1)$ changes as follows:

In (3.1), the facts \bar{g}_i , $i \in \{1, 2, 3\}$, are replaced by $\bar{g_a_i} \Rightarrow \bar{g}_i$; In (3.2.), the clauses for $\mathbf{a'}$ and f, h are added, $i \in \{1, 2, 3\}$:

$$\bar{f_0} \Rightarrow \bar{g_a_1}. \quad \bar{f_1} \Rightarrow \bar{g_a_2}. \quad \bar{f_2} \Rightarrow \bar{g_a_3}. \quad \bar{h_0} \Rightarrow \bar{g_a_1}. \quad \bar{h_1} \Rightarrow \bar{g_a_2}. \quad \bar{h_2} \Rightarrow \bar{g_a_3}.$$

In (3.3), the clauses for $\mathbf{a'}$ and g are added:

$$\bar{g_a_2} \Rightarrow \bar{g_a_1}. \quad \bar{g_a_3} \Rightarrow \bar{g_a_2}.$$

In this encoding $\bar{sat'}(I_2)$ of the modified instance I_2 , we no longer have a fact $\bar{g_3}$ in (3.1) and thus the above derivation of a contradiction for the truth value of $\bar{b_3}$ in any model of $\bar{sat'}(I_2)$ is not applicable. In fact, $\bar{sat'}(I_2)$ is satisfiable, and its least model is

$$M = \{\bar{b_0}, \bar{c_0}, \bar{d_0}, \bar{f_0}, \bar{g_0}, \bar{b_a_1}, \bar{c_a_1}, \bar{b_a_1'}, \bar{g_a_1'}, \bar{b_1}, \bar{c_1}, \bar{g_1}, \bar{b_a_2}\}.$$

Then, we have $\bar{C}_M = \{b, c, d, f, g, h\}$, $\bar{\ell}_M(b) = \bar{\ell}_M(c) = \bar{\ell}_M(g) = 2$ and $\bar{\ell}_M(d) = \bar{\ell}_M(f) = 1$, which leads to a single 3-maintaining control K such that $K(s) = \mathbf{a}$ for $s \in \{b, c, d, f\}$ and $K(g) = \mathbf{a'}$. Note that since K is defined on every state except h , it 3-maintains every set S w.r.t. every E which includes h . As for $S = \{b\}$, $K(c)$ and $K(d)$ could remain undefined, since they are not in the closure of b (which can be easily detected) at the price of losing robustness with respect to enlarging S . There is an alternative solution in which $K(b) = \mathbf{a'}$ instead of $K(b) = \mathbf{a}$. Here $K(s)$ can not be made undefined on any $s \neq h$.

5.3. Genuine algorithm

From the encoding to Horn SAT above, we can distill a direct algorithm to construct a k -maintainable control, if one exists. The algorithm mimics the steps which a SAT solver might take in order to solve $sat'(I)$. It uses counters $c[s]$ and $c[s_a]$ for each state $s \in S$ and possible agent action a in state s , which range over $\{-1, 0, \dots, k\}$ and $\{0, 1, \dots, k\}$, respectively. Intuitively, value i of counter $c[s]$ (at a particular step in the computation) represents that so far $\bar{s_0}, \dots, \bar{s_i}$ are assigned true, and that at least $i + 1$ steps are needed from s to reach E ; in particular, $i = -1$ represents that no $\bar{s_i}$ is assigned true yet. Similarly, value i for $c[s_a]$ (at a particular step in the computation) represents that so far $\bar{s_a_1}, \dots, \bar{s_a_i}$ are assigned true (in particular, $i = 0$ that no $\bar{s_a_i}$ is assigned true yet), and that at least $i + 1$ steps are needed from s to reach E starting with a .

Starting from an initialization, the algorithm updates by demand of the clauses in $\bar{sat'}(I)$ the counters (i.e., sets propositions true) using a command $upd(c, i)$ which is short for “if $c < i$ then $c := i$,” towards a fixpoint. If a counter violation is detected, corresponding to violation of a clause $\bar{s_0} \rightarrow \perp$ for $s \in S \cap E$ in (1) or $\bar{s_k} \rightarrow \perp$ for $s \in S \setminus E$ in (4), then no control is possible. Otherwise, a control is constructed from the counters.

The detailed algorithm is shown in Fig. 9.

It can be easily adjusted if we simply want to output a non-deterministic control such that each of its refinements is a k -maintainable control, leaving a choice about the refinement to the user. Alternatively, we can implement in Step 4 such a choice based on preference information. The following proposition states that the algorithm works correctly and runs in polynomial time.

Proposition 9. *Algorithm k -CONTROL solves problem k -MAINTAIN, and terminates for any input I in polynomial time. Furthermore, it can be implemented to run in $O(k\|I\|)$ time.*

Algorithm k -CONTROL

Input: A system $A = (S, \mathcal{A}, \Phi, poss)$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$ of agent actions, sets of states $E, S \subseteq S$, an exogenous function exo , and an integer $k \geq 0$.

Output: A control K which k -maintains S with respect to E , if any such control exists. Otherwise, output that no such control exists.

(Step 1) Initialization

- (i) Set $\Phi_{exo} = \{(s, a, s') \mid s \in S, a \in exo(s), s' \in \Phi(s, a)\}$, $\Phi_{poss}^E = \{(s, a, s') \mid s \in S \setminus E, a \in poss(s), s' \in \Phi(s, a)\}$, and for every $s \in S$, $poss_{ag}(s) = \mathcal{A}_{ag} \cap poss(s)$.
- (ii) For every s in E , set $c[s] := -1$.
- (iii) For every s in $S \setminus E$, set $c[s] := k$ if $poss_{ag}(s) = \emptyset$; otherwise, set $c[s] := 0$.
- (iv) For every s in $S \setminus E$ and $a \in poss_{ag}(s)$, set $c[s_a] := 0$.

(Step 2) Repeat the following steps until there is no change or $c[s] = k$ for some $s \in S \setminus E$ or $c[s] \geq 0$ for some $s \in S \cap E$:

- (i) For any $(s, a, s') \in \Phi_{exo}$ such that $c[s'] = k$ do $upd(c[s], k)$.
- (ii) For any $(s, a, s') \in \Phi_{poss}^E$ such that $c[s'] = i$:
if $0 \leq i < k$ then do $upd(c[s_a], i + 1)$, elseif $i = k$ then do $upd(c[s_a], k)$.
- (iii) For any state $s \in S \setminus E$ such that $poss_{ag}(s) \neq \emptyset$ and $i = \min\{c[s_a] \mid a \in poss_{ag}(s)\}$ do $upd(c[s], i)$.

(Step 3) If $c[s] = k$ for some $s \in S \setminus E$ or $c[s] \geq 0$ for some $s \in S \cap E$, then output that S is not k -maintainable w.r.t. E and halt.

(Step 4) Output any control $K : S \setminus E \rightarrow \mathcal{A}_{ag}$ defined on all states $s \in S \setminus E$ with $c[s] < k$ and such that $K(s) \in \{a \in poss_{ag}(s) \mid c[s_a] = \min_{b \in poss_{ag}(s)} c[s_b] < k\}$.

Fig. 9. Algorithm for problem k -MAINTAIN.

Proof. The correctness of the algorithm follows from Theorem 7 and the fact that k -CONTROL mimics, starting from facts in (5) and (3.1), the computation of the least model of $\overline{\text{sat}}'(I)$ by a standard fix-point computation. As for the polynomial time complexity, since counters are only increased, and the loop in Step 2 is reentered only if at least one counter has increased in the latest run, it follows that the number of iterations is polynomially bounded. Since the body of Step 2 and each other step is polynomial, it follows that k -CONTROL runs in polynomial time.

For the more detailed account, note that bitmaps for S , E and \mathcal{A} (if not available in the input) can be generated in time $O(|S| + |\mathcal{A}|)$. In (i) of Step 1, the sets Φ_{exo} and Φ_{poss}^E can be constructed in time $O(|\Phi| + |exo|)$ and $O(|\Phi| + |poss| + |S|)$, respectively, using an auxiliary array for random access to $\Phi(a, s)$ in case if the functions are given by their graphs (cf. proof of Corollary 8). Constructing $poss_{ag}(s)$ for all $s \in S$ takes $O(|poss|)$ time, and (ii)–(iv) of Step 1 is feasible in time $O(|S| + |poss|)$.

Using flags to signal changes to counters $c[s]$, $c[s_a]$, and auxiliary counters for $\min(c[s_a] \mid a \in poss_{ag}(s))$, the number of calls of *upd* in Step 2 is $O(k(|\Phi_{exo}| + |\Phi_{poss}| + |S|))$, and each call takes $O(1)$ time. The loop condition can be checked in $O(m)$ time where m is the number of changes in the loop. Hence, the total time for Step 2 is $O(k\|I\|)$. Step 3 is $O(1)$ if a flag is set in Step 2 indicating the reason for the loop exit. Finally, in Step 4, a control K can be easily output in time $O(|poss|)$. In total, the time is $O(k\|I\|)$. \square

Thus, for k bounded by a constant, k -CONTROL can be implemented to run in linear time. We remark that further improvements are possible. For example, states may be eliminated beforehand which will not be reachable from any state in S under any control that is eventually constructed. This can be done efficiently by computing an upper bound of $\text{Closure}(S, K_{A,exo})$ in which all possible actions at any state are merged into a single action. Similarly, we can efficiently prune all states which can not reach E within k steps in linear time. This can e.g. be achieved by a slight extension of algorithm K -CONTROL, in which flags *final*[s] and *final*[s_a] for the states s and actions $a \in poss_{ag}(s)$ signal whether the counters $c[s]$ and $c[s_a]$ correspond to the shortest distance to E , and in Step 2 (ii) and (iii) we chose next always some s_a (respectively s) such that its *final* flag can be switched from false to true and $c[s_a]$ (resp. $c[s]$) is smaller than k . We leave further discussion and refinements for future work.

5.4. Generic maintaining controls

By the results in the previous subsections, we can solve problem MAINTAIN, which is analogous to k -MAINTAIN but k is not in the input and can be arbitrarily chosen, in time $O(|S|\|I\|)$, that is, in time quadratic in the size of the input. This follows from the fact that k -maintainability of S w.r.t. E for some arbitrary but finite $k \geq 0$ is equivalent to k -maintainability of S w.r.t. E where $k = |S|$ is the number of states.

However, we can take advantage of the property that the exact number of steps to reach E does not matter (as long as it is finite), and design a more efficient (linear time) algorithm, which proceeds in two phases. In the first phase, those states s are determined from which E is reachable by an a -path of arbitrary length, and all other states are pruned. In the second phase, those states are iteratively pruned which are taken by some exogenous action to a state without such an a -path, or where each action a leads to a pruned state.

We can obtain a genuine linear-time algorithm for solving problem MAINTAIN by adapting the algorithm k -CONTROL such that it implements the two phases, where the counters $c[s]$ and $c[s_a]$ only range over a fixed domain independent of k . However, we skip the discovery process and go straight to a simple algorithm, which is shown in Fig. 10. We refer to this algorithm as ω -control or ω -maintaining control.

It implements the phases 1 and 2 in the steps 2 and 3, respectively. If in Step 4 a maintaining control is found to exist, Step 5 extracts such a control from the data structures. Like for k -maintainability, this requires some care since a naïve extraction does not work (in particular, cycles may cause problems). The following result, whose proof is omitted, states that the algorithm works properly.

Proposition 10. *Algorithm ω -CONTROL solves problem MAINTAIN, and terminates for any input I in polynomial time. Furthermore, it can be implemented to run in time $O(\|I\|)$, i.e., in linear time.*

6. Encoding k -maintainability for an answer set solver

In this section, we use the results of the previous section to show how computing a k -maintainable control can be encoded as finding answer sets of a non-monotonic logic program. More precisely, we describe an encoding to non-monotonic logic programs under the Answer Set Semantics [33], which can be executed on one of the available answer set solvers such as DLV [29,45] or Smodels [53,65]. These solvers support the computation of answer sets (models) of a given program, from which solutions (in our case, k -maintaining controls) can be extracted.

The encoding is generic, i.e., given by a *fixed program* which is evaluated over the instance I represented by input facts $F(I)$. It makes use of the fact that non-monotonic logic programs can have multiple models, which correspond to different solutions, i.e., different k -maintainable controls.

Algorithm ω -CONTROL**Input:** A system $A = (S, \mathcal{A}, \Phi, \text{poss})$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$ of agent actions, sets of states $E, S \subseteq \mathcal{S}$, an exogenous function exo .**Output:** A control K which maintains S with respect to E , if any such control exists. Otherwise, output that no such control exists.(Step 1) $X := E$.(Step 2) Repeat until there is no change to X :

$$X := X \cup \{s \mid \exists a \in \mathcal{A}_{ag}: \forall s', s' \in \Phi(s, a), s' \in X\}.$$

(Step 3) Repeat until there is no change to X :

$$X := X \setminus \{s \mid \exists a \in \text{exo}(s): \exists s', s' \in \Phi(s, a), s' \notin X\};$$

$$X := X \setminus \{s \mid \forall a \in \mathcal{A}_{ag}: \exists s', s' \in \Phi(s, a), s' \notin X\}.$$

(Step 4) If $S \setminus X \neq \emptyset$ then output that S is not maintainable w.r.t. E and halt.

(Step 5) Construct the control going backwards from the goal states in the following manner.

(i) Initialize counters: for all $s \in X$ and $a \in \mathcal{A}_{ag}$ do $c[s_a] := |\Phi(s, a)|$.(ii) For every state $s \in E$ do put (s, nop) in a queue Q .(iii) While Q is not empty doPop an element (s, x) from Q ;if $s \notin E$ then $K(s) := x$;for all transitions (s', a, s) such that $s \in \Phi(s', a)$ and $s' \in X$ do $c[s'_a] := c[s'_a] - 1$;if $c[s'_a] = 0$ and $K(s')$ is undefined then put (s', a) in Q .**Fig. 10.** Algorithm for problem MAINTAIN.

In the following, we first describe how a system is represented in a logic program, and then we develop the logic programs for both deterministic and general, non-deterministic domains. We shall follow here the syntax of the DLV system; the changes needed to adapt the programs to other answer set solvers such as Smodels are minor.

6.1. Input representation

The input I of problem k -MAINTAIN, can be represented by facts $F(I)$ as follows.

- The system $A = (S, \mathcal{A}, \Phi, \text{poss})$ can be represented using predicates `state`, `transition`, and `poss` by the following facts:
 - `state(s)`, for each $s \in S$;
 - `action(a)`, for each $a \in \mathcal{A}$;
 - `transition(s, a, s')`, for each $s, s' \in S$ and $a \in \mathcal{A}$ such that $s' \in \Phi(s, a)$;
 - `poss(s, a)`, for each $s \in S$ and $a \in \mathcal{A}$ such that $a \in \text{poss}(s)$.
- the set $\mathcal{A}_{ag} \subseteq \mathcal{A}$ of agent actions is represented using a predicate `agent` by facts `agent(a)`, for each $a \in \mathcal{A}_{ag}$;
- the set of states S is represented by using a predicate `start` by facts `start(s)`, for each $s \in S$;
- the set of states E is represented by using a predicate `goals` by facts `goal(s)`, for each $s \in E$;
- the exogenous function exo is represented by using a predicate `exo` by facts `exo(s, a)` for each $s \in S$ and $a \in \text{exo}(s)$;

The integer k is represented as constant `k`.

Example 7. Coming back to Example 3, the input I is represented as follows:

```
state(b). state(c). state(d). state(f). state(g). state(h).
action(a). action(a1). action(e).
trans(b,a,c). trans(b,a1,f). trans(c,a,d). trans(d,a,h).
trans(f,a,h). trans(f,e,g).
poss(b,a). poss(b,a1). poss(f,a). poss(f,e).
poss(c,a). poss(d,a).
agent(a). agent(a1).
start(b). goal(h).
exo(f,e).
const k=3.
```

6.2. Deterministic transition function Φ

The following is a program, executable on the DLV engine, for deciding the existence of a k -control. In addition to the predicates for the input facts $F(I)$, it employs a predicate $n_path(X, I)$, which intuitively corresponds to \bar{X}_I , and further auxiliary predicates.

```
% Define range of 0,1,...,(k) for stages.
range(0..k).

% Rule for (0).
n_path(X,I) :- state(X), range(I), I < k, n_path(X,J), J = I+1.

% Rule for (1).
:- n_path(X,0), goal(X), start(X).

% Rule for (2)
n_path(X,k) :- trans(X,A,Y), exo(X,A), n_path(Y,k).

% Rules for (3)
n_path(X,I) :- state(X), not goal(X), range(I),
I>0, not some_pass(X,I).\smallskip
some_pass(X,I) :- range(I), I>0, trans(X,A,Y), agent(A),
poss(X,A), not n_path(Y,J), I=J+1.

% Rule for (4)
:- n_path(X,k), start(X), not goal(X).

% Rule for (5)
n_path(X,0) :- state(X), not goal(X).
```

The predicate $range(I)$ specifies the index range from 0 to k , given by the input $limit(k)$. The rules encoding the clause groups (0)–(2) and (4), (5) are straightforward and self explanatory. For (3), we need to encode rules with bodies of different size depending on the transition function Φ , which itself is part of the input. We use that the antecedent of any implication (3) is true if it is not falsified, where falsification means that some atom s'_{i-1} , $s' \in PS(s)$, is false; to assess this, we use the auxiliary predicate $some_pass(X, I)$.

To compute the non-deterministic control K^+ , we may add the rule:

```
% Define  $\bar{C}_M$ 
cbar(X) :- state(X), not n_path(X,k).

%Define state level L
level(X,I) :- cbar(X), not n_path(X,I), I > 0, n_path(X,J), I=J+1.
level(X,0) :- cbar(X), not n_path(X,0).

% Define non-deterministic control k_plus
k_plus(X,A) :- agent(A), trans(X,A,Y), poss(X,A), level(X,I),
level(Y,J), J<I, not goal(X).
```

In $cbar(X)$, we compute the states in \bar{C}_M , and in $level(X, I)$ the level $\ell_{\bar{M}}(s)$ of each state $s \in \bar{C}_M$ ($= C_{\bar{M}}$ for the corresponding model \bar{M} of $sat(I)$). The non-deterministic control K_M^+ is then computed in $k_plus(X, A)$.

Finally, by the following rules we can non-deterministically generate any control which refines K_M^+ :

```
% Selecting a control from k_plus.
control(X,Y) :- k_plus(X,Y), not exclude_k_plus(X,Y).

exclude_k_plus(X,Y) :- k_plus(X,Y), control(X,Z), Y<>Z.
```

The first rule enforces that any possible choice for $K(s)$ must be taken unless it is excluded, which by the second rule is the case if some other choice has been made. In combination the two rules effect that one and only one element from $K_M^+(s)$ is chosen for $K(s)$.

Example 8. If the input representation of Example 5 is in a file `exa3.dlv` and the above program, denoted by Π_{det} , in a file `det.dlv`, the DLV engine can be invoked e.g. by

```
dlv exa3.dlv det.dlv -N=3 -filter=control
```

which outputs the controls; here `-N=3` sets the range of integers dynamically supported by the engine to 3, and `-filter=control` effects that the answer sets are clipped to the predicate `control`. In the particular case, the output on the call is (apart from system version information)

```
{control(b,a), control(c,a), control(d,a)}
```

yielding the unique control which exists in this case. If we would add a further agent action a_2 to the action set, and extend the transition function by $\Phi(b, a_2) = c$, then a call of DLV for the respective representation would yield

```
{control(b,a2), control(c,a), control(d,a)}
{control(b,a), control(c,a), control(d,a)}
```

corresponding to the two alternative controls which emerge, since the agent can take either action a or action a_2 in state a .

6.3. Non-deterministic transition function Φ

As for deciding the existence of a k -maintaining control, the only change in the code for the deterministic case affects Step (3). The modified code is as follows, where $n_apath(X, A, I)$ intuitively corresponds to $\overline{X_A}_I$.

```
% Rules for (3); different from above
% (3.1)
n_path(X,I) :- state(X), not goal(X), range(I), I>0, not some_apass(X,I).

some_apass(X,I) :- range(I), I>0, agent(A), poss(X,A), not n_apath(X,A,I),
                  not goal(X).

% (3.2)
n_apath(X,A,I) :- agent(A), trans(X,A,Y), poss(X,A), range(I), I>0,
                  n_path(Y,J), I=J+1, not goal(X).

% (3.3)
n_apath(X,A,I) :- agent(A), poss(X,A), range(I), I>0, I<k,
                  n_apath(X,A,J), J=I+1, not goal(X).
```

Here, $some_apass(X, A, I)$ plays for encoding (3.1) a similar role as $some_pass(X, I)$ for encoding (3) in the deterministic encoding.

To compute the non-deterministic control K_M^+ , we may then add the following rules:

```
% Define  $\bar{C}_M$ 
cbar(X) :- state(X), not n_path(X,K), limit(K).

% Define state action level, alevel ( $\geq 1$ )
alevel(X,I) :- alevel_leq(X,I), I=J+1, range(J), not level_leq(X,J).

alevel_leq(X,I) :- cbar(X), not goal(X), poss(X,A), agent(A), I>0,
                  range(I), not n_apath(X,A,I).

% Define non-deterministic control k_plus
k_plus(X,A) :- agent(A), alevel(X,I), poss(X,A), not n_apath(X,A,I).
```

Here, the value of $\bar{\ell}_M(s)$ is computed in $alevel(X, I)$, using the auxiliary predicate $alevel_leq(X, I)$ which intuitively means that $\bar{\ell}_M(X) \leq I$.

For computing the controls refining K_M^+ , we can add the two rules for selecting a control from k_plus from the program for the deterministic case.

Example 9. Let us revisit the instance I_1 in Example 6. We get the DLV representation of I_1 by adding the fact `trans(c,a,f)` to the representation for I . Assuming that it is in a file `exa4.dlv` and the program Π_{ndet} in a file `ndet.dlv`, a call

```
dlv exa4.dlv ndet.dlv -N=3 -filter=control
```

yields no output (apart from some system version print), which is correct. On the other hand, if we consider the input I_2 for the variant of Example 6 (with agent action \mathbf{a}' possible in g and $\Phi(g, \mathbf{a}') = \{f, h\}$), then the output is

```
{control(b,a1), control(c,a), control(d,a), control(f,a), control(g,a1)}
```

(where $a1$ encodes \mathbf{a}'). Again, this is a correct result.

6.4. Layered use of negation

An important note at this point is that the programs Π_{det} and Π_{ndet} do not necessarily have models which correspond to the least models of the Horn theories $\bar{s}at(I)$ and $\bar{s}at'(I)$, respectively. The reason is that the use of negation `not some_pass(X,I)` and resp. `not some_apass(X,I)` may lead through cycles in recursion. Thus, not each control computed is necessarily maximal (even though the maximal controls will be computed in some models). Furthermore, because of cyclic negation it is not a priori clear that the part of the program deciding the existence of a control is evaluated by DLV in polynomial time. However, consistency (i.e., existence of an answer set) is guaranteed whenever $\bar{s}at(I)$ resp. $\bar{s}at'(I)$ has a model.

It is possible to modify Π_{det} such that the use of negation in recursion cycles is eliminated, by using standard coding methods to evaluate the body of the rule in (3). Namely, introduce for Π_{det} a predicate `all_true` and replace `not some_pass(X,I)` in the code for (3) with `all_true(X,I)`, which is defined such that `all_true(s,i)` represents that every $s'_{i-1} \in PS(s)$ is assigned true, which can be checked using a linear ordering \leq on $PS(s)$. However, we refrain from this here.

Notice that in the case where $PS(s)$ has size bounded by a constant c , we can use a predicate `ps` of arity $c+1$ to represent $PS(s) = \{s^{(1)}, \dots, s^{(l)}\}$ by a single fact `ps(s, s^{(1)}, \dots, s^{(l)}, \dots, s^{(l)})` where $s^{(l)}$ is reduplicated if $l < c$. It is then easy to express the clause (3).

We can similarly modify Π_{ndet} such that the use of negation in recursion cycles is eliminated, where we use a linear ordering on $\mathcal{A}_{ag} \cap \text{poss}(s)$ (or simply on \mathcal{A}_{ag} , assuming that there are not many agent actions overall). Finally, we can also use for the program Π_{det} simply an ordering of \mathcal{A}_{ag} , since the deterministic transformation $\Phi(s, a)$ is a (partial) surjective mapping of \mathcal{A} onto $PS(s)$, which guarantees that via $\mathcal{A} \cap \text{poss}(s)$ each $s' \in PS(s)$ can be accessed through Φ .

The modified programs use negation only in a stratified manner, and thus will be evaluated by DLV in guaranteed polynomial time in the size of the DLV representation of $\bar{s}at(I)$ and $\bar{s}at'(I)$, respectively.

6.5. State descriptions by variables

In many cases, states of a system are described by a vector of values for parameters which are variable over time. It is easy to incorporate such state descriptions into the LP encoding from above, and to evaluate them on answer set solvers provided that the variables range over finite domains. In fact, if any state s is given by a (unique) vector $s = \langle s^1, \dots, s^m \rangle$ $m > 0$, of values s^i , $1 \leq i \leq m$, for variables X_i ranging over nonempty domains, then we can represent s as fact `state(v_1^1, \dots, v_{r_1}^1)` and use a vector X_1, \dots, X_m of state variables in the DLV code, in place of a single variable, X . No further change of the programs from above is needed.

Similarly, we can easily accommodate actions $a(P_1, P_2, \dots, P_m)$ with parameters P_1, \dots, P_m (which is important) from a finite set if desired. However, here the rule defining `exclude_k_plus(X,Y)` should be replaced by all rules emerging if the atom $Y \leftrightarrow Z$ in the body is replaced by $Y_i \leftrightarrow Z_i$, $i \in \{1, \dots, m\}$ (assuming that Y and Z are replaced by Y_1, \dots, Y_m and Z_1, \dots, Z_m , respectively).

Another possibility to handle state descriptions by variables would be to implement a coding scheme, which maps each vector $s = \langle s^1, \dots, s^m \rangle$ into an integer $i(s)$, represented by fact `code(i(s), s^1, \dots, s^m)`.

Furthermore, we point out that the input need not consist merely of facts, but may also involve rules to define the predicates of the input representation more compactly. Finally, the facts for `action` can be dropped, since they are not referenced by any rule in programs Π_{det} and Π_{ndet} .

For illustration, we consider the buffer example from Section 3.

Example 10. Recall that states in the buffer example are given by pairs of integers $\langle i, j \rangle$ where i and j are the numbers of objects in buffer b_1 and b_2 , respectively. We thus use variables X_1, X_2 and Y_1, Y_2 in place of X and Y , respectively.

For buffer capacity of 3, $S = \{\langle 0, 0 \rangle\}$, $E = \{\langle 0, j \rangle \mid 1 \leq j \leq 3\}$, and $k = 6$, the input can be represented as follows:

```
state(X1,X2) :- #int(X1), #int(X2), X1 <= 3, X2 <= 3.

start(0,0).

goal(0,X2) :- state(0,X2).

trans(X1,X2,m_12,Y1,Y2) :- state(X1,X2), state(Y1,Y2), X1=Y1+1, Y2=X2+1.
```



```

trans(X1,X2,m_21,Y1,Y2) :- state(X1,X2), state(Y1,Y2), Y1=X1+1, X2=Y2+1.
trans(X,X2,proc,X,Y2) :- state(X,X2), state(X,Y2), X2=Y2+1.
trans(X1,X,ins,Y1,X) :- state(X1,X), state(Y1,X), Y1=X1+1.

```

```

poss(X1,X2,m_12) :- state(X1,X2), 1 <= X1, X2 <= 2.
poss(X1,X2,m_21) :- state(X1,X2), 1 <= X2, X1 <= 2.
poss(X1,X2,proc) :- state(X1,X2), 1 <= X2.
poss(X1,X2,ins) :- state(X1,X2), X1 <= 2.

```

```

agent(m_12). agent(m_21). agent(proc). exo(ins).

```

```

const k = 6.

```

Here, equalities $X1=0$ for $X1,X2$ in the rule defining goal and $X1=Y1$ in the definition of `trans(X,X2,proc,X,Y2)` etc are pushed through.

Invoking DLV, assuming the representation is stored in file `exa-buffer.dlv` and the expanded version of Π_{det} in a file `det2.dlv`, with

```

dlv exa-buffer.dlv det2.dlv -N=6 -filter=control

```

yields 13 models, of which encode different controls. Among the maximal controls is

```

{control(1,0,m_12), control(1,1,m_12), control(1,2,m_12), control(1,3,proc),
 control(2,0,m_12), control(2,1,m_12), control(2,2,proc), control(2,3,proc),
 control(3,0,m_12), control(3,1,proc), control(3,2,proc),
 control(3,3,proc)}

```

which is defined on all states outside E , and thus constitutes a 6-maintaining control for the whole system.

7. Computational complexity

In this section, we consider the complexity of constructing k -maintainable controls under various assumptions. To this end, we first describe the problems analyzed and give an overview of the complexity results. After that, the results are established in a separate subsection; the reader who is not interested in the technical proofs might safely skip it.

7.1. Problems considered and overview of results

Following the common practice, we consider here the decision problem associated with k -MAINTAIN, which we refer to as k -MAINTAINABILITY: Given a system $A = (\mathcal{S}, \mathcal{A}, \Phi, poss)$, a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$ of agent actions, sets of states $E, S \subseteq \mathcal{S}$, an exogenous function `exo`, and an integer $k \geq 0$, decide whether S is k -maintainable with respect to E in A . Furthermore, we also consider MAINTAINABILITY, which has the same input except k and asks whether S is maintainable with respect to E in A .

We consider the problems in two different input settings, in line with the previous sections:

Enumerative representation: The constituents of an instance I are explicitly given, i.e., the sets $(\mathcal{A}, \mathcal{S}, \mathcal{A}_{ag}, S$, and $E)$ in enumerative form and the functions $(\Phi(a, s), poss(s),$ and `exo`) by their graphs in tables.

State variables representation: A system state s is represented by a vector $s = (v_1, \dots, v_m)$ of values for variables f_1, \dots, f_m ranging over given finite domains D_1, \dots, D_m , while \mathcal{A} and \mathcal{A}_{ag} are given in enumerative form. We assume that polynomial-time procedures for evaluating the following predicates are available:

- $in_Phi(s, a, s')$, $in_poss(s, a)$, and $in_exo(s, a)$ for deciding $s' \in \Phi(s, a)$, $a \in poss(s)$, and $a \in exo(s)$, respectively.
- $in_S(s)$ and $in_E(s)$ for deciding whether $s \in S$ and $s \in E$, respectively.

Orthogonal to this, we also consider (1) general k versus constant k , in order to highlight the complexity of small windows of opportunity for maintenance; (2) absence of exogenous actions, to see what cost intuitively is caused by an adversary; and (3) non-deterministic versus deterministic actions.

The results of the complexity analysis are compactly summarized in Tables 1 and 2, in which unless stated otherwise, the entries stand for completeness results under logspace reductions. We assume that the reader is familiar with the classes **P** (polynomial time), **EXP** (exponential time), **L** (logarithmic workspace), **NL** (non-deterministic logarithmic work space), **co-NP** (co-non-deterministic polynomial time), and **PSPACE** (polynomial space) appearing in the tables, and refer to [57] and references therein for further background on complexity. By **LH** we denote the logarithmic time hierarchy [11,38], which is given by $\mathbf{LH} = \bigcup_{i \geq 0} \Sigma_i^{\log}$, where Σ_i^{\log} denotes the decision problems solvable on an alternating Turing machine in

Table 1Complexity of deciding $(k-)$ MAINTAINABILITY under enumerative representation (logspace completeness)

| +/- exogenous actions | k -MAINTAINABILITY | | MAINTAINABILITY |
|-----------------------|----------------------------------|------------------------------------------------------------|--------------------------------------|
| | given k | constant $k \geq 1$ | |
| deterministic | P / NL (Th. 11/15) | P / in LH (\subset L) (Th. 11/16) | P / NL (Co. 12/Th. 15) |
| non-deterministic | P (Th. 11/13) | P / in LH (\subset L) (Th. 11/16) | P (Co. 12/Th. 13) |

Table 2Complexity of deciding $(k-)$ MAINTAINABILITY under state variables representation (logspace completeness)

| +/- exogenous actions | k -MAINTAINABILITY | | MAINTAINABILITY |
|-----------------------|----------------------------------------|---------------------------------------|--------------------------------------------|
| | given k | constant $k \geq 1$ | |
| deterministic | EXP / PSPACE (Th. 18/21) | EXP / co-NP (Th. 18/22) | EXP / PSPACE (Co. 19/Th. 21) |
| non-deterministic | EXP (Th. 18/20) | EXP / co-NP (Th. 18/22) | EXP (Co. 19/Th. 20) |

logarithmic time with at most $i-1$ alternations between existential and universal states, starting in an existential state. Note that **LH** is strictly included in **L**. A more refined complexity assessment is given in Section 7.2. However, we refrain here from providing a sharp complexity characterization of the problems classified within **LH** in terms of completeness under a suitable notion of reduction, since they are not central to the maintainability issue under an “adversarial” environment.

Under enumerative representation (Table 1), $(k-)$ MAINTAINABILITY has the same complexity as Horn SAT, which is **P**-complete [57]. In fact, this holds also for the case of constant $k = 1$ and the restriction that all actions are deterministic and that there is a single exogenous action. Thus, even in the simplest setting with an adversary according to the dimensions above, the problem already harbors its full complexity; excluding non-deterministic actions and/or fixing k does not make the problems simpler. Intuitively, this is because with the help of exogenous actions, one can simulate non-determinism and split sequences of agent maintenance actions into small segments.

On the other hand, when exogenous actions are excluded (listed under “-”), $(k-)$ MAINTAINABILITY is always easier when the actions are deterministic or the maintenance window is small (k is constant). In summary, the results show that exogenous actions can *not* be compiled efficiently away (with reasonable complexity) to an instance of maintainability under a small maintenance window, and that non-deterministic actions are indispensable for such a compilation.

The reason is that in absence of exogenous actions, k -MAINTAINABILITY is akin to a graph reachability resp. planning problem (for the latter, see Section 8.3). Indeed, define for a fixed system $A = (\mathcal{S}, \mathcal{A}, \Phi, poss)$, a set of agent action $\mathcal{A}_{ag} \subseteq \mathcal{A}$, and sets $E, S \subseteq \mathcal{S}$ of states the predicates $r_i(s)$, $i \geq 0$, on $s \in \mathcal{S}$ inductively by

$$\begin{aligned}
 r_0(s) &= s \in E, \\
 r_{i+1}(s) &= s \in E \vee \exists a \in \mathcal{A}_{ag} \cap poss(s) \\
 &\quad \forall s' \in \mathcal{S} (s' \in \Phi(s, a) \Rightarrow r_i(s')), \quad \text{for } i \geq 0.
 \end{aligned} \tag{1}$$

Informally, $r_i(s)$ expresses that some state in E can be reached from s within i agent actions, and it holds that S is k -maintainable with respect to E , exactly if $r_k(s)$ holds for every s in S (as proved in Lemma 1 below). The predicate $r_k(s)$ is definable in first-order predicate logic with a suitable relational vocabulary (using the predicates given for enumerative representation). As well-known, the first-order definable properties are those which can be decided in **LH** [11,38]. Since **LH** is considered to contain problems which have much lower complexity than hard problems in **P**, the effect of exogenous actions is drastic in complexity terms.

Under state variables representation (Table 2), the complexity of the problems, with few exceptions increases by an exponential. This increase is intuitively explained by the fact that state variables permit in general an exponentially smaller input representation, which must be unpacked for solving the problem. The exception for constant k in absence of exogenous functions, where the complexity increases from within **LH** to **co-NP**, is intuitively explained by the fact that the quantifier “ $\exists a \in \mathcal{A}_{ag} \cap poss(s)$ ” in Eq. (1), as opposed to “ $\forall s' \in \mathcal{S}$ ”, ranges over a polynomial set of values (in the input size), and thus can be deterministically eliminated. Exogenous actions cannot be compiled efficiently away in the same cases as under enumerative representation.

For practical concerns, we can draw from the results above the following conclusions. While k -MAINTAINABILITY is tractable under enumerative representation, it is because of its **P**-completeness not amenable to efficient parallel computation and not solvable within poly-logarithmic workspace under widely believed complexity hypotheses. However, if exogenous actions are absent and the maintenance window has size bounded by a constant, the problem can be solved in constant time using a polynomial number of processors as follows from membership in **LH** (see [38]).

The **EXP**-completeness results for state variables representation imply that the problems are provably intractable, and that exponential time and, by current methods, also exponential workspace is needed to solve them. Thus, a polynomial-time reduction to popular propositional logic engines such as SAT/UNSAT or QBF solvers (see <http://www.satcompetition.org/>, <http://www.qbflib.org/> for state-of-the-art systems) is infeasible in general. Only in the “cheapest” cases (where the size of

the maintenance window is bounded by a constant and exogenous actions are absent), a polynomial-time reduction of k -MAINTAINABILITY to SAT/UNSAT solvers is feasible; a polynomial-time reduction of $(k-)$ MAINTAINABILITY to QBF solvers is only feasible in deterministic domains and in absence of exogenous actions. When exogenous actions are possible, the full complexity shows up and one has to resort to more expressive engines such as answer set solvers (as discussed in Section 6.5) for instance. We remind, however, that the results in Table 2 are worst-case complexity results, and that under further constraints the problems may be solvable with polynomial resources. A detailed study of this issue remains for future work.

7.2. Enumerative representation

We start with the case of enumerative representation. Our first result is the following.

Theorem 11 (Problem k -MAINTAINABILITY is **P**-complete (under logspace reductions)). *The **P**-hardness holds under the restriction that $k = f(A, S, E)$ is any function of A, S , and E such that $f(A, S, E) \geq 1$ (in particular, for fixed $k \geq 1$), even if in addition all actions are deterministic and there is only one exogenous action.*

Proof. The membership of k -MAINTAINABILITY in **P** follows from Corollary 8.

We prove **P**-hardness under the stated restriction by a reduction from deciding logical entailment $\pi \models q$ of a propositional atom q from a propositional Horn logic program (PHLP) π , which is a set of rules of the form

$$b_0 \leftarrow b_1, \dots, b_n, \quad n \geq 0, \quad (2)$$

and each b_i is a propositional atom from an underlying atom set At ; b_0 is the head and b_1, \dots, b_n is the body of the rule.

As well-known, $\pi \models q$ holds iff there is a sequence of rules r_1, r_2, \dots, r_m , $m \geq 1$, from π where r_i is of form $b_{i0} \leftarrow b_{i1}, \dots, b_{in_i}$, such that $\{b_{i1}, \dots, b_{in_i}\} \subseteq \{b_{10}, \dots, b_{i-1,0}\}$, for all $i \in \{1, \dots, m\}$ (thus in particular, $1_n = 0$) and $b_{m0} = q$, called a *proof of q from π* . Informally, q is derived by successive application of the rules r_1, \dots, r_m , where r_i “fires” after all previous rules r_1, \dots, r_{i-1} have fired.

A natural idea is to represent backward rule application r_m, r_{m-1}, \dots, r_1 through agent actions; for a rule r of form (2), there is an agent action a_r which applied to a state s_{b_0} representing b_0 , brings the agent non-deterministically to any state s_{b_i} representing b_i , $i \in \{1, \dots, n\}$. Given a state s_q encoding q , $S = \{s_q\}$ is maintainable w.r.t. a set of states E encoding the facts in π if q has a proof from π . However, this does not account for the restriction that $k = f(A, S, E)$ for any such f . The key for this is to establish the result for the extremal case where $k = 1$ is constant (i.e., for 1-MAINTAINABILITY) and then to extend it to the general case.

Using a constrained rule format in π and an exogenous action, we can emulate non-deterministic agent actions and sequences of agent actions with some coding tricks by alternating sequences of deterministic agent and exogenous actions, such that provability of q from π corresponds to 1-maintainability of S w.r.t. a set E in a system A constructible in logarithmic workspace from q and π .

Without loss of generality, we assume that each rule has either zero or two atoms in the body (i.e., $n = 0$ or $n = 2$ in (2)). We construct from π and q a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, sets of states S and E , a set $\mathcal{A}_{ag} \subseteq \mathcal{A}$, and a function exo as follows:

1. S : For each atom f in π and rule $r \in \pi$, f^0, \dots, f^m and r^1, \dots, r^m are states in S . Furthermore, if the body of r is u, v then $(u, v)^0, \dots, (u, v)^{m-1}$ are states in S .
2. $\mathcal{A} = \{a_r \mid r \in \pi\} \cup \{e\}$.
3. Φ : For any rule $r \in \pi$ with head f , $\Phi(a_r, f^i) = \{r^i\}$ for $i \in \{1, \dots, m\}$ and $\Phi(a_r, (f, v)^i) = \{r^i\}$, for $(f, v)^i \in S$, $i \in \{1, \dots, m-1\}$. If moreover r has body u, v , then $\Phi(e, r^i) = \{(u, v)^{i-1}\}$, and $\Phi(e, (u, v)^{i-1}) = \{v^{i-1}\}$, for $i \in \{1, \dots, m-1\}$. In all other cases, $\Phi(a, s) = \emptyset$.
4. poss : For each state s , $\text{poss}(s) = \{a \in \mathcal{A} \mid \Phi(a, s) \neq \emptyset\}$.
5. $E = \{r^1, \dots, r^m \mid r \in \pi\}$.
6. $S = \{q^m\}$.
7. $\mathcal{A}_{ag} = \mathcal{A} \setminus \{e\}$.
8. exo : for all rules $r \in \pi$ of form $f \leftarrow u, v$, $\text{exo}(r^i) = \{e\}$ for $i \in \{1, \dots, m\}$ and $\text{exo}((u, v)^j) = \{e\}$ for $j \in \{1, \dots, m-1\}$. For all other states s , $\text{exo}(s) = \emptyset$.

The transition diagram for the system constructed for $\pi = \{a \leftarrow b, c; b \leftarrow; c \leftarrow\}$ is shown in Fig. 11. Intuitively, the state f^i encodes that f can be derived from π with a proof of length at most i . This is propagated in backward rule application. Each agent action a_r selects a rule r to prove an atom f ; if the rule has a body u, v , the exogenous action pushes the agent to prove both u (from (u, v)) and v within decreased recursion depth.

We claim that $\pi \models q$ iff there exists some 1-maintaining control K for S with respect to E in A .

Suppose first that $\pi \models q$. We then construct a 1-maintaining control K for S with respect to E as follows. Let $P = r_1, \dots, r_k$ be a proof of q from π such that, without loss of generality, all rules r_i have different heads. Set $D = \{q^m\}$ and iterate the following until D remains unchanged: For each $f^i \in D$ resp. $(u, v)^i \in D$, $i \geq 0$, let r_j be the rule with head f

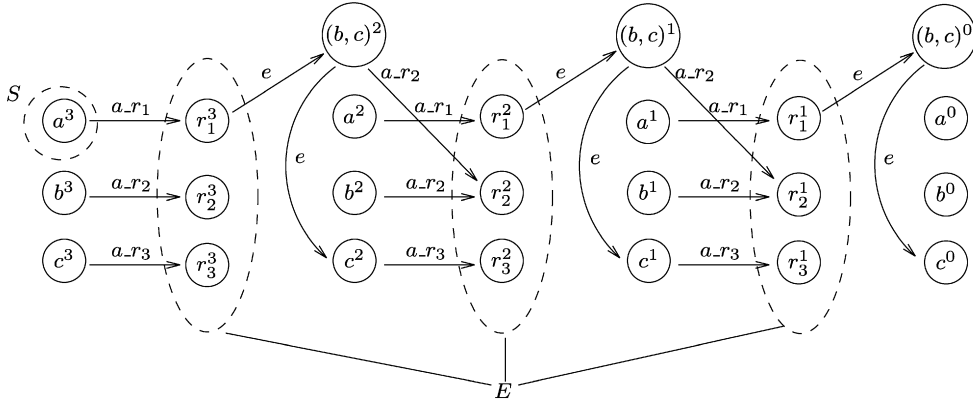


Fig. 11. Transition diagram of the system for $\pi = \{a \leftarrow b, c; b \leftarrow; c \leftarrow\}$ and $q = a$ (S and E encircled).

resp. u in P . Define $K(f^i) = \{a_{r_j}\}$ resp. $K((u, v)^i) = \{a_{r_j}\}$, and add, if r_j has body u', v' the states $(u, v)^{i-1}$ and v'^{i-1} to D . Since P is a proof of q from π , the rule r_j always exists, and for each state s in $\text{Closure}(S, A_{K, \text{exo}}) \setminus E (=D)$, $K(s)$ is defined and $\Phi(K(s), s)$ yields some state in E . Hence, K is a 1-maintaining control for S with respect to E in A .

Conversely, suppose K is a 1-maintaining control for S with respect to E in A . Without loss of generality, $K(s)$ is undefined for all states $s \in E$. An easy induction on $i \geq 1$ shows that for each $f^i \in \text{Closure}(S, A_{K, \text{exo}})$ resp. $(u, v)^i \in \text{Closure}(S, A_{K, \text{exo}})$, it holds that $\pi \models f$ resp. $\pi \models u$ and $\pi \models v$. For $i = 1$, suppose first $K(f^1) = a_{r_1}$. Rule r_1 must have form $f \leftarrow$; otherwise, some states $(u, v)^0, v^0$ would be in $\text{Closure}(S, A_{K, \text{exo}})$, which contradicts that K is a 1-maintaining control. Hence, $\pi \models f$. Next suppose $K((u, v)^1) = a_{r_1}$. Then, for similar reasons, r_1 must be of form $u \leftarrow$, hence $\pi \models u$. Furthermore, $v^1 \in \text{Closure}(S, A_{K, \text{exo}})$ and as already established $\pi \models v$. For $i > 1$, suppose $K(f^i) = a_{r_i}$. Then either r_i is of form $f \leftarrow$ and thus $\pi \models f$, or of form $f \leftarrow u, v$. In the latter case, $(u, v)^{i-1} \in \text{Closure}(S, A_{K, \text{exo}})$ and hence, by the induction hypothesis, $\pi \models u$ and $\pi \models v$. Consequently, $\pi \models f$. Similarly, if $K((u, v)^i) = a_{r_i}$, then either r_i is of form $u \leftarrow$ or of form $u \leftarrow u', v'$ and $(u', v')^{i-1} \in \text{Closure}(S, A_{K, \text{exo}})$, which by the induction hypothesis implies $\pi \models u'$ and $\pi \models v'$, thus $\pi \models u$. Since $v^i \in \text{Closure}(S, A_{K, \text{exo}})$, as already established $\pi \models v$. Consequently, $\pi \models f$. This proves the statement for $i > 1$, and concludes the induction. Since $q^m \in \text{Closure}(S, A_{K, \text{exo}})$, we have $\pi \models q$. This proves our claim.

Notice that A , S and E can be constructed in logarithmic workspace from π and q . This proves **P-hardness** of 1-MAINTAINABILITY. An easy observation is that every agent action in the system A leads to some state in the set E described. Hence, S is 1-maintainable with respect to E in A iff S is k -maintainable with respect to E in A for any $f(A, S, E)$ such that $f(A, S, E) \geq 1$. Hence, **P-hardness** under the stated restriction follows. \square

The following result is immediate from this result and the fact that maintainability is equivalent to k -maintainability where $k = |S|$ is the number of states.

Corollary 12. MAINTAINABILITY is **P-complete**. The **P-hardness** holds even if all actions are deterministic and there is only one exogenous action.

The following result states a further **P-complete** restriction of the above problems.

Theorem 13. k -MAINTAINABILITY and MAINTAINABILITY with no exogenous actions are **P-complete**.

Proof. Membership in **P** was established above. The **P-hardness** follows from Theorem 11 by merging the (single) exogenous action e into the agent actions as follows: For each state s such that $e \in \text{exo}(s)$, redefine every action $a \in \text{poss}(s) \cap \mathcal{A}_{\text{ag}}$ by $\Phi(s, a) := \Phi(s, a) \cup \Phi(s, e)$. It is easy to see that given S and E , S is $|S|$ -maintainable w.r.t. E in the resulting system A' iff S is $|S|$ -maintainable w.r.t. E in A . Furthermore, A' is computable in logspace from A . This implies the result. \square

The hardness results above are at the border of the hardness frontier, in the sense that in the absence of exogenous actions and, in case of MAINTAINABILITY also non-determinism, the problems are no longer **P-hard**. The following lemma gives a useful characterization of k -maintainability for this purpose.

Lemma 14. Given a system $A = (S, \mathcal{A}, \Phi, \text{poss})$, a set of agents action $\mathcal{A}_{\text{ag}} \subseteq \mathcal{A}$, and a set of states E , a set of states S is k -maintainable with respect to E in absence of exogenous actions (i.e., exo is void), $k \geq 0$, iff $r_k(s)$ as in (1) holds for all $s \in S$.

Proof. For the only if direction, consider any 1-maintaining control K which without loss of generality is undefined on every $s \in E$. For every state $s \in \text{Closure}(S, A_{K, \text{exo}}) = \text{Closure}(S, A_K)$, let d_s be the distance of s from E under K , i.e., the

largest i such that $\sigma = s_0, s_1, \dots, s_i \in \text{Unfold}_k(s, A, K)$ where $s_0 = s$. By an easy induction on $d_s \geq 0$, we obtain using $K(s)$ as witness for a in (1), that $r_{d_s}(s), r_{d_s+1}(s), \dots, r_k(s)$ must hold for s . Hence, $r_k(s)$ holds for every $s \in S$.

Conversely, let for each $s \in S$ be i_s the least integer i such that $r_i(s)$ holds. If $i_s > 0$, then define $K(s) := a$ for some arbitrary action $a \in \mathcal{A}_{ag} \cap \text{poss}(s)$ witnessing (1) for $i + 1 = i_s$, otherwise (i.e., if $i_s = 0$ or $r_i(s)$ does not hold for any $i \geq 0$) let $K(s)$ undefined. Then, K is a k -maintaining control for S with respect to E , since by definition of the relations r_i , for each $s \in \text{Closure}(S, \mathcal{A}_K)$, and $\sigma = s_0, s_1, \dots, s_l \in \text{Unfold}_k(s, A, K)$ such that $s_0 = s$ it holds that $l \leq k$ and $s_l \in E$ (recall that, as tacitly assumed, $\Phi(a, s) \neq \emptyset$ for each $a \in \text{poss}(a)$). Hence, S is k -maintainable with respect to E . \square

We then establish the following result.

Theorem 15. k -MAINTAINABILITY and MAINTAINABILITY for systems with only deterministic actions and no exogenous actions are **NL**-complete.

Proof. In this case, deciding $r_i(s)$ for given $s \in S$ and $i \geq 0$ is in **NL**: If $s \notin E$, a proper a in (1) and $s' = \Phi(s, a)$ can be guessed and, recursively, $r_{k-1}(s')$ established, maintaining a counter i . This is feasible in logarithmic workspace in the representation size of A . By looping through all $s \in S$, it thus follows from Lemma 14 that deciding whether S is k -maintainable with respect to E , where $k \leq |S|$, is non-deterministically feasible in logarithmic workspace. This implies **NL**-membership of k -MAINTAINABILITY and MAINTAINABILITY. The hardness follows from a simple reduction of the well-known **NL**-complete REACHABILITY problem [57] to k -resp. MAINTAINABILITY: Given a directed graph $G = (V, E)$ and nodes $s, t \in V$, decide whether there is a directed path from s to t in G . Define $A = (S, \mathcal{A}, \Phi, \text{poss})$ such that $S = \mathcal{A} = V$, $\Phi(v, w) = w$, and $\text{poss}(v) = \{w \mid v \rightarrow w \in E\}$. Then, for $\mathcal{A}_{ag} = \mathcal{A}$, $S = \{s\}$ is $|V|$ -maintainable w.r.t. $E = \{t\}$ in A iff there is a directed path from s to t in G . Clearly, A is constructible in logarithmic workspace from G . This shows the **NL**-hardness. \square

In case of constant k , Eq. (1) is decidable by a straightforward deterministic recursive procedure in logarithmic workspace, even under non-determinism, since the recursion depth is bounded by a constant and each recursion level requires only logarithmic work space. Hence, k -MAINTAINABILITY is decidable in logarithmic space. A finer grained analysis that it is within the class Π_{k+1}^{\log} of the logarithmic time hierarchy, which is a much better upper bound and makes completeness for logspace (under suitable reductions) fairly unlikely.

We assume that the input I of k -MAINTAINABILITY for fixed k , is a relational structure \mathcal{M}_I with universe $U(\mathcal{M}_I) = S \cup \mathcal{A}$, and relations over $U(\mathcal{M}_I)$ for the predicates $\text{in_Phi}(s, a, s')$, $\text{in_poss}(s, a)$, $\text{in_exo}(s, a)$, $\text{in_S}(s)$ and $\text{in_E}(s)$ from above, and relations for the additional predicates $\text{ag_act}(a)$, $\text{in_S}(s)$, and $\text{in_A}(a)$ representing membership $a \in \mathcal{A}_{ag}$, $s \in S$ and $a \in \mathcal{A}$ for each $s, a \in U(\mathcal{M}_I)$, respectively. The structure \mathcal{M}_I is encoded in a standard way by a bit-string [38].

Theorem 16. Problem k -MAINTAINABILITY for systems without exogenous actions is in Π_{2k+1}^{\log} ($= \text{co-}\Sigma_{2k+1}^{\log}$), if $k \geq 0$ is constant.

Proof. Any first-order formula $\psi_1 \vee Qx\psi_2$ (resp. $\psi_1 \wedge Qx\psi_2$) such that ψ_1 has no free variables and $Q \in \{\exists, \forall\}$, is logically equivalent to $Qx(\psi_1 \vee \psi_2)$ (resp. $Qx(\psi_1 \wedge \psi_2)$). Exploiting this, $r_k(s)$ in (1) can be written, using the vocabulary from above, as a first-order formula $\phi_k(x)$ in prenex form $\exists x_1 \forall x_2 \exists x_3 \dots Q_k x_k \psi(x_1, \dots, x_k, x)$ where $\psi(x_1, \dots, x_k, x)$ is quantifier-free, such that for any element $s \in U(\mathcal{M}_I)$ of an input structure \mathcal{M} , the sentence $\text{in_S}(s) \wedge \phi_k(s)$ is true on \mathcal{M} iff $r_k(s)$ holds. Hence, by Lemma 14, k -maintainability of S w.r.t. E in A is definable by a Π_{k+1}^{\log} prenex sentence $\forall x_0 \exists x_1 \dots Q_k x_k \psi'(x_0, x_1, \dots, x_k)$, where $\psi'(x_0, x_1, \dots, x_k)$ is quantifier-free, on the above vocabulary. Whether a fixed such sentence is false on a given structure \mathcal{M}_I can be decided by an alternating Turing machine, starting in an existential state, in logarithmic time using k alternations [11,38]. Hence, the problem is in $\text{co-}\Sigma_{k+1}^{\log} = \Pi_{2k+1}^{\log}$. \square

We remark that the hardness results in this section can be further strengthened to the case where only 2 agent actions are available, but leave a proof of this to the interested reader.

7.3. State variables

The following is an easy lemma, which in combination with the results in the previous subsection implies most upper bounds in Table 2.

Lemma 17. For any instance of k -MAINTAINABILITY resp., MAINTAINABILITY in which states are represented by variables, the corresponding instance in ordinary (enumerative) form can be generated in polynomial workspace.

Using this lemma, we then prove the following result.

Theorem 18. Under state representation by variables, k -MAINTAINABILITY is **EXP**-complete. The **EXP**-hardness holds under the restriction that $k = f(A, S, E)$ is any function of A, S , and E such that $f(A, S, E) \geq 1$ (in particular, for fixed $k \geq 1$), even if in addition all actions are deterministic and there is only one exogenous action.

Proof. Membership in **EXP** follows easily from Lemma 17 and Theorem 11. The **EXP**-hardness is shown by a reduction from deciding inference $\pi \models p(t)$ of a ground atom $p(\bar{c})$ from a function-free Horn logic program π with variables (i.e., a datalog program), which consists of rules of the form

$$p_0(\bar{t}_0) \leftarrow p_1(\bar{t}_1), \dots, p_n(\bar{t}_n), \quad n \geq 0, \quad (3)$$

where each p_i is the name of a predicate of arity $a_i \geq 0$ and $\bar{t}_i = t_{i,1}, \dots, t_{i,n}$ is a list of constants and variables $t_{i,j}$; $p_0(\bar{t}_0)$ is the head and $p_1(\bar{t}_1), \dots, p_n(\bar{t}_n)$ the body of the rule.

It holds that $\pi \models p(\bar{c})$ iff there is a sequence rules r_i of the form $p_{i_0}(\bar{t}_{i_0}) \leftarrow p_{i_1}(\bar{t}_{i_1}), \dots, p_{i_n}(\bar{t}_{i_n})$ and substitutions θ_i for r_i , i.e., a mappings from the variables in r_i to the set of constants C_π in π , such that $\{p_{i_1}(\bar{t}_{i_1}\theta_i), \dots, p_{i_n}(\bar{t}_{i_n}\theta_i)\} \subseteq \{p_{1_0}(\bar{t}_{1_0}\theta_1), \dots, p_{i-1_0}(\bar{t}_{i-1_0}\theta_{i-1})\}$, for all $i \in \{1, \dots, m\}$ (thus in particular, $1_n = 0$) and $p_{m_0}(\bar{t}_{m_0}\theta_m) = p(\bar{c})$, called a *proof of $p(\bar{c})$ from π* . Informally, $p(\bar{c})$ is derived by successive application of the rule instances $r_1\theta_1, \dots, r_m\theta_m$, like in a propositional logic program.

Deciding whether $\pi \models p(t)$ is well-known to be **EXP**-complete, cf. [22]. The construction is similar in spirit to the one in proof of Theorem 11 but more involved.

To prove **EXP**-hardness of k -MAINTAINABILITY under the given restriction, we first focus on 1-MAINTAINABILITY, and we describe how to reduce $\pi \models p(\bar{c})$ in logarithmic workspace to deciding 1-maintainability of a set of states S w.r.t. a set of states E in an agent system A .

Without loss of generality, we make the following assumptions on π and $p(\bar{c})$:

- The set of constants occurring in π , C_π , is $\{0, 1\}$;
- each rule r in π has either zero or two atoms in the body;
- all rules in r are safe, i.e., each variable X occurring in the head of a rule r also occurs in the body;
- π uses only one predicate, p ;
- $c = (0, 0, \dots, 0)$.

Any problem $\pi \models p(\bar{c})$ can be transformed to an equivalent one of this form in logarithmic workspace.

Similar as in the propositional case, the idea is to represent a reversed proof $r_m, \theta_m, \dots, r_1\theta_1$ of $p(\bar{c})$ from π through agent actions, and model backward rule applications through agent actions; note that m ranges from 1 to 2^{a_p} , where a_p is the arity of p (thus m requires a_p bits). The problem here which makes this more complex is the fact that we must, for each rule r_i , also take θ_i into account. If r_i has a nonempty body, the candidates for θ_i are systematically generated by alternating agent and exogenous actions. For each possible such θ_i , the derivation of the body atoms $p(\bar{t}_{i_2}\theta_i)$ and $p(\bar{t}_{i_1}\theta_i)$ is then explored.

More precisely, for each ground atom $p(\bar{c})$, and $m \in \{0, \dots, 2^{a_p}\}$, we have a state $(\bar{c}, m, \text{prove})$ outside E which intuitively says that $p(\bar{c})$ is derivable within m ($0 \leq m \leq 2^{a_p}$) steps. For each rule r in π , there is an agent action a_r , which is possible on $(\bar{c}, m, \text{prove})$ if $m > 0$ and $p(\bar{c})$ unifies with the head $p(\bar{t})$ of r , and it results in the state $(\bar{c}, m, r, \text{apply})$, which is in E . For r of form $p(\bar{t}) \leftarrow p(\bar{t}_1), p(\bar{t}_2)$, two phases are now established: (1) the selection of a substitution θ for the variables X in r , and (2) the generation of states $(\bar{c}_1, m-1, \text{prove})$ and $(\bar{c}_2, m-1, \text{prove})$, where $\bar{c}_1 = \theta_1$ and $\bar{c}_2 = \theta_2$, for the recursive test.

As for 1) an exogenous action e pushes the agent from $(\bar{c}, m, r, \text{apply})$ to a state $(\bar{c}, m, (0, 0, \dots, 0), r, \text{sel}_\theta)$. Here $(0, 0, \dots, 0)$ is the substitution θ : $X_1 = 0, \dots, X_k = 0$ to all variables in r . By executing an agent action inc_θ on this state, this vector is incremented to $(0, 0, \dots, 0, 1)$, resulting in a state $(\bar{c}, m, (0, 0, \dots, 0, 1), r, \text{inc}_\theta)$ in E , from which e pushes the agent to a state $(\bar{c}, m, (0, 0, \dots, 1), r, \text{sel}_\theta)$, where $X_n = 1$ in θ . Here again inc_θ is possible, leading to a state $(\bar{c}, m, (0, 0, \dots, 1, 0), r, \text{inc}_\theta)$ in E from which e pushes the agent to the state $(m-1, t, (0, 0, \dots, 1, 0), r, \text{sel}_\theta)$. Here again an inc action is possible for the agent etc.

In each state $(\bar{c}, m, \theta, r, \text{sel}_\theta)$ such that $p(t\theta) = \bar{c}$, the agent might alternatively take the action *choose*, which brings her to the state $(\bar{c}, m, \theta, r, \text{chosen}_\theta)$ in E , which closes phase 1. The exogenous action e pushes the agent from this state to the state $(m, \bar{t}_1\theta, \bar{t}_2\theta, \text{do_split})$ out of E . From this state, e pushes the agent further to the state $(\bar{t}_1\theta, m-1, \text{prove})$, and the agent must take at $(m, \bar{t}_1\theta, \bar{t}_2\theta, \text{do_split})$ the action *split*, which brings her to the state $(\bar{t}_2\theta, m-1, \text{goto_prove})$ in E , from which e pushes the agent to $(\bar{t}_2\theta, m-1, \text{prove})$. Fig. 12 gives a summary of the steps in graphical form.




Fig. 12. Schematic transition diagram for backward application of rule $r: p(\bar{t}) \leftarrow p(\bar{t}_1), p(\bar{t}_2)$ with substitution θ to prove $p(\bar{c})$.

In this way, the derivation of $p(0, 0, \dots, 0)$ from π is encoded to deciding 1-maintainability of $S = \{(2^d, (0, 0, \dots, 0), \text{prove})\}$ with respect to the set of states E described above. Note that to prove $p(\bar{c})$ from π via rule r , only one instance of $r\theta$ must be chosen; the 1-maintaining control has to single out this θ , by proper placement of the action chosen_θ . The proof of correctness is along the lines of the respective one in Theorem 11.

Given the regular structure of the states and the easy checks and manipulations that need to be done for determining applicability of actions and determining the successor state, respectively, it is not difficult to see that a representation of the above 1-MAINTAINABILITY instance using state variables can be compiled from π and $p(0, 0, \dots, 0)$ in logarithmic work space (in particular, that the polynomial-time procedures for deciding the membership predicates $\text{in_Phi}(s, a, s')$, $\text{in_poss}(s, a)$, $\text{in_exo}(s, a)$, $\text{in_S}(s)$, and $\text{in_E}(s)$ can be provided in polynomial time). Note that this instance employs only deterministic actions, and there is a single exogenous action. This establishes **EXP**-hardness for 1-MAINTAINABILITY.

Furthermore, for A and E as constructed, each agent action results in a state in E . Thus, k -maintainability of S w.r.t. E in A , for any $k = f(A, S, E)$ such that $f(A, S, E) \geq 1$, is equivalent to 1-maintainability of S w.r.t. E in A . Hence, the reduction shows **EXP**-hardness of k -MAINTAINABILITY under the stated restriction. \square

Corollary 19. *Under state representation by variables, MAINTAINABILITY is **EXP**-complete. The **EXP**-hardness holds even if all actions are deterministic and there is only one exogenous action.*

Using Theorem 18 instead of Theorem 11, we can prove the following result similarly as Theorem 13:

Theorem 20. *Under state representation by variables and in absence of exogenous actions, k -MAINTAINABILITY and MAINTAINABILITY are **EXP**-complete.*

For the case without exogenous actions and with only deterministic actions, we have lower complexity:

Theorem 21. *Under state representation by variables, k -MAINTAINABILITY and MAINTAINABILITY for systems with only deterministic actions and no exogenous actions are **PSPACE**-complete.*

Proof. By well-known standard methods, a computation composed of a **PSPACE** computation A piped into an **NL** computation B (which is **NPSPACE** in the size of the input for A) can be redesigned as an **NPSPACE** computation. Since **NPSPACE** = **PSPACE**, membership of the problems in **PSPACE** thus follows from Lemma 17 and Theorem 15.

The **PSPACE**-hardness can be shown e.g. by a straightforward reduction from propositional STRIPS planning [16]. Rather than to introduce STRIPS here, we give for completeness sake a simple reduction from SUCCINCT REACHABILITY [57], which is the version of REACHABILITY where $G = (V, E)$ is such that the nodes v are given by the binary vectors $v = (v_1, \dots, v_n)$, $n \geq 1$, on $\{0, 1\}$ and the problem input consists of a Boolean circuit C_G with $2n$ inputs $v_1, \dots, v_n, w_1, \dots, w_n$ which outputs true iff $v \rightarrow w \in E$, and $s = (0, 0, \dots, 0)$ and $t = (1, 1, \dots, 1)$. We construct from this an instance of k -MAINTAINABILITY resp. MAINTAINABILITY as follows: $\mathcal{S} = V \times V$, described by $2n$ binary variables f_1, \dots, f_{2n} ; $\mathcal{A} = \{\text{inc}, \text{arc}\} = \mathcal{A}_{\text{ag}}$; $\Phi(v \times w, \text{inc}) = v \times w'$ such that $w' = w + 1$ modulo 2^n , and $\Phi(v \times w, \text{arc}) = w \times (0, 0, \dots, 0)$ if $v \rightarrow w$ in G and $\Phi(v \times w, \text{arc}) = v \times w$ otherwise; $\text{poss}(s) = \mathcal{A}$, for each state s . Then, the state $s = (1, 1, \dots, 1) \times (0, 0, \dots, 0)$ is $|\mathcal{S}|$ -maintainable with respect to $E = \{(1, 1, \dots, 1) \times (1, 1, \dots, 1)\}$ in A iff $(1, 1, \dots, 1)$ is reachable from $(0, 0, \dots, 0)$ in G . A state variable representation of A can be easily generated from the circuit C_G in logarithmic workspace. This implies **PSPACE**-hardness of the problems. \square

If the maintenance window is bounded by a constant, the problem is easier.

Theorem 22. *Under state representation by variables, k -MAINTAINABILITY for systems without exogenous actions and constant $k \geq 0$ is co-NP-complete.*

Proof. For a given $s \in \mathcal{S}$, falsity of $r_k(s)$ can be proved by exhibiting (assuming $s \notin E$), for each $a \in \mathcal{A}_{\text{ag}} \cap \text{poss}(s)$ a witness $w(s, a) \in \mathcal{S}$ such that $w(s, a) \in \Phi(s, a)$ and $r_{k-1}(w(s, a))$ is false, which in recursion can be proved similarly. For constant k , this leads to $O(|\mathcal{A}_{\text{ag}}|^k)$ many guesses $w(s, a)$, which is polynomial in the size of the input. By Lemma 14, it thus follows that deciding the complement of k -MAINTAINABILITY is in **NP**. This proves membership in co-NP.

The co-NP-hardness, for every $k \geq 0$, is a simple consequence that under representation by state variables, deciding whether $S \subseteq E$ is co-NP-complete (this can be shown, e.g., by a simple reduction from propositional unsatisfiability). \square

8. Discussion and conclusion

In this paper, we gave a formal characterization of maintenance goals and distinguished it from the notions of stabilizability and temporal goals of the form $\square \Diamond f$ (over all valid trajectories). We present several motivating examples that illustrate the need for our notion of maintainability. The basic idea being that for certain kinds of maintenance it is important that the maintaining agent be given a window of non-interference from the environment so that it can do the

maintenance. To formalize this we need to distinguish between the agent's actions and the environment's actions. In our formalization we define the notion of k -maintainability, where k refers to the maximum window of opportunity necessary for the maintenance. We then gave polynomial time algorithms to compute k -maintainable controls, which are linear-time for small k , and we analyzed the complexity of determining k -maintainability under various assumptions. One interesting aspect of our polynomial time algorithm is the approach that led to its finding: the use of a SAT encoding, and complexity results regarding the special Horn sub-class of propositional logic.

We next report on some experiments which we have carried out, and we then discuss related work, before we conclude.

8.1. Experimental results

The different methods for constructing k -maintaining controls in the previous sections have been implemented, in order to compare them on an experimental basis. More specifically, programs generating the SAT encoding and the Horn SAT encoding, the latter written as a propositional logic program, have been implemented in Java, and also the algorithm k -CONTROL. For SAT solving, zChaff has been employed and for evaluating the Horn SAT instances the answer set solvers DLV and Smodels plus its preprocessor Lparse. The logic programs with variables given in Section 6 are ready for use in DLV as described, and only minor adjustments are needed for Smodels. The encodings and implementations are, together with the descriptions of the domains, available at <http://www.public.asu.edu/~jzhao6/k-maintain/>.

To evaluate the performance of the implementations, we conducted experiments on the buffer domain described in Section 3, varying the size of the buffers, max, and the number k for a maintaining control. In all experiments, there was a single goal state, (0, 0) (i.e., both buffers are empty), and a single initial state, which in was either (1, 1) or (3, 5); note that in both cases, the smallest k for which a k -maintaining control exists is $k = 2 * \text{max} + 1$.

The results of the experiments are shown in Table 3. They were collected on a Dell desktop with an Intel Pentium 4 (2.53 GHz), 512 MB main memory, and 753 MB swap space, under Slackware 11.0 including Linux kernel 2.4.33.3. We used a pre-release (beta-version) of DLV version 2007-10-11,⁸ Lparse 1.0.17, and Smodels 2.32. In the table, the leftmost column shows the buffer size max and the parameter k ; the rightmost column tells whether a solution exists or not. The times reported are for deciding whether a k -maintaining control exists, i.e., computing a model, from which a control can be efficiently extracted (in fact, in linear time); also for the genuine algorithm, output of a control in Step 4 is fast.

For the SAT and the Horn SAT encodings, the column "instance" shows the timings for generating the respective instances in Section 6.3 from the input representation in Section 6.1, and the other columns show the timings for solving the instances using a respective solver; the numbers in parentheses is the total of instance generation and solving. Since Lparse, the preprocessor of Smodels, already solves the Horn SAT instances, the call of Smodels has actually been omitted. The column "genuine algorithm" shows the results for the implementation of Algorithm k -CONTROL, and the columns right of it the results for the logic programming encodings, both the one for deterministic and for non-deterministic transition function. Here the problem input was not in explicit form by facts but described by rules as in Section 6.5. However, generating the factual representation takes only short time (between 14 and 115 ms for larger problem size) and is negligible compared to the time required for solving the instance.

The experiments show some interesting results. Among the SAT and Horn SAT solving methods, the logic programming engines perform overall better than the SAT solver. This is explained by the fact that generating the SAT instance for the solver takes much longer than generating the corresponding Horn logic program. One possible reason is that the input format of the SAT solver requires that strings are mapped to integers which represent propositional variables. This has been done using a hashmap with vectors; a better design of the data structure may save some time on this translation. On the other hand, the SAT solver needs less time to solve an instance than the LP engines. We remind, though, that the latter are committed to compute a special model, while the SAT solver may compute an arbitrary model of the instance.

The SAT solver scales reasonably well on the instances, and similarly DLV which is within a small factor in most cases. Moreover, DLV also scales well regarding the overall time, while Smodels showed some weakness here. Both SAT solving and Horn logic programs have limitations when the constructed instances become larger, since the memory may be exhausted (however, as for memory the experimental setting was modest).

The genuine algorithm is faster than the SAT solving method on larger instances, but slower than Horn SAT solving. The latter may be explained by the fact that the internal Java data structures used (vectors and maps) are not optimal. However, the implementation scales well with respect to k .

The logic programming encodings with variables behave similar to the Horn SAT encodings, and apart from one case the timings for the non-deterministic LP encoding in Lparse+Smodels are comparable to those for the Horn SAT method using DLV. The non-deterministic LP encoding in DLV is in many cases faster than the one of Lparse+smodels, but shows a less smooth scaling in some cases. This discrepancy may be explained by the different heuristics which is used by the systems in the model search. The deterministic encodings are, in both cases, faster than the non-deterministic counterparts, but the speedup is limited by a small factor (between 2 and 3).

We remark that extracting an actual control out of the model computed an LP engine, by adding the rules in Section 6.2 to the program does not scale well for large k in general. The reason is that the rule for k_plus has, for each single

⁸ Kindly provided by the DLV team. The official release 2007-10-11 behaves similarly.

Table 3Experimental results for deciding k -maintainability in the Buffer Domain (times in secs)

| start state (1,1), goal state (0,0) | | | | | | | | | | | | | | |
|-------------------------------------|---------------|----------|------------|----------|--------------------|-----------|-------|----------|----------------------|--------|--------|------------------|--------|--------------------|
| Problem max, k | SAT encoding | | | Horn SAT | | | | | Genuine Algorithm | DLV | | Lparse + Smodels | | solution yes/no |
| | instance | (zChaff) | | instance | (Lparse + Smodels) | (DLV) | | | | det | ndet | det | ndet | |
| 10,5 | 0.648 | 0.021 | (0.669) | 0.465 | 0.435 | (0.900) | 0.118 | (0.583) | 1.819 | 0.094 | 0.157 | 0.257 | 0.405 | no |
| 10,10 | 1.439 | 0.062 | (1.501) | 0.533 | 0.842 | (1.375) | 0.203 | (0.736) | 3.127 | 0.171 | 0.308 | 0.365 | 0.691 | no |
| 10,15 | 3.320 | 0.079 | (3.399) | 0.613 | 1.211 | (1.824) | 0.271 | (0.884) | 4.269 | 0.227 | 0.485 | 0.486 | 0.955 | no |
| 10,20 | 6.543 | 0.100 | (6.643) | 0.650 | 1.676 | (2.326) | 0.380 | (1.030) | 5.576 | 0.282 | 0.607 | 0.615 | 1.245 | no |
| 10,21 | 7.421 | 0.209 | (7.630) | 0.665 | 1.606 | (2.271) | 0.599 | (1.264) | 5.598 | 0.399 | 0.866 | 0.769 | 1.536 | yes |
| 10,25 | 12.335 | 0.276 | (12.611) | 0.688 | 1.770 | (2.458) | 0.743 | (1.431) | 6.235 | 0.523 | 1.771 | 0.971 | 1.619 | yes |
| 10,30 | 19.467 | 0.346 | (19.813) | 0.761 | 1.656 | (2.417) | 0.865 | (1.626) | 5.440 | 2.022 | 4.444 | 1.141 | 1.929 | yes |
| 10,35 | 26.236 | 0.398 | (26.634) | 0.801 | 1.809 | (2.610) | 1.007 | (1.808) | 5.412 | 2.916 | 6.664 | 1.284 | 2.226 | yes |
| 10,40 | 35.321 | 0.417 | (35.738) | 0.851 | 2.034 | (2.885) | 1.220 | (2.071) | 5.422 | 4.310 | 10.935 | 1.445 | 2.534 | yes |
| 10,45 | 45.195 | 0.463 | (45.658) | 0.883 | 2.111 | (2.994) | 1.455 | (2.338) | 5.418 | 5.481 | 14.787 | 1.628 | 2.831 | yes |
| start state (3,5), goal state (0,0) | | | | | | | | | | | | | | |
| 20,5 | 7.933 | 0.121 | (8.054) | 2.347 | 1.602 | (3.949) | 0.425 | (2.772) | 23.177 | 0.334 | 0.655 | 0.934 | 1.531 | no |
| 20,10 | 35.043 | 0.178 | (35.221) | 2.532 | 3.196 | (5.728) | 0.814 | (3.346) | 43.569 | 0.623 | 1.363 | 1.595 | 2.765 | no |
| 20,15 | 88.729 | 0.249 | (88.978) | 2.735 | 4.240 | (6.975) | 1.152 | (3.887) | 67.770 | 0.915 | 2.046 | 2.229 | 4.022 | no |
| 20,20 | 164.208 | 0.332 | (164.540) | 2.933 | 6.299 | (9.232) | 1.624 | (4.557) | 82.192 | 1.192 | 2.878 | 2.862 | 5.294 | no |
| 20,25 | 263.623 | 0.409 | (264.032) | 3.207 | 7.981 | (11.188) | 1.955 | (5.162) | 102.264 | 1.561 | 3.752 | 3.569 | 6.573 | no |
| 20,30 | 397.018 | 0.519 | (397.537) | 3.250 | 9.709 | (12.959) | 2.359 | (5.609) | 120.628 | 1.844 | 4.704 | 4.187 | 7.848 | no |
| 20,35 | 564.108 | 0.595 | (564.703) | 3.523 | 11.138 | (14.661) | 3.166 | (6.689) | 144.418 | 2.274 | 5.932 | 4.856 | 9.108 | no |
| 20,40 | 684.858 | 0.631 | (685.489) | 3.691 | 12.757 | (16.448) | 3.560 | (7.251) | 151.763 | 2.724 | 6.489 | 5.536 | 10.381 | no |
| 20,45 | 852.107 | 1.990 | (854.097) | 3.758 | 13.119 | (16.877) | 8.353 | (12.111) | 162.128 | 3.907 | 14.954 | 7.523 | 12.224 | yes |
| 20,50 | 1073.237 | 2.315 | (1075.552) | 3.889 | 13.650 | (17.539) | 7.885 | (11.774) | 172.677 | 24.345 | 22.828 | 8.309 | 13.655 | yes |
| 20,55 | 1323.519 | 2.526 | (1326.045) | 4.071 | 13.923 | (17.994) | 8.247 | (12.318) | 193.645 | 48.505 | 35.543 | 9.398 | 15.083 | yes |
| 20,60 | 1630.561 | 3.448 | (1634.009) | 4.353 | 28.818 | (33.171) | 8.795 | (13.148) | 152.911 | 79.994 | 55.327 | 10.496 | 16.530 | yes |
| 20,5 | 7.646 | 0.125 | (7.771) | 2.483 | 1.639 | (4.122) | 0.432 | (2.915) | 22.734 | 0.355 | 0.642 | 1.047 | 1.533 | no |
| 20,10 | 33.814 | 0.174 | (33.988) | 2.564 | 3.146 | (5.710) | 0.860 | (3.424) | 41.956 | 0.623 | 1.349 | 1.611 | 2.790 | no |
| 20,15 | 86.259 | 0.248 | (86.507) | 2.767 | 4.661 | (7.428) | 1.206 | (3.973) | 59.228 | 0.932 | 2.012 | 2.249 | 4.026 | no |
| 20,20 | 158.087 | 0.328 | (158.415) | 2.983 | 6.252 | (9.235) | 1.656 | (4.639) | 78.784 | 1.220 | 2.803 | 2.896 | 5.344 | no |
| 20,25 | 252.355 | 0.398 | (252.753) | 3.113 | 7.781 | (10.894) | 1.925 | (5.038) | 95.766 | 1.596 | 3.592 | 3.573 | 6.671 | no |
| 20,30 | 383.141 | 0.468 | (383.609) | 3.307 | 9.242 | (12.549) | 2.319 | (5.626) | 115.793 | 1.897 | 4.588 | 4.210 | 7.944 | no |
| 20,35 | 525.430 | 0.547 | (525.977) | 3.444 | 10.853 | (14.297) | 3.115 | (6.559) | 132.549 | 2.247 | 5.834 | 4.850 | 9.106 | no |
| 20,40 | 696.514 | 0.643 | (697.157) | 3.607 | 13.342 | (16.949) | 3.388 | (6.995) | 150.920 | 2.653 | 6.285 | 5.619 | 10.363 | no |
| 20,45 | 893.684 | 3.417 | (897.101) | 3.839 | 13.781 | (17.620) | 7.721 | (11.560) | 161.092 | 3.593 | 13.752 | 7.393 | 12.216 | yes |
| 20,50 | 1072.981 | 2.831 | (1075.812) | 3.938 | 13.250 | (17.188) | 7.552 | (11.490) | 172.896 | 11.786 | 23.132 | 7.548 | 13.655 | yes |
| 20,55 | 1327.165 | 2.541 | (1329.706) | 3.981 | 13.547 | (17.528) | 7.681 | (11.662) | 185.470 | 20.229 | 38.767 | 8.924 | 15.047 | yes |
| 20,60 | 1601.926 | 3.335 | (1605.261) | 4.327 | 15.091 | (19.418) | 8.055 | (12.382) | 153.073 | 33.750 | 62.340 | 10.415 | 16.558 | yes |
| 30,5 | 51.410 | 0.219 | (51.629) | 10.176 | 3.557 | (13.733) | 0.867 | (11.043) | 125.925 | 0.871 | 1.424 | 2.628 | 3.904 | no |
| 30,10 | 206.631 | 0.364 | (206.995) | 10.569 | 6.913 | (17.482) | 1.862 | (12.431) | 235.732 | 1.732 | 3.003 | 4.428 | 7.215 | no |
| 30,15 | 496.045 | 0.541 | (496.586) | 10.754 | 10.107 | (20.861) | 3.025 | (13.779) | 332.056 | 2.718 | 5.593 | 6.323 | 10.395 | no |
| 30,20 | 866.907 | 0.723 | (867.630) | 11.153 | 13.788 | (24.941) | 3.668 | (14.821) | 444.367 | 4.090 | 6.784 | 8.228 | 13.693 | no |
| 30,25 | 1388.404 | 0.880 | (1389.284) | 11.595 | 26.746 | (38.341) | 4.664 | (16.259) | 535.193 | 3.929 | 9.013 | 10.308 | 17.057 | no |
| 30,30 | 1971.139 | 1.065 | (1972.204) | 12.306 | 25.319 | (37.625) | 7.637 | (19.943) | 646.284 | 4.718 | 12.413 | 12.072 | 20.498 | no |
| 30,35 | 2685.617 | 1.267 | (2686.884) | 12.522 | 124.525 | (137.047) | 7.333 | (19.855) | 742.747 | 5.846 | 14.634 | 13.823 | 23.927 | no |
| 30,40 | 3408.562 | 1.663 | (3410.225) | 13.252 | 665.565 | (678.817) | 7.675 | (20.927) | 863.109 | 7.185 | 16.401 | 16.230 | 27.347 | no |
| 30,45 | 4272.727 | 1.818 | (4274.545) | 18.476 | 681.040 | (699.516) | 8.421 | (26.897) | 940.408 | 8.269 | 19.305 | 17.985 | 30.798 | no |
| 30,50 | out of memory | | | 31.191 | 593.518 | (624.709) | 9.450 | (40.641) | 1065.352 | 9.656 | 21.914 | 19.646 | 34.278 | no |
| 30,55 | out of memory | | | | | | | | 1155.366 | 11.058 | 25.889 | 21.517 | 37.816 | no |
| 30,60 | out of memory | | | | | | | | 1262.594 | 13.481 | 31.304 | 23.919 | 41.146 | no |
| 30,65 | out of memory | | | | | | | | 1319.128 | 17.987 | 53.466 | 29.790 | 56.084 | yes |
| 30,70 | out of memory | | | | | | | | 1387.423 | 43.161 | 89.959 | 32.501 | 80.523 | yes |

possible transition (s, a, s') , i.e., $s, s' \in \mathcal{S}$ and $a \in \mathcal{A} \cap \text{poss}(s)$ such that $s' \in \Phi(s, a)$, quadratically many ground instances in k . Although only a single such instance is relevant, the solver can not predetermine it in the grounding phase. However, this problem can be easily circumvented by keeping the additional rules in a separate post-processing program, and feeding into it the model computed by the main program. Then, a control is quickly output, in time which is largely dominated by the time for model computation (in the worst case, in a few seconds).

ω -maintaining controls: We have also conducted some experiments with different methods for deciding the existence of an ω -maintaining control. More specifically, following the two-phase approach in Section 5.4, we have considered propositional logic programs, a genuine algorithm, and logic programs with variables; note that a SAT encoding for the two phases is not straightforward, since we need to compute transitive closure in both phases and phase 2 accesses the complement of the output of phase 1; a simple realization requires layered use of negation for that, which is done in all logic programming encodings (for generic programs with variables, unstratified negation is avoided using the method in Section 6.4).

Table 4

Experimental results for deciding maintainability in the Buffer Domain (times in secs)

| s; g | sol method/max | 10 | 20 | 30 | 40 | 50 | 60 | 70 | 80 | 90 | 100 |
|---------------------|----------------------|-------------|-------------|-------------|--------------|--------------|--------------|---------------|---------------|---------------|----------------|
| (1,1); yes (0,0) | propositional Lparse | 0.333+0.056 | 1.192+0.186 | 4.578+0.362 | 10.182+0.677 | 23.992+1.041 | 52.712+2.148 | 102.632+3.717 | 181.020+5.801 | 297.172+9.154 | 446.141+13.399 |
| | propositional DLV | 0.333+0.018 | 1.192+0.104 | 4.578+0.270 | 10.182+0.439 | 23.992+0.693 | 52.712+1.054 | 102.632+1.134 | 181.020+1.694 | 297.172+2.134 | 446.141+2.535 |
| | genuine algorithm | 1.034 | 20.266 | 152.918 | 629.925 | 1835.164 | 4576.097 | 9762.176 | – | – | – |
| | Lparse+Smodels | 0.044 | 0.226 | 0.631 | 1.522 | 3.278 | 6.206 | 10.949 | 17.919 | 28.304 | 42.432 |
| | DLV | 0.124 | 0.398 | 0.922 | 1.815 | 4.071 | 5.781 | 8.770 | 12.832 | 17.687 | 23.770 |
| (9,1); yes (5,5) | propositional Lparse | 0.342+0.094 | 1.174+0.224 | 4.975+0.496 | 10.765+0.675 | 25.335+1.237 | 55.492+2.154 | 102.401+3.672 | 180.949+5.808 | 301.249+8.514 | 445.678+12.074 |
| | propositional DLV | 0.342+0.065 | 1.174+0.136 | 4.975+0.330 | 10.765+0.456 | 25.335+0.672 | 55.492+0.984 | 102.401+1.172 | 180.949+1.534 | 301.249+2.103 | 445.678+2.446 |
| | genuine algorithm | 0.738 | 15.898 | 127.646 | 553.806 | 1315.840 | 2813.434 | 5097.606 | – | – | – |
| | Lparse+Smodels | 0.048 | 0.200 | 0.610 | 1.497 | 3.234 | 7.140 | 10.909 | 19.647 | 28.932 | 43.868 |
| | DLV | 0.056 | 0.388 | 0.903 | 1.725 | 3.626 | 5.656 | 8.653 | 12.048 | 17.505 | 23.106 |
| (3,2); no (4,4) | propositional Lparse | 0.357+0.072 | 1.184+0.213 | 4.647+0.376 | 10.191+0.630 | 23.965+1.054 | 52.574+2.074 | 102.743+3.640 | 180.601+6.251 | 302.436+8.869 | 445.867+12.019 |
| | propositional DLV | 0.357+0.017 | 1.184+0.129 | 4.647+0.222 | 10.191+0.316 | 23.965+0.507 | 52.574+0.807 | 102.743+1.006 | 180.601+1.282 | 302.436+2.392 | 445.867+2.177 |
| | genuine algorithm | 0.698 | 15.702 | 125.848 | 542.273 | 1281.064 | 2772.359 | 8973.392 | – | – | – |
| | Lparse+Smodels | 0.049 | 0.180 | 0.591 | 1.503 | 3.384 | 6.553 | 11.122 | 18.343 | 28.787 | 43.159 |
| | DLV | 0.040 | 0.259 | 0.617 | 1.322 | 2.414 | 4.035 | 6.286 | 9.292 | 13.152 | 17.803 |
| (1,9); no (7,4) | propositional Lparse | 0.345+0.090 | 1.101+0.217 | 4.588+0.374 | 10.173+0.688 | 24.085+1.162 | 52.655+2.092 | 102.092+3.633 | 180.900+5.771 | 298.540+8.529 | 446.833+11.969 |
| | propositional DLV | 0.345+0.017 | 1.101+0.093 | 4.588+0.196 | 10.173+0.297 | 24.085+0.497 | 52.655+0.800 | 102.092+1.013 | 180.900+1.291 | 298.540+1.844 | 446.833+2.170 |
| | genuine algorithm | 0.511 | 12.519 | 112.653 | 498.534 | 1523.882 | 3943.009 | 8735.724 | – | – | – |
| | Lparse+Smodels | 0.051 | 0.184 | 0.607 | 0.503 | 3.512 | 6.325 | 11.194 | 18.351 | 28.862 | 43.260 |
| | DLV | 0.066 | 0.214 | 0.581 | 1.280 | 2.377 | 4.005 | 6.219 | 9.190 | 12.922 | 18.442 |

Phase 2 is, in fact, realized in the logic programs via computing first the set of all states which must be pruned, and then complementing it.

The results of our experiments are summarized in Table 4. There, *s* is the single start state and *g* the single goal state. For each propositional LP program, both the times for creating it from the factual representation and for solving it are shown (in ms); a dash “–” means timeout after 10,000 seconds. The first two instances have a solution, while the latter two are not. As can be seen, the LP encodings with variables scale well, with a linear time trend (note that the instance size of the problem is quadratic in the buffer size). Unsurprisingly, the propositional LP encodings are evaluated faster than the LP encodings with variables, but a lot of time is needed to construct the respective programs; in total, using non-ground programs is much faster. The implementation of the genuine algorithm does not scale well, and is way slower than the LP encodings. However, similar as in the case of *k*-maintainability, the implementation does not use optimal data and storage structures and therefore has room for improvement.

We remark that the propositional LP encoding can be constructed faster (in the experiments, up to almost four times), if one exploits that the domain is deterministic. Furthermore, a variant of the DLV encoding with variables, in which pruning in phase 2 is focused to the states having an *a*-path to the goal set (determined in phase 1), shows very similar performance.

With respect to our methodological approach, we can observe the following. Using SAT solving and logic programming to construct *k*-maintaining and ω -maintaining controls is, at least in this example, not only of theoretical interest, but also shows value to obtain some quick implementations which perform reasonably well. Genuine algorithms that are extracted from SAT or logic programming solutions have potential, but they require implementation and optimization efforts in order to become highly efficient and significantly faster than the first shot methods. In particular, if the state space is not too large and the maintenance window small, a logic programming based approach is attractive, where we may keep in mind that further preferences or constraints may be imposed on solutions in a declarative manner.

8.2. Relation with earlier work on control synthesis

In this section, we compare our work with earlier work on control synthesis [1,10,19,41,49,60]. We start with the papers by Barbeau et al. [10] and Kabanza et al. [41] from the AI literature.

8.2.1. Relation with Barbeau et al.'s and Kabanza et al.'s work

Barbeau et al. [10] and Kabanza et al. [41] present a method to synthesize reactive plans⁹ based on progressing formulas in a linear temporal logic, called Metric Temporal Logic (MTL), which allows to specify durations in the temporal operators. If we ignore the durations of the operator then their specification language is the linear temporal logic LTL with operators \bigcirc , \square , \diamond and \mathcal{U} . Similar to ours, they allow actions to have non-deterministic transitions. However, [41] does not allow exogenous actions. Although [10] allows exogenous actions, it requires that exogenous actions and agent actions be interleaved. Thus to use their formulation one has to translate our formalism to theirs. For example one can compile away

⁹ Their notion of reactive plans are slightly different from our notion of control functions; in their case a state, which they call a world state, may have many associated plan states. For our type of goal, control functions in our sense can however be easily extracted from their “reactive plans” without a blowup.

exogenous actions¹⁰ from our formulation and, as mentioned earlier, introduce a fluent *interfered* to remember that a transition was made immediately made due to an exogenous action. The resulting transition function will have non-deterministic transitions and thus would need the specification $A_\pi(\Box(\phi_S \Rightarrow \Box(\text{Step}[k](\text{interfered} \vee \phi_E))))$ in the language of [9] to express k -maintainability. However, the control generation algorithm in [10,41] is able to construct a control for all goals of the form $A_\pi f$, where f is an LTL formula and the algorithm focuses on f .

Their algorithm is based on progressing an LTL formula. The intuition behind progression is that given a sequence of states $\sigma = s_0, s_1 \dots$ and an LTL formula f , the progression of f with respect to a state s_i satisfies the property: $(\sigma, i) \models f$ iff $(\sigma, i + 1) \models \text{progression}(f, s_i)$.

A partial definition of progression, as needed for our illustration, is as follows:

- for a proposition p , $\text{progression}(p, s)$ is true if p is true in s else is false.
- $\text{progression}(\neg f, s) = \neg \text{progression}(f, s)$
- $\text{progression}(f \wedge g, s) = \text{progression}(f, s) \wedge \text{progression}(g, s)$
- $\text{progression}(f \vee g, s) = \text{progression}(f, s) \vee \text{progression}(g, s)$
- $\text{progression}(\Box f, s) = \text{progression}(f, s) \wedge \Box f$
- $\text{progression}(\bigcirc f, s) = f$.

To illustrate how their algorithm works, we consider an important fragment of the LTL part of our goal: $\Box(\text{Step}[k](\text{interfered} \vee \phi_E))$. To simplify further, let φ denote $\text{interfered} \vee \phi_E$, and let s_φ denote a state where φ is true and $s_{\neg\varphi}$ denote a state where φ is not true. In the following, let $\text{progression}(f, s_1, s_2, \dots, s_n)$ denote $\text{progression}(\text{progression}(\dots \text{progression}(f, s_1), \dots), s_n)$. We now illustrate how progression works with respect to the above goal:

- $\text{progression}(\Box \text{Step}[k](\varphi), s_\varphi) = \Box \text{Step}[k](\varphi) = \text{Step}[k](\varphi) \wedge \bigcirc \Box \text{Step}[k](\varphi)$
- $\text{progression}(\Box \text{Step}[k](\varphi), s_{\neg\varphi}) = \text{Step}[k-1](\varphi) \wedge \Box \text{Step}[k](\varphi) = \text{Step}[k-1](\varphi) \wedge \bigcirc \Box \text{Step}[k](\varphi)$
- $\text{progression}(\Box \text{Step}[k](\varphi), s_{\neg\varphi}, s_{\neg\varphi}) = \text{Step}[k-2](\varphi) \wedge \bigcirc \Box \text{Step}[k](\varphi)$
- $\text{progression}(\Box \text{Step}[k](\varphi), s_{\neg\varphi}, s_{\neg\varphi}, s_{\neg\varphi}) = \text{Step}[k-3](\varphi) \wedge \bigcirc \Box \text{Step}[k](\varphi)$

and so on. As evident from the above the progression may lead to k different formulas. The general algorithm in [10,41] introduces decomposition due to the \vee connective in $\text{Step}[k](\varphi)$ and can have 2^k formulas labeling a state, leading to a search space of $2^k \times |I|$. Since one does not need to worry about cycles, one can avoid double exponential search.

However, noticing that our goal specification does not have unbounded eventualities, one can also modify the algorithm to avoid decomposition and thus restrict to a search space of $k \times |I|$. After that one needs to do an efficient search in this search space as the search algorithm given in [10,41] does not seem to do efficient backtracking.

In conclusion, although the algorithm in [10,41] as presented does not find a k -maintainable control in polynomial time, certain modifications make that algorithm more efficient. This illustrates the importance of the work [10,41]. It suggests the research direction of exploring various sub-classes of LTL goal specifications and identifying appropriate modifications or simplifications of the algorithms in [10,41] so that they lead to a more efficient control finding method. In slight contrast, our approach has been to find an efficient algorithm for a specific goal through logical manipulations. It led us to finding efficient and different algorithms for other goal specifications such as strong planning, strong cyclic planning and weak planning, which we discuss in [4].

8.2.2. Relation with other work on control synthesis

Most of the other works on control synthesis dates back earlier, around the time when various temporal logics were proposed for program specification purposes.

Clarke and Emerson [19] consider specifications in the branching time temporal logic CTL and present an algorithm to construct synchronization skeletons of concurrent programs from scratch. The algorithm constructs a finite model of the formula using a tableau-based procedure, and factors the control skeletons of the individual processes out from the global flow graph defined by the model. There is no complexity analysis given, but the authors mention that the algorithm is potentially exponential. Furthermore, the algorithm assumes a closed system, in which neither environment actions nor

¹⁰ One compilation involves the following: (i) For each state s a new state called s_{int} is created. The states s and s_{int} are equivalent with respect to all fluents except the newly introduced fluent *interfered*, which is false in the former and true in the later. (ii) All edges in the original transition diagram due to agent actions are kept. (iii) If there is a transition from s to s' due to an agent action a in the original transition diagram, then a transition from s_{int} to s' due to agent action a is introduced. (iv) For all transitions between s and s' due to a series of exogenous actions, if there is an edge from s'' to s due to an agent action a , then a transition from s'' to s'_{int} is introduced. (v) The set of initial states is enlarged as follows: If s is an initial state and there is a transition from s to s' due to a series of exogenous actions then s'_{int} is made an initial state.

However, the complexity results in Table 1, and the following discussions suggest that such a compilation has a high worst-case effort, and is not doable in non-deterministic logspace. This remains true with respect to any other compilation where the new k depends only on the old k .

It may be noted that with respect to our buffer example the above compilation is not linear in the size of the input; indeed, the original transition has $O(n^2)$ many nodes and arcs (where n is the buffer size), while the compiled one has $O(n^2)$ many nodes but $\Omega(n^3)$ many arcs, which come in through transitive closure of exogenous actions. Thus, in that example doing the compilation and then solving the LTL planning problem, to solve k -maintainability (for a fixed k), is not a linear time algorithm.

specific agent actions (which constrains the model building) are respected. However, this is crucial for the construction of a k -maintaining control; applied to the CTL specification of such a control for S w.r.t. E , the algorithm may return, e.g., a model which has only states corresponding to states outside S .

Independently, Manna and Wolper [49] consider the construction of the synchronization part of communicating processes from specifications in the linear time temporal logic PTL. As they mention, the main differences between their approach and [19] is the usage of PTL instead of CTL, and that [19] synthesizes shared memory programs. Manna and Wolper's synthesis method employs a tableau-style satisfiability algorithm, which is essentially the one in [12] restricted to linear time operators and modified to their assumption that in each state exactly one atomic proposition is true. The algorithm either declares the specification to be unsatisfiable or constructs a "model graph" from which all possible models of the PTL specification can be extracted. In the construction, PTL formulas are decomposed using the identities

$$\Box f \equiv f \wedge \Box f, \quad \Diamond f \equiv f \vee \Diamond f, \quad \text{and} \quad f_1 \cup f_2 \equiv f_2 \vee (f_1 \wedge \Box f_1 \cup f_2),$$

which enables to model a flow graph by talking about the current and the follow up state in an execution. As Manna and Wolper argue, the size of the model graph is at most exponential in the size of the specification formula. In a final step, the control code is generated from the model graph. Special care is here applied to so called eventuality conditions, which are conditions of the form $\Diamond f$ whose satisfaction must not be indefinitely postponed (this aspect was not addressed in detail in [19]). Since like [19], the method assumes a closed system without an environment, it is not readily applicable for constructing a k -maintainable control.

Pnueli and Rosner [60] give an algorithm to synthesize a reactive module with input x from and output y to an environment, where the values are from a finite domain, specified by a linear temporal formula $\varphi(x, y)$. The running time of the algorithm, which is constructed using automata-theoretic results, is at most double exponential in the length of the given specification; by its nature, it is difficult to say how it would perform for the construction of a k -maintaining control.

Abadi et al. [1] introduce a notion of realizability which is somewhat more general than the notion of implementability in [60], since it also considers an environment whose behavior is restricted. They distinguish environment and agent activity in changes to the environment, and define specifications at an abstract level as sets of behaviors, which are alternating sequences of states and active parties (which is either the environment, the agent, or none of them), satisfying some conditions. Realizability is given if a subset of the behaviors can be generated by the runs of a "computer", which is a restricted (possibly non-recursive) partial function from prefixes of behaviors to states; as noted, realizability is a necessary but not sufficient condition for the existence of a real implementation. A weaker notion of realizability, which takes into account that the implementor knows exactly how the environment behaves (in a deterministic manner), is considered and shown to be equivalent to realizability for an important class of specifications. Finally, the issue of realizability for finite-state transition systems P_t (represented by automata) equipped with a further restriction on the infinite behaviors of the systems, given by a (finite-state) Büchi automaton P_i , is considered. As pointed out, this is different from consistency of the behaviors of P_t and P_i . Applying, like Rosner and Pnueli, automata-theoretic results, it is argued that deciding realizability is in **EXPTIME** and **PSPACE**-hard under logspace reductions. An important note is that for realizability, Abadi et al. view non-deterministic transitions as optional outcomes of actions from which an implementor of a realization can choose an arbitrary subset. For our notion of control, it would be required that if a particular agent action is chosen at a state, then all outcomes of that action are chosen. Since a realization must merely be compatible with the infinite behaviors, but need not manifest all of them, enforcing such a choice via the infinite behaviors is infeasible. Hence, the method of [1] seems not always applicable for constructing a k -maintaining control in our setting.

Recently there have been some works in developing polynomial time control generation algorithms (in the size of the state space) for particular classes of LTL goals. In particular, the paper [59] presents a cubic algorithm in the size of the state space to automatically construct controls for $GR(1)$ goals which are of the form $(\Box \Diamond p_1 \wedge \dots \wedge \Box \Diamond p_m) \Rightarrow (\Box \Diamond q_1 \wedge \dots \wedge \Box \Diamond q_n)$, where p_i s and q_j s are propositional formulas. This complexity is still higher than our algorithm; however, it remains open to find out if our approach can lead to a more efficient algorithm for $GR(1)$ goals.

8.3. Other related work

Besides the related works we already mentioned such as stabilizability and temporal logic, the notion of maintenance has appeared in AI in many other papers. For example, in [55], Ortiz discusses maintenance actions. His notion of maintenance is stronger than both the notion of stabilizability and our notion as he requires the formula that is maintained to be true throughout. The notion of maintenance is also related to the notion of 'execution monitoring' which is studied in the context of robot programs in [23]. In 'execution monitoring' the world is monitored and if a discrepancy is found between the prediction made by the agent and the real world, then new plans are made to recover from the discrepancy. A deliberative architecture for maintenance can be extrapolated from the notions in [5], where an agent executes a cycle of *observe*; *assimilate*; *(re)plan_from_current_situation*; *execute_part_of_the_plan*.

In a series of papers [27,28,69], Wooldridge and Dunne have formalized the problem of constructing agent control functions and analyzed its complexity in a rich framework, for various kinds of tasks such as 'achievement' tasks (where the agent has to bring about a certain goal condition), 'maintenance' tasks (where the agent has to avoid that some goal condition is ever satisfied during execution), and combinations thereof [27]. Their notion of boolean task specification allows a seamless combination of 'achievement' and 'maintenance' goals. In their notion a goal specification is a propositional

formula where each proposition corresponds to a set of states. The intuitive meaning of a goal specification p would be to reach for sure one of the states corresponding to p . This corresponds to an achievement goal, while a specification $\neg p$ intuitively means that the agent should avoid any state in p . They refer to the later as a ‘maintenance task with bad states’. Thus their notion of ‘maintenance’ differs from ours. While we are concerned with the hindrance posed by the adversary explicitly, and take into account the number of steps when the adversary is not interfering, they do not take exogenous actions into account explicitly. They allow non-deterministic effect of actions, and although one can partially take exogenous actions into account by a straightforward compilation, keeping the count of the window of non-interference is not straightforward. Their control policies are richer than ours. In their framework, action effects and the selection of the agent action by the control may depend on the history of the execution. Under restriction to history-independent state transitions and reactive agents, finding controls for achievement tasks in their framework corresponds to finding maintaining controls with an unbounded window of opportunity in our framework. In [27] they do a comprehensive complexity analysis of agent design in their framework. There is a natural correspondence between some of our complexity results and theirs. In particular, our Theorems 15 and 21 correspond to the respective results in [27].

In AI planning, the seminal STRIPS approach [32] has been one of the most influential approaches. We briefly recall that in STRIPS, states are modeled as sets of propositional atoms and actions as operators which, given that a precondition in terms of a conjunction of literals is true on the current state, transform it to a successor state by removing atoms from a delete list and adding atoms from an add list. A plan for achieving a goal, described by a conjunction of atoms γ , from an initial state S_0 is a sequence of operators op_1, \dots, op_n which takes the agent from S_0 to a state where γ holds. STRIPS planning has been generalized in several directions, such as conditional effects, non-deterministic actions, or planning under incomplete information and partial observability using conditional and conformant plans, respectively, and a number of papers has considered the computation and complexity of planning in such settings, e.g., [6,16,18,31,64].

However, like in the framework of Wooldridge and Dunne, in none of these works agent actions and exogenous actions are viewed separately, and thus they are best compared to our framework in absence of exogenous functions. Furthermore, plans per se are conceived as *action strategies* (cf. [64]) in which, in principle, different actions might be taken by the agent if during plan execution the same state is entered again; however, such looping is a priori excluded if the goal must be achieved under all contingencies.

Daniele et al. [21] introduce the notion of strong cyclic planning which has some similarity with our notion of maintainability. In particular, both accept the possibility that an environment can sometimes be belligerent and in that case one needs to differentiate between an agent that is trying and an agent that is not. However, they encode the environmental interference through non-determinism while we allow explicit representation of environmental actions. Daniele et al. and later Cimatti et al. [18] consider constructing universal plans akin to our policies, with different semantics (weak, strong and strong cyclic) for goal achievement, based on OBDD methods and algorithms. In particular, in absence of exogenous actions our maintaining controls correspond to what they call strong solutions for a planning problem. For further discussion, we refer to [4].

As for complexity, Theorem 21, corresponds to the classical result of Bylander [16] that deciding plan existence in propositional STRIPS is **PSPACE**-complete, while Theorem 20 corresponds to Littman’s result that conditional planning for STRIPS with non-deterministic actions is **EXPTIME**-complete [46,64]. In conditional planning, via conditions on the current state branching to subplans is possible, such that an appropriate plan is followed depending on the state evolution. Branching might be modeled by actions and the conditional planning problem, with loops disregarded, as the problem of constructing a maintaining control.

In other related work, Jensen et al. [39,40] consider the somewhat dual problem of developing policies that achieve a given goal while there are interferences from the environment. In their model, environment actions and actions of multiple agents are combined to a joint action, by which the system is transferred from the current state to one out of a set of possible successor states. With such non-deterministic transitions, Jensen et al. aim at modeling both an adversarial environment and infrequent errors which make an otherwise deterministic action non-deterministic. In [39], they consider constructing policies coping with arbitrarily many interferences of the environment (but without action failure) by an extension of OBDD-based universal planning, and in [40] they consider generating policies which tolerate up to a given number n of errors modeled as “secondary action effects” (caused by improper action execution or environment interference), by reducing it to a so called strong planning problem, which is solved using OBDD based methods. For arbitrarily many environment interferences as in [39], the problem is basically very similar to our problem of unbounded maintainability, but interference in goal states has different significance and goal achievement is not guaranteed because of possible loops. A formal connection between k -maintainable controls and n -fault tolerant policies, if any, remains open. Intuitively, n -fault tolerant plans are easier to construct, since the number of errors that have occurred can be recorded in plan construction and when the limit n is reached, the problem boils down to an ordinary planning problem. For k -maintaining controls, however, each environment interference (even at a goal state) causes a restart which pushes the agent to a new initial state.

Outside of AI, our notion of k -maintenance is very closely related to the notion of self-stabilization in [24] which is used in characterizing fault-tolerant systems. There the concern is about proving correctness of (hand developed) self-stabilization protocols and achieving self-stabilization for various distributed algorithms such as mutual exclusion. Our algorithm here can be thought of as an algorithm that automatically generates a self-stabilization protocol. Although, this is a new dimension to the existing work on self-stabilization, further research is needed to compare assumptions made in our formulation and the ones in the self-stabilization literature, and overcome them. In particular, often in the self-stabilization literature the

global states are composed of local states of various distributed elements and a particular element does not have the access to the complete global state. In those cases one can not directly use the kind of global policies generated by the algorithm in this paper. We elaborate more on this in Appendix A.

8.4. Further work and open issues

There are several directions for further research extending the work of this paper. One direction concerns other classes of goal specifications, apart from k -maintainability and maintainability, for which controls can be synthesized in polynomial time. Another direction are more general execution models, for instance by taking action duration into account. In such a scenario, the maintenance goal may be formulated as requirement that the agent reaches some desired state always within a given time frame, if she is not disturbed by the environment. Preliminary investigations suggest that the results in this paper can be extended to handle this setting, and that alternatively the algorithm of Barbeau et al. [10] and Kabanza et al. [41] can be used for efficient computation, as long as durations are not long (binary number representation may lead to an increase in complexity).

The intractability results for the problems under state variable representations challenges methods and techniques for handling the problem in practice. Suitable heuristics may therefore be researched that allow to solve the problems in many cases in polynomial time, and, in a refined complexity analysis, meaningful tractable cases should be singled out. Furthermore, the issue of computing optimal k -maintenance controls efficiently, in the sense that k is as small as possible (which is trivially polynomially solvable in the enumerative setting), is an interesting issue for variable state representation.

Another issue concerns investigating computational transformations between maintenance and planning. By the complexity results in [46] and this paper, transformations between k -MAINTAINABILITY and conditional planning are feasible in polynomial time. It would be interesting to study different transformations, and to assess possible benefits of these transformations for solving k -MAINTAINABILITY and planning by cross-utilizing different algorithms and implementations (e.g. [18] for planning in non-deterministic domains). In particular a transformation similar to the one in the proof of Theorem 13, with an additional parameter that counts the number of agent actions since the last exogenous action, can¹¹ be used to compile out exogenous actions and transform finding k -maintainable policies to finding strong cyclic plans [18]. On the other hand, encodings similar to the one in Section 5.2 for obtaining strong cyclic plans through linear-time Horn logic programming can be designed. For more details and results on the latter, we refer to [4].

Acknowledgements

This work was partially supported by FWF (Austrian Science Fund) projects P-16536-N04 and Z29-N04, the European Commission under grant IST 2001-37004 WASP, the NSF (National Science Foundation of USA) grant numbers 0070463, and 0412000, NASA grant number NCC2-1232, and contracts from ARDA and DTO. We would like to acknowledge W. Cushing for his feedback on an earlier draft and S. Gupta and M. Gouda for their clarifications on self-stabilization. Furthermore, we acknowledge comments by J. Rintanen on the ICAPS'04 paper and are grateful for his pointers to related work. We owe special thanks to F. Kabanza for helping us understand his co-authored papers [10,41] and making observations regarding how the algorithms in those papers could be modified to make them more efficient for specific goal specifications. We furthermore appreciate the constructive comments of the reviewers to improve the presentation, the suggestion of an LTL formulation in Section 4.1 and the suggestion of a linear time algorithm for generic maintainability. Finally, we are indebted to J. Zhao for implementing some of the algorithms and running the experiments, and appreciate the support of the DLV team.

Appendix A. Self-stabilization and related notions in distributed computing

Earlier we remarked that our notion of maintainability has similarities with Dijkstra's notion of self-stabilization [24]. In this appendix, we elaborate on the relationship between our notions and Dijkstra's as well as similar notions in distributed computing [2].

A.1. Dijkstra's notion of self-stabilization

Dijkstra considers in [24] a connected graph (in which a majority of edges are missing) with a finite state machine placed at each node; machines placed in directed connected nodes are called each other's neighbors. For each node (or machine), a Boolean function over the state of the machine and the states of its neighbors, called a "privilege", is defined. There is a central daemon that can select one of the "privileges" (over the whole graph) that is true. The machine whose privilege is selected can move its state to another state based on a policy which is a function from the state of the machine and its neighbor's states to a move that results in a state transition in the machine under consideration. If for such a machine more than one privilege is present, the new state may also depend on the privilege selected.

¹¹ This transformation increases the number of states by k times. It is unknown if there exist a transformation that can eliminate exogenous actions without increasing the number of states, and yet is able to model the notion of k -maintainability.

There is a global criterion, telling whether the system as a whole is in a “legitimate” state. It is required that

- (i) in each legitimate state one or more privileges will be present;
- (ii) in each legitimate state each possible move will bring the system to a legitimate state;
- (iii) each privilege must be present in at least one legitimate state; and
- (iv) for any pair of legitimate states there exists a sequence of moves transferring the system from one to the other.

The system is called “self-stabilizing” if and only if, regardless of the initial state and the privilege selected each time for the next move, at least one privilege will always be present and the system is guaranteed to find itself in a legitimate state after a finite number of moves.

Comparison to our notion of maintainability Our notion of a state corresponds to Dijkstra’s global state (which is composed of the local states), and our goal states are his legitimate states. Analogous to our actions are his “moves”, and analogous to our policies are his selections of privileges that are present. Our policies select an action that is to be executed based on the current state. In Dijkstra’s model, although “moves” are selected based on the current global state, the selection is finer; one of the privileges that are present is selected. The transition due to moves is similar to our transition between states due to action. But Dijkstra’s transition is also finer grained because the changes happen only to the local state of the machine whose privilege is selected, and this change depends on the local state (prior to the move) of the machine, the local state of the neighbors, and the privilege selected. We have a specific set of initial states and we use exogenous actions to compute a closure of states that may be reached. In Dijkstra’s definition of self-stabilization all states are possible initial states. This is the same as having the closure to be the set of all states. Modulo the above differences, Dijkstra’s notion of self-stabilization is the same as our notion of finite maintainability.

If one thinks in terms of global states, then essentially self-stabilization is the same as finite maintainability. Our major contribution is that, unlike research in the area of self-stabilization protocols, where such protocols are invented by people and self-stabilizability of systems using those protocols is proven, we give a method to automatically come up with the protocols (or “policies” in our terms). Nonetheless, a large body of research in self-stabilization is with respect to distributed systems where one does not have access to the global state in deciding what action to take. Thus a direction for future research is to find algorithms similar to ours that can generate protocols (policies) in the distributed domain.

A.2. Arora and Gouda’s notion of closure, convergence, and fault-tolerance

Arora and Gouda’s abstraction [2] is more closer to ours than Dijkstra’s. We start with their definitions and terminology.

- A program consists of a set of variables and a set of processes.
- Each process consists of a set of condition-statement pairs (B, st) , where B is a Boolean expression over program variables, and st updates zero or more program variables and always terminates upon execution.
- A state of a program p is defined by the value for each variable of p .
- For any process of a program p , using the set of condition-statement pairs we can construct a mapping from states to sets of statements, where s maps to a set containing st iff there is a (B, st) in the process such that B evaluates to true in s . This mapping is referred to as the policy corresponding to that process. If a program has a single process, then the policy corresponding to that process is also referred to as the policy of that program.
- A state predicate of p is a Boolean expression over the variables of p .
- A condition-statement pair (B, st) is enabled at a state iff B evaluates to true at that state.
- A process is said to be enabled at a state iff some pair (B, st) of that process is enabled at that state.
- Closure: A state predicate S is closed in a program p iff for each pair (B, st) in each process of p , executing st starting from a state where $B \wedge S$ holds results in a state where S holds.
- A computation of p is a sequence of states that satisfies the following three conditions:
 - (i) for each pair of consecutive states c, d in the sequence, there exists a pair (B, st) in some process of p such that B holds at c and executing st starting from c results in d ;
 - (ii) the sequence is maximal, i.e., the sequence is either infinite or it is finite and no condition-statement pair is enabled in the last state; and
 - (iii) if any process j of p is continuously enabled along the sequence, then eventually some action of j is chosen for execution.
- Convergence: Let S and T be state predicates of p . T is said to converge to S in p iff
 - (a) S and T are closed in p and
 - (b) in each computation of p starting at any state where T holds, there exists a state where S holds.
- Fault-tolerance: Let S be a closed state predicate of p and let F be a set of condition-statement pairs on variables of p . Then p is F -tolerant of S iff there exists a state predicate T of p such that:
 - (a) T holds at every state where S holds;
 - (b) for each pair (B, st) in F , executing st starting from a state where $B \wedge T$ holds results in a state where T holds; and
 - (c) T converges to S in p .

The notions of closure, convergence, and fault tolerance of Arora and Gouda are very close to similar notions in our definition of maintainability. We now compare their notions to ours more in detail.

Comparison to our notions Arora and Gouda's notion of states is similar to our notion of states, and their notion of statements is similar to our notion of actions, except that statements cause deterministic transition between states and are executable in all states, while actions can have non-deterministic effects and are executable in some states only. Their notion of a process (consisting of a set of condition-statement pairs) is similar to our policy, and their notion of a program consists of a set of variables (that define the states) and a set of processes. Their notion of closure is the same as our notion of closure when we assume the absence of exogenous actions and the presence of a single process. (We assume the presence of a single process in the rest of our comparison.) Their notion of a computation is similar to our notion of a sequence of states obtained using the Unfold function, except that we have no fairness condition, which we do not need since we have only a single policy under consideration) and that policy dictates the execution of only one action in any state.

Arora and Gouda's notion of " T converges to S " (when there is a single process) is similar to finite maintainability (in the absence of exogenous actions), where T and S correspond to the closure of our initial and final states, respectively. In their notion of fault tolerance, captured by " p is F -tolerant of S ," the fault set F corresponds to our exogenous actions and when they can occur, S corresponds to the set of final states, and p reflects the state-space and the policy under consideration. In Arora and Gouda's definition of fault tolerance, T corresponds to our closure of a set of initial states, but has the additional requirement that the set of final states S is a subset of T . Thus " p is F -tolerant of S " corresponds to our notion of the policy of p maintaining a set of initial states (whose closure contains S) with respect to S in presence of exogenous actions described by F .

While our notion of maintainability is similar to the notion of tolerance under a class of faults in [2], Arora and Gouda do not give an algorithm to automatically construct a policy (or process in their terms) that will result in the tolerance, while we give algorithms which generate policies, if such policies exist, that guarantee maintainability.

References

- [1] M. Abadi, L. Lamport, P. Wolper, Realizable and unrealizable specifications of reactive systems, in: Proc. 16th International Conference on Automata, Languages, and Programming (ICALP 89), in: LNCS, vol. 372, Springer, 1989, pp. 1–17.
- [2] A. Arora, M.G. Gouda, Closure and convergence: A foundation of fault-tolerant computing, IEEE Transactions on Software Engineering 19 (11) (1993) 1015–1027.
- [3] F. Bacchus, F. Kabanza, Planning for temporally extended goals, Annals of Mathematics and Artificial Intelligence 22 (1998) 5–27.
- [4] C. Baral, T. Eiter, J. Zhao, Using SAT and LP to design polynomial-time algorithms for planning in non-deterministic domains, in: Proc. 20th National Conference on Artificial Intelligence (AAAI '05), AAAI Press, 2005, pp. 578–583.
- [5] C. Baral, M. Gelfond, A. Provetti, Representing actions: Laws, observations, and hypothesis, Journal of Logic Programming 31 (1997) 201–243.
- [6] C. Baral, V. Kreinovich, R. Trejo, Computational complexity of planning and approximate planning in the presence of incompleteness, Artificial Intelligence 122 (1–2) (2000) 241–267.
- [7] C. Baral, V. Kreinovich, R. Trejo, Computational complexity of planning with temporal goals, in: B. Nebel (Ed.), Proc. 17th International Joint Conference on Artificial Intelligence (IJCAI-01), Morgan Kaufmann, 2001, pp. 509–514.
- [8] C. Baral, T. Son, Relating theories of actions and reactive control, Electronic Transactions on Artificial Intelligence 2 (3–4) (1998) 211–271.
- [9] C. Baral, J. Zhao, Goal specification in presence of non-deterministic actions, in: R.L. de Mántaras, L. Saitta (Eds.), Proc. 16th European Conference on Artificial Intelligence (ECAI 2004), IOS Press, 2004, pp. 273–277.
- [10] M. Barbeau, F. Kabanza, R. St-Denis, Synthesizing plan controllers using real-time goals, in: Proc. 14th International Joint Conference on Artificial Intelligence (IJCAI-95), 1995, pp. 791–800.
- [11] D. Barrington, N. Immerman, H. Straubing, On uniformity within NC^1 , Journal of Computer and System Sciences 41 (1990) 274–306.
- [12] M. Ben-Ari, Z. Manna, A. Puneli, The temporal logic of branching time, in: Proc. 8th Symposium on Principles of Programming Languages, 1981, pp. 164–176.
- [13] P. Bertoli, A. Cimatti, M. Pistore, Strong cyclic planning under partial observability, in: ECAI, 2006, pp. 580–584.
- [14] P. Bertoli, M. Pistore, Planning with extended goals and partial observability, in: S. Zilberstein, J. Koehler, S. Koenig (Eds.), ICAPS, 2004, pp. 270–278.
- [15] R. Brooks, A robust layered control system for a mobile robot, IEEE Journal of Robotics and Automation 2 (1) (1986) 14–23.
- [16] T. Bylander, The computational complexity of propositional strips planning, Artificial Intelligence 69 (1994) 165–204.
- [17] S. Ceri, J. Widom, Deriving production rules for constraint maintenance, in: P.M.G. Apers, G. Wiederhold (Eds.), Proc. 15th International Conference on Very Large Data Bases (VLDB-90), 1990, pp. 566–577.
- [18] A. Cimatti, M. Pistore, M. Roveri, P. Traverso, Weak, strong, and strong cyclic planning via symbolic model checking, Artificial Intelligence 147 (1–2) (2003) 35–84.
- [19] E. Clarke, E. Emerson, Design and synthesis of synchronization skeletons using branching-time temporal logic, in: Proc. Workshop on Logic of Programs, in: LNCS, vol. 131, Springer, 1981, pp. 52–71.
- [20] E. Clarke, E. Emerson, A. Sistla, Automatic verification of finite-state concurrent systems using temporal logic specifications, ACM Transactions on Programming Languages and Systems 8 (2) (1986) 244–263.
- [21] M. Daniele, P. Traverso, M. Vardi, Strong cyclic planning revisited, in: Proc. 5th European Conference on Planning (ECP'99), in: LNCS/LNAI, vol. 1809, Springer, 1999, pp. 35–48.
- [22] E. Dantsin, T. Eiter, G. Gottlob, A. Voronkov, Complexity and expressive power of logic programming, ACM Computing Surveys 33 (3) (2001) 374–425.
- [23] G. De Giacomo, R. Reiter, M. Soutchanski, Execution monitoring of high-level robot programs, in: Proc. Sixth Conference on Principles of Knowledge Representation and Reasoning (KR-98), 1998, pp. 453–465.
- [24] E.W. Dijkstra, Self-stabilizing systems in spite of distributed control, CACM 17 (11) (1974) 644–843.
- [25] W. Dowling, J.H. Gallier, Linear-time algorithms for testing the satisfiability of propositional Horn theories, Journal of Logic Programming 3 (1984) 267–284.
- [26] M. Drummond, Situation control rules, in: Proc. First International Conference on Principles of Knowledge Representation and Reasoning (KR-89), 1989, pp. 103–113.

- [27] P. Dunne, M. Laurence, M. Wooldridge, Complexity results for agent design problems, *Annals of Mathematics, Computing & Teleinformatics* 1 (1) (2003) 19–36.
- [28] P. Dunne, M. Wooldridge, Optimistic and disjunctive agent design problems, in: C. Castelfranchi, Y. Lésperance (Eds.), *Proc. 7th International Workshop on Intelligent Agents (ATAL VII)*, in: LNCS, vol. 1986, Springer, 2001, pp. 1–14.
- [29] T. Eiter, W. Faber, N. Leone, G. Pfeifer, Declarative problem-solving using the DLV system, in: J. Minker (Ed.), *Logic-Based Artificial Intelligence*, Kluwer, 2000, pp. 79–103.
- [30] E. Emerson, Temporal and modal logics, in: J. van Leeuwen (Ed.), *Handbook of Theoretical Computer Science*, vol. B, Elsevier, 1990, Chapter 16.
- [31] K. Erol, V. Subrahmanian, D. Nau, Complexity, decidability and undecidability results for domain-independent planning, *Artificial Intelligence* 76 (1995) 75–88.
- [32] R.E. Fikes, N.J. Nilsson, Strips: A new approach to the application of theorem proving to problem solving, *Artificial Intelligence* 2 (3–4) (1971) 189–208.
- [33] M. Gelfond, V. Lifschitz, Classical negation in logic programs and disjunctive databases, *New Generation Computing* 9 (1991) 365–385.
- [34] M. Gelfond, V. Lifschitz, Representing action in extended logic programs, in: *Proc. Joint International Conference and Symposium on Logic Programming (JICSLP'92)*, MIT Press, 1992, pp. 559–573.
- [35] M. Ghallab, D. Nau, P. Traverso, *Automated Planning—Theory and Practice*, Morgan Kaufmann, 2004.
- [36] M.L. Ginsberg, Universal planning: An (almost) universally bad idea, *AI Magazine* 10 (4) (1989) 40–44.
- [37] A. Harding, M. Ryan, P.-Y. Schobbens, A new algorithm for strategy synthesis in itl games, in: N. Halbwachs, L.D. Zuck (Eds.), *TACAS*, in: LNCS, vol. 3440, Springer, 2005, pp. 477–492.
- [38] N. Immerman, *Descriptive Complexity*, Springer, 1999.
- [39] R.M. Jensen, M.M. Veloso, M.H. Bowling, OBDD-based optimistic and strong cyclic adversarial planning, in: *Proc. 6th European Conference on Planning (ECP-01)*, 2001.
- [40] R.M. Jensen, M.M. Veloso, R.E. Bryant, Fault tolerant planning: Toward probabilistic uncertainty models in symbolic non-deterministic planning, in: S. Zilberstein, J. Koehler, S. Koenig (Eds.), *Proc. 14th International Conference on Automated Planning and Scheduling (ICAPS 2004)*, 2004, pp. 335–344.
- [41] F. Kabanza, M. Barbeau, R. St-Denis, Planning control rules for reactive agents, *Artificial Intelligence* 95 (1) (1997) 67–113.
- [42] L.P. Kaelbling, S.J. Rosenschein, Action and planning in embedded agents, in: Maes [47], pp. 35–48.
- [43] C. Kuratowski, *Topology I*, Academic Press, New York, 1966.
- [44] U.D. Lago, M. Pistore, P. Traverso, Planning with a language for extended goals, in: *AAAI/IAAI*, 2002, pp. 447–454.
- [45] N. Leone, G. Pfeifer, W. Faber, T. Eiter, G. Gottlob, S. Perri, F. Scarcello, The DLV system for knowledge representation and reasoning, *ACM Transactions on Computational Logic* 7 (3) (2006) 499–562.
- [46] M.L. Littman, Probabilistic propositional planning: Representations and complexity, in: *Proc. 14th National Conference on Artificial Intelligence and 9th Innovative Applications of Artificial Intelligence Conference (AAAI/IAAI 1997)*, 1997, pp. 748–754.
- [47] P. Maes (Ed.), *Designing Autonomous Agents: Theory and Practice from Biology to Engineering and Back*, MIT Press, 1990.
- [48] Z. Manna, A. Pnueli, *The Temporal Logic of Reactive and Concurrent Systems Specification*, Springer, 1992.
- [49] Z. Manna, P. Wolper, Synthesis of communicating processes from temporal logic specifications, *ACM Transactions on Programming Languages and Systems* 6 (1) (1984) 68–93.
- [50] M. Minoux, LTUR: A simplified linear time unit resolution for Horn formulae and computer implementation, *Information Processing Letters* 29 (1988) 1–12.
- [51] M. Nakamura, C. Baral, Invariance, maintenance and other declarative objectives of triggers—a formal characterization of active databases, in: J. Lloyd et al. (Ed.), *Proc. First International Conference on Computational Logic (CL 2000)*, in: LNAI, vol. 1861, Springer, 2000, pp. 1210–1224.
- [52] M. Nakamura, C. Baral, M. Bjærelund, Maintainability: A weaker stabilizability like notion for high level control, in: *Proc. 17th National Conference on Artificial Intelligence and Twelfth Conference on Innovative Applications of Artificial Intelligence (AAAI/IAAI 2000)*, AAAI Press, 2000, pp. 62–67.
- [53] I. Niemelä, P. Simons, T. Syrjänen, Smodels: A system for answer set programming, in: C. Baral, M. Truszczyński (Eds.), *Proc. 8th International Workshop on Non-Monotonic Reasoning (NMR'2000)*, 2000, available at <http://xxx.lanl.gov/abs/cs.AI/0003033>.
- [54] R. Niyogi, S. Sarkar, Logical specification of goals, in: R.K. Ghosh, D. Misra (Eds.), *Proc. 3rd International Conference on Information Technology (CIT 2001)*, Tata McGraw-Hill, 2000, pp. 77–82.
- [55] C. Ortiz, A commonsense language for reasoning about causation and rational action, *Artificial Intelligence* 111 (2) (1999) 73–130.
- [56] O. Ozveren, A. Willsky, P. Antsaklis, Stability and stabilizability of discrete event dynamic systems, *Journal of ACM* 38 (3) (1991) 730–752.
- [57] C.H. Papadimitriou, *Computational Complexity*, Addison-Wesley, 1994.
- [58] K. Passino, K. Burgess, *Stability Analysis of Discrete Event Systems*, John Wiley and Sons, 1998.
- [59] N. Piterman, A. Pnueli, Y. Sa'ar, Synthesis of reactive(1) designs, in: *VMCAI*, 2006, pp. 364–380.
- [60] A. Pnueli, R. Rosner, On the synthesis of a reactive module, in: *Proc. 16th Annual ACM Symposium on Principles of Programming Languages (POPL 1989)*, 1989, pp. 179–190.
- [61] P. Ramadge, W. Wonham, Modular feedback logic for discrete event systems, *SIAM Journal of Control and Optimization* 25 (5) (1987) 1202–1217.
- [62] P. Ramadge, W. Wonham, Supervisory control of a class of discrete event process, *SIAM Journal of Control and Optimization* 25 (1) (1987) 206–230.
- [63] R. Reiter, *Knowledge in Action: Logical Foundation for Describing and Implementing Dynamical Systems*, MIT Press, 2001.
- [64] J. Rintanen, Complexity of planning with partial observability, in: S. Zilberstein, J. Koehler, S. Koenig, (Eds.), *Proc. 14th International Conference on Automated Planning and Scheduling (ICAPS 2004)*, 2004, pp. 345–354.
- [65] P. Simons, I. Niemelä, T. Soinen, Extending and implementing the stable model semantics, *Artificial Intelligence* 138 (2002) 181–234.
- [66] E. Sontag, Stability and stabilization: Discontinuities and the effect of disturbances, in: F. Clarke, R. Sternand (Eds.), *Proc. NATO Advanced Study Institute*, Kluwer, 1998, pp. 551–598.
- [67] D. Weld, O. Etzioni, The first law of robotics (a call to arms), in: *Proc. Twelfth National Conference on Artificial Intelligence (AAAI-94)*, AAAI Press, 1994, pp. 1042–1047.
- [68] J. Widom, S. Ceri (Eds.), *Active Database Systems: Triggers and Rules For Advanced Database Processing*, Morgan Kaufmann, 1996.
- [69] M. Wooldridge, The computational complexity of agent design problems, in: *Proc. Fourth International Conference on Multi-Agent Systems (ICMAS 2000)*, IEEE Press, 2000, pp. 341–348.